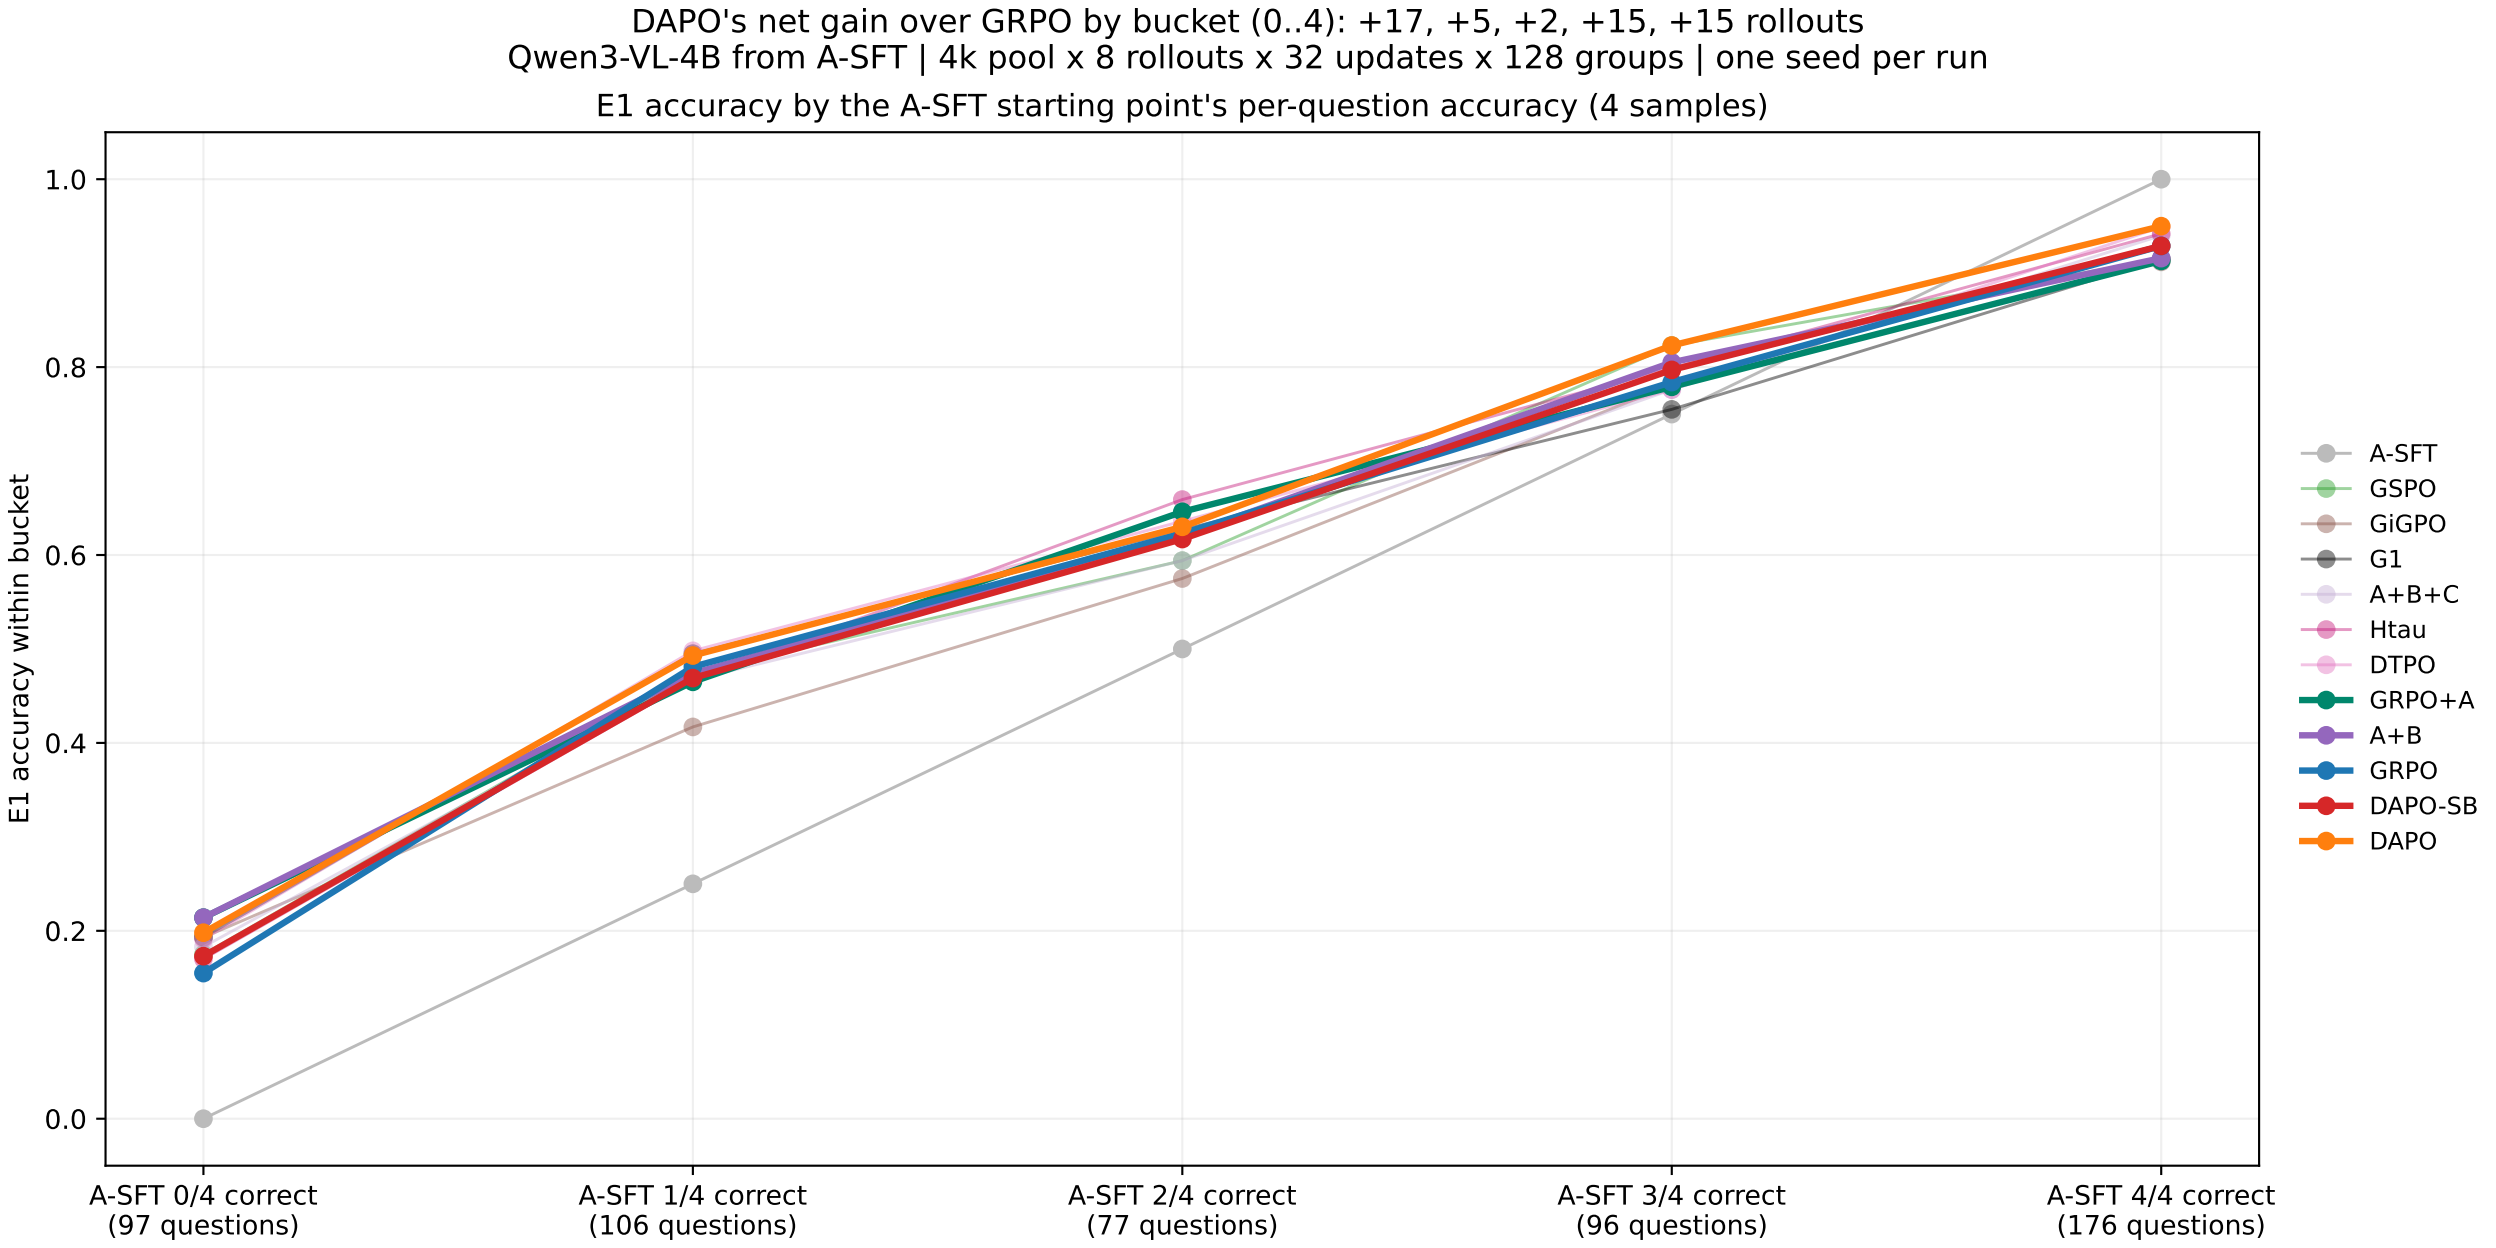
DAPO's net gain over GRPO by bucket (0..4): +17, +5, +2, +15, +15 rollouts
Qwen3-VL-4B from A-SFT | 4k pool x 8 rollouts x 32 updates x 128 groups | one seed per run
E1 accuracy by the A-SFT starting point's per-question accuracy (4 samples)
E1 accuracy within bucket
1.0
0.8
0.6
0.4
0.2
0.0
A-SFT 0/4 correct
(97 questions)
A-SFT 1/4 correct
(106 questions)
A-SFT 2/4 correct
(77 questions)
A-SFT 3/4 correct
(96 questions)
A-SFT 4/4 correct
(176 questions)
A-SFT
GSPO
GiGPO
G1
A+B+C
Htau
DTPO
GRPO+A
A+B
GRPO
DAPO-SB
DAPO