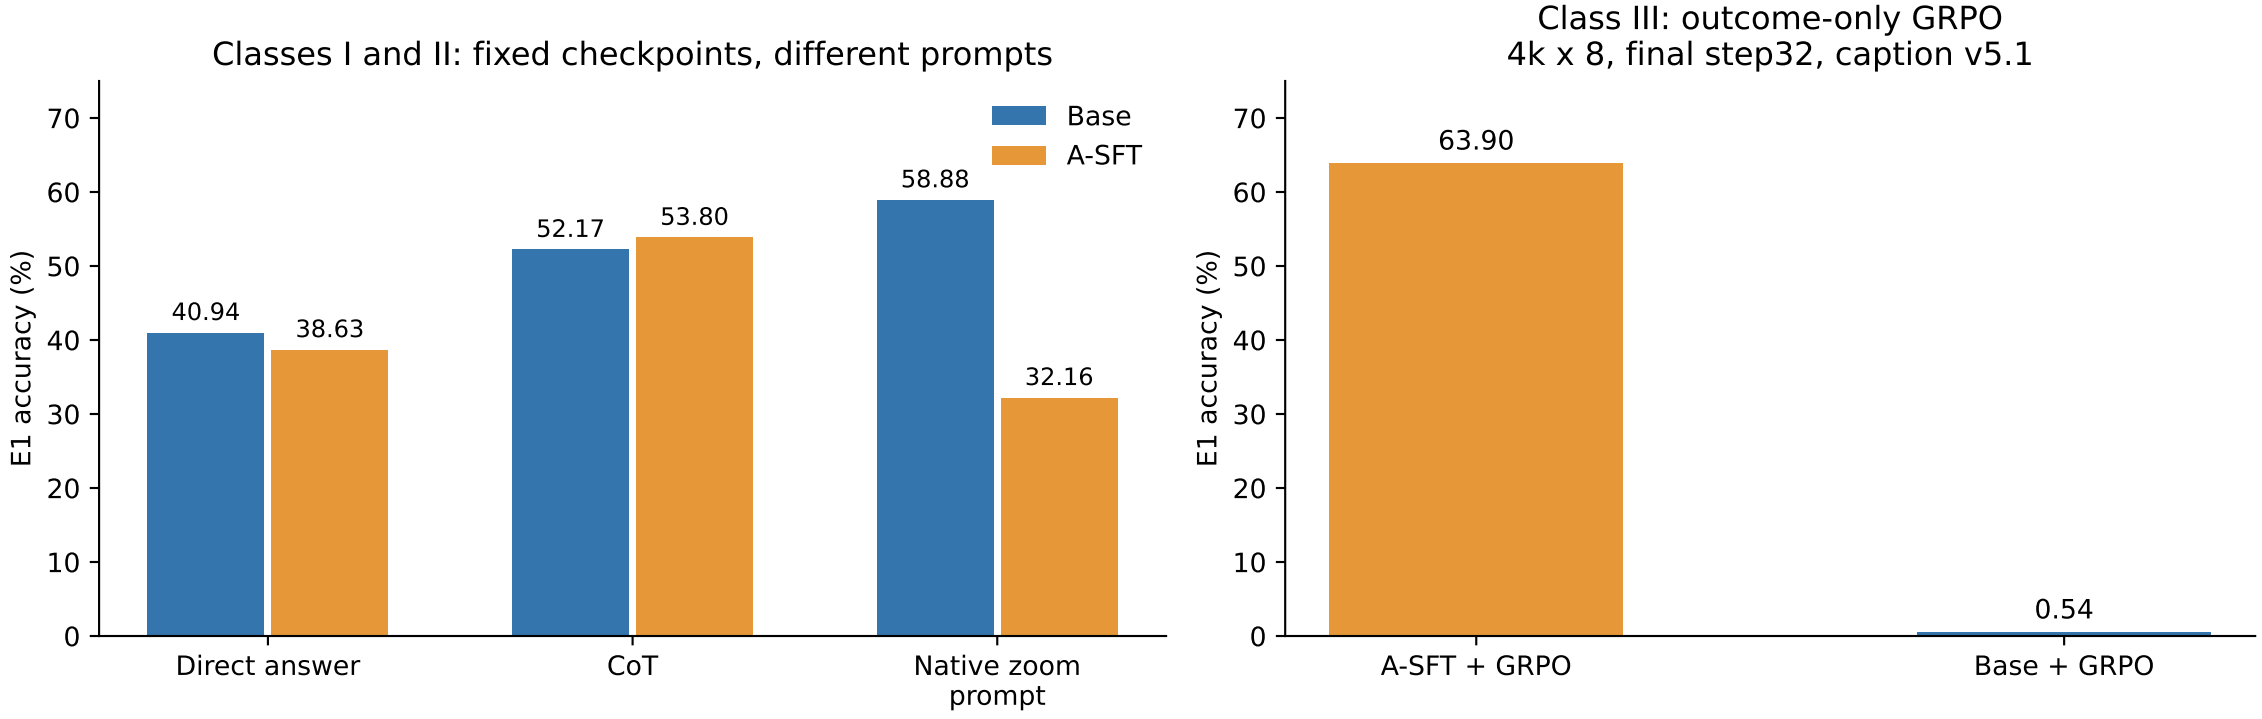
Existing E1 baselines: 552 questions x 4 samples per condition
Classes I and II: fixed checkpoints, different prompts
E1 accuracy (%)
Base
A-SFT
40.94
38.63
52.17
53.80
58.88
32.16
Direct answer
CoT
Native zoom prompt
Class III: outcome-only GRPO
4k x 8, final step32, caption v5.1
63.90
0.54
A-SFT + GRPO
Base + GRPO
All failures retained. E1 is diagnostic, not blind. Native zoom uses model-generated tool text.