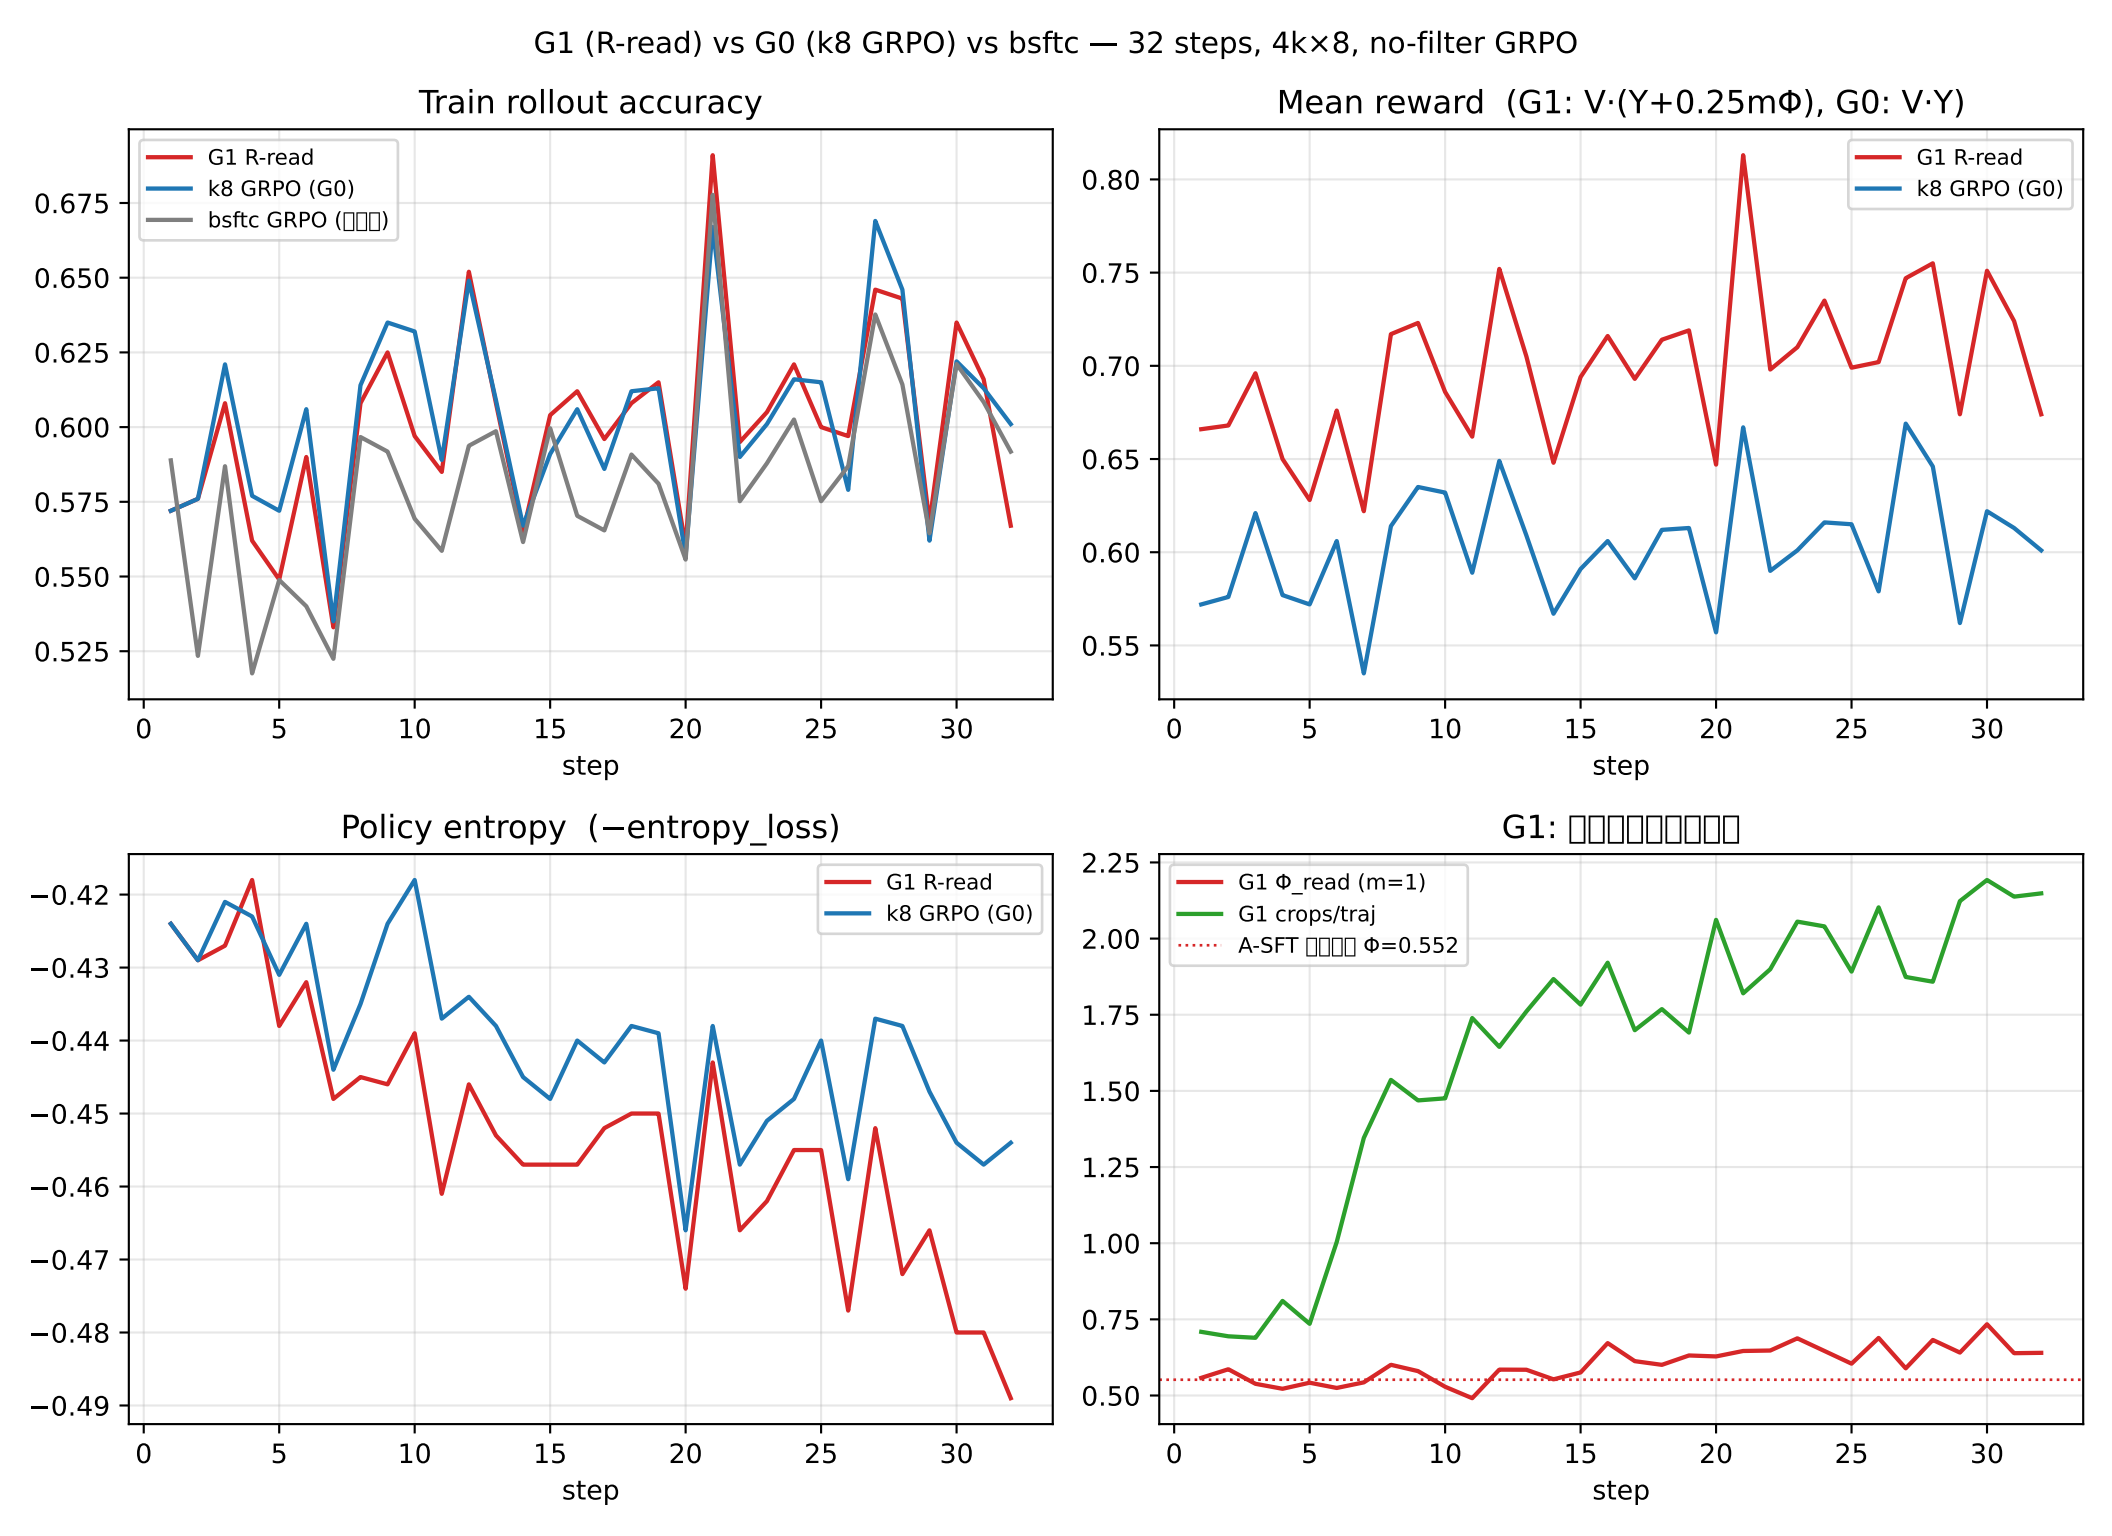
G1 (R-read) vs G0 (k8 GRPO) vs bsftc — 32 steps, 4k×8, no-filter GRPO
Train rollout accuracy
G1 R-read
k8 GRPO (G0)
bsftc GRPO (□□□)
step
Mean reward (G1: V·(Y+0.25mΦ), G0: V·Y)
G1 R-read
k8 GRPO (G0)
step
Policy entropy (−entropy_loss)
G1 R-read
k8 GRPO (G0)
step
G1: □□□□□□□□□
G1 Φ_read (m=1)
G1 crops/traj
A-SFT □□□□ Φ=0.552
step