G1 (R-read): reading reward is learned (Phi .55 -> .65, same-question probe .69) but end-task accuracy is flat (E1 .6395 = G0 .6390)

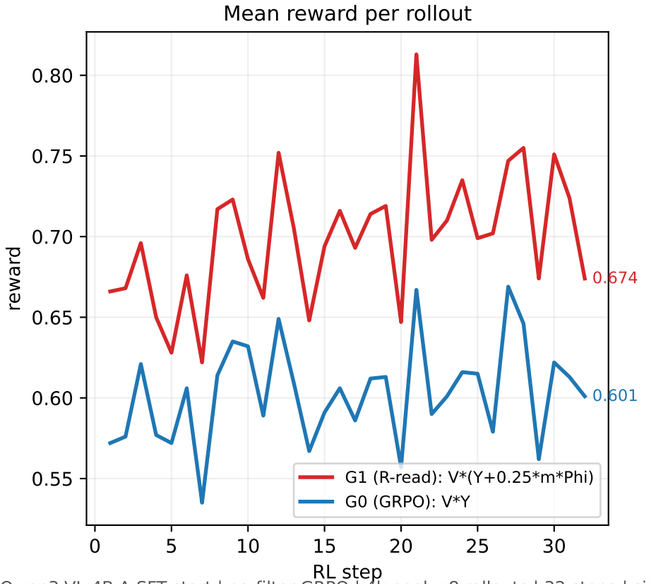

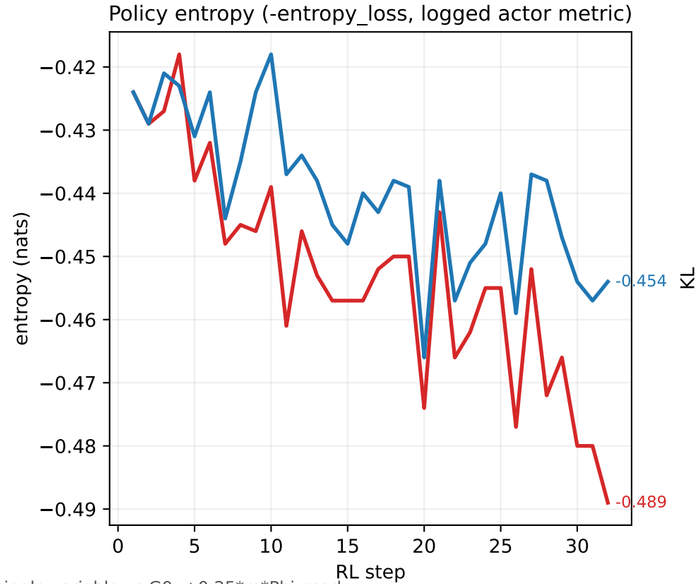

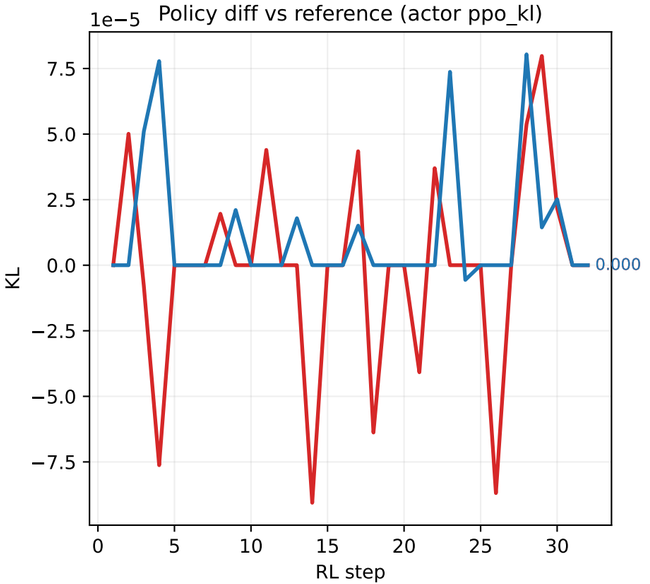

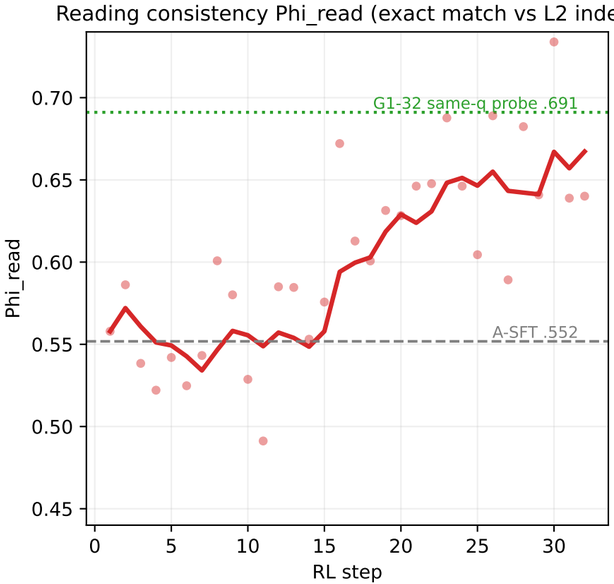

Qwen3-VL-4B A-SFT start | no filter GRPO | 4k pool x 8 rollouts | 32 steps | single variable vs G0: +0.25*m*Phi_read