NATIVE ZOOM / A-SFT EVALUATION

Coordinate normalization improved accuracy; a larger generation budget left correctness unchanged.

Same checkpoint and E1 questions • 552 questions × 4 samples per condition

A Answer accuracy

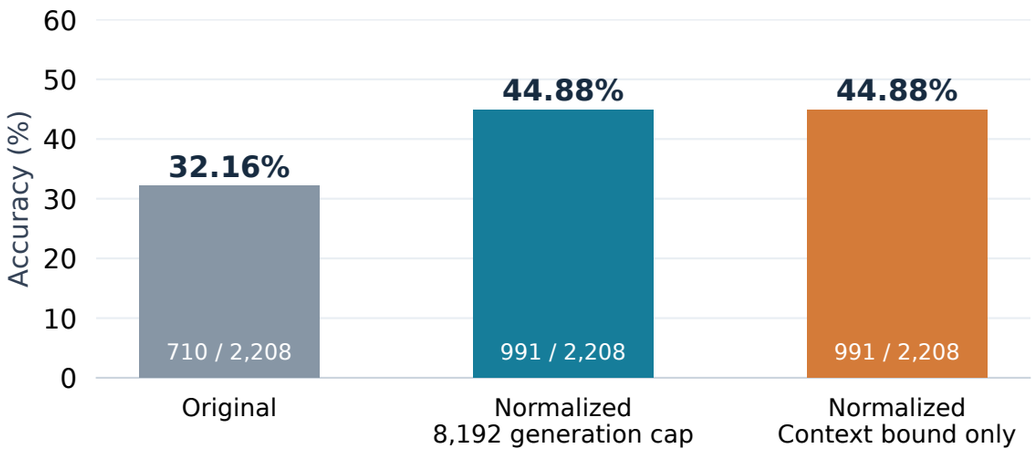

B Executed crops

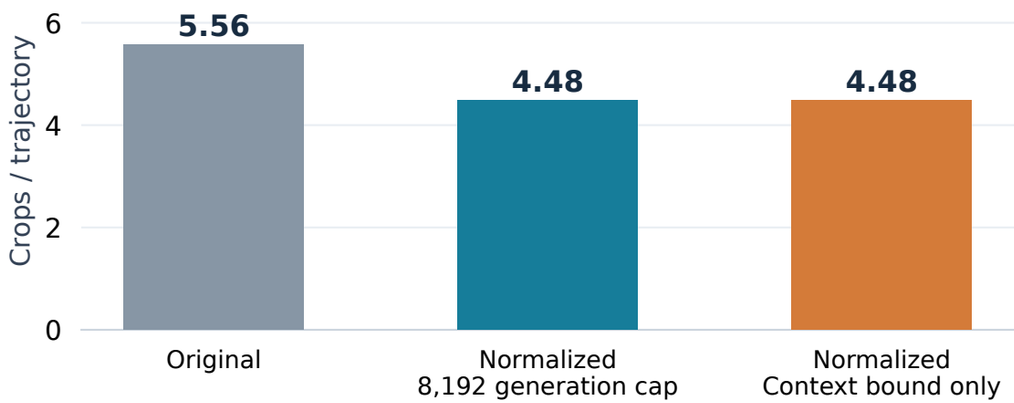

C Generated tokens

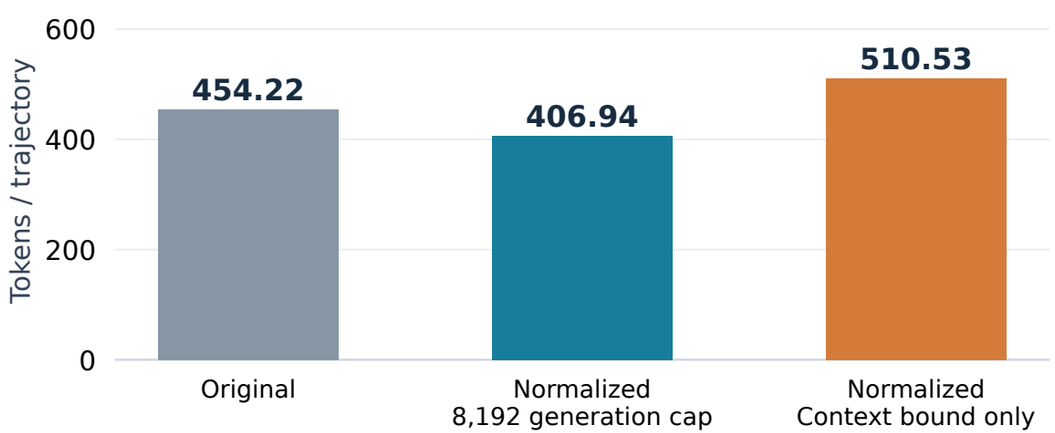

D Total processed tokens

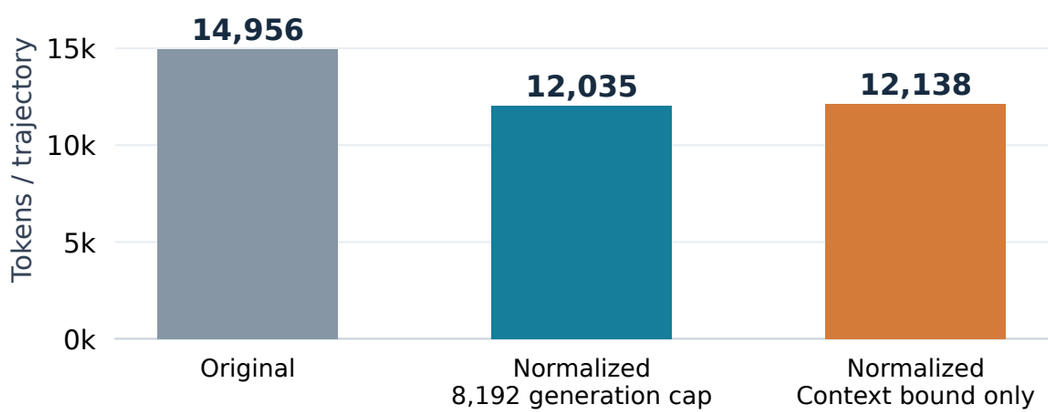

All conditions retain 8 crops and a 32,768-token context. Failures remain in the denominator; E1 is not blind.

Processed tokens include repeated prompts, images and generation across turns. Context-bound generation is not unlimited.