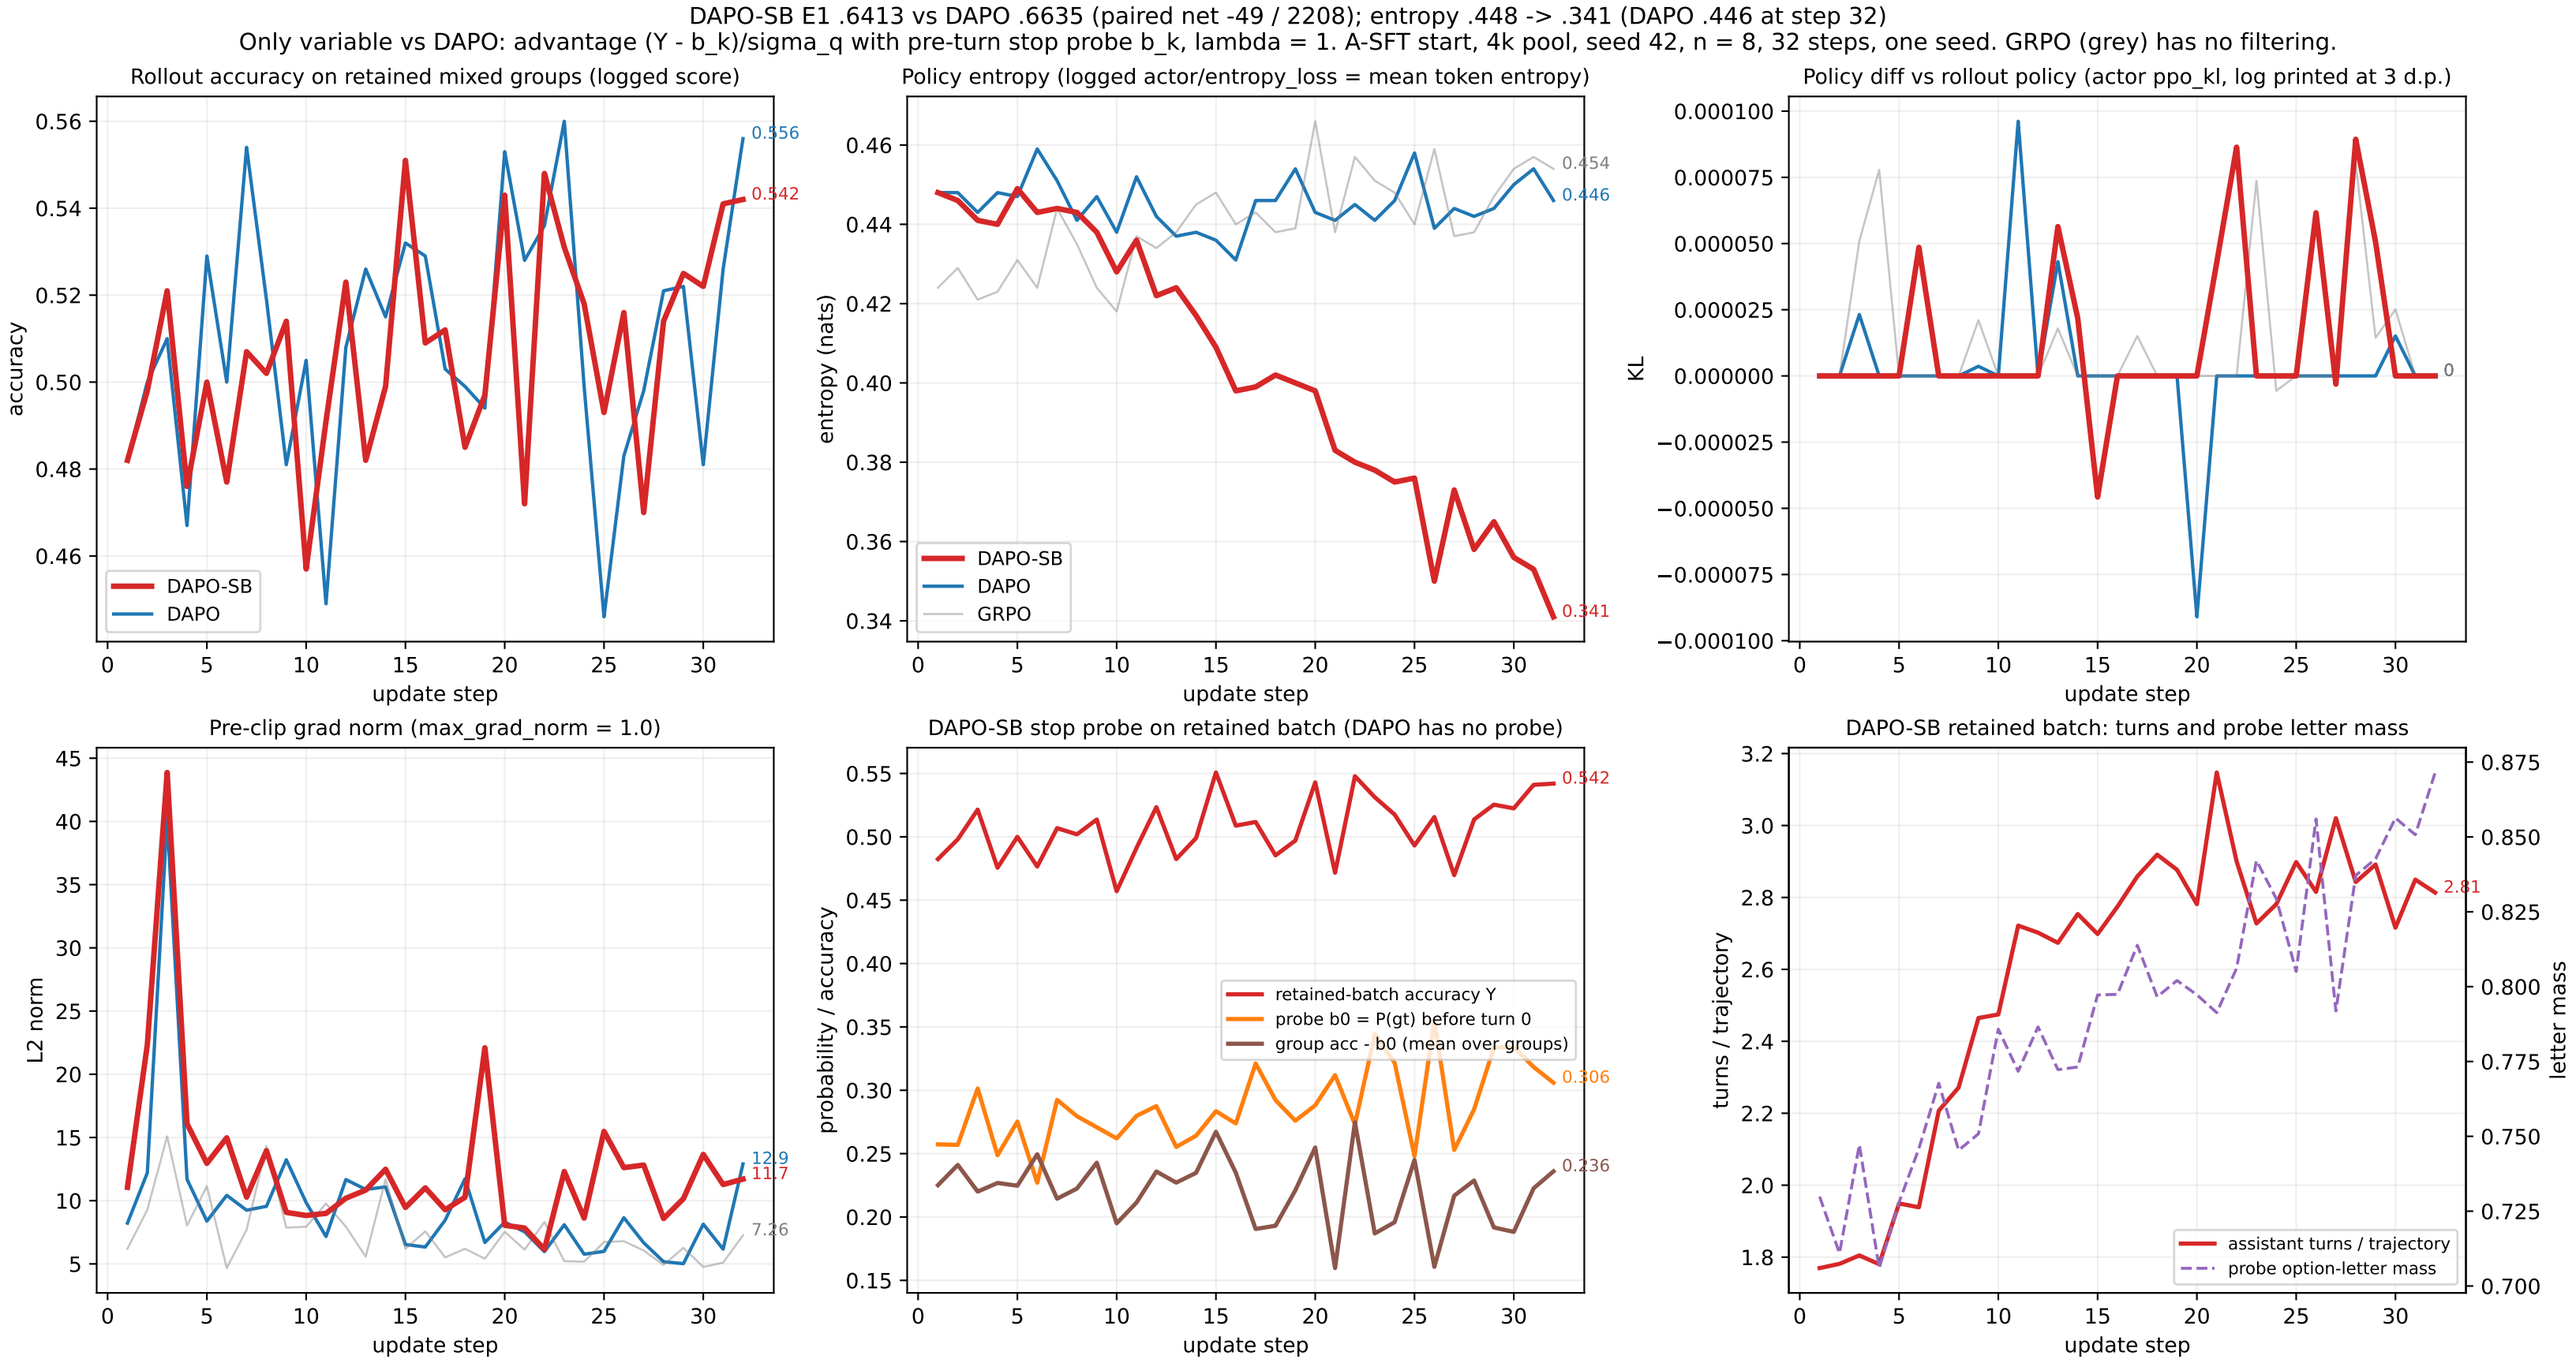
DAPO-SB E1 .6413 vs DAPO .6635 (paired net -49 / 2208); entropy .448 -> .341 (DAPO .446 at step 32)
Only variable vs DAPO: advantage (Y - b_k)/sigma_q with pre-turn stop probe b_k, lambda = 1. A-SFT start, 4k pool, seed 42, n = 8, 32 steps, one seed. GRPO (grey) has no filtering.
Rollout accuracy on retained mixed groups (logged score)
accuracy
update step
DAPO-SB
DAPO
0.556
0.542
Policy entropy (logged actor/entropy_loss = mean token entropy)
entropy (nats)
update step
DAPO-SB
DAPO
GRPO
0.454
0.446
0.341
Policy diff vs rollout policy (actor ppo_kl, log printed at 3 d.p.)
KL
update step
0
Pre-clip grad norm (max_grad_norm = 1.0)
L2 norm
update step
12.9
11.7
7.26
DAPO-SB stop probe on retained batch (DAPO has no probe)
probability / accuracy
update step
retained-batch accuracy Y
probe b0 = P(gt) before turn 0
group acc - b0 (mean over groups)
0.542
0.306
0.236
DAPO-SB retained batch: turns and probe letter mass
turns / trajectory
letter mass
update step
assistant turns / trajectory
probe option-letter mass
2.81