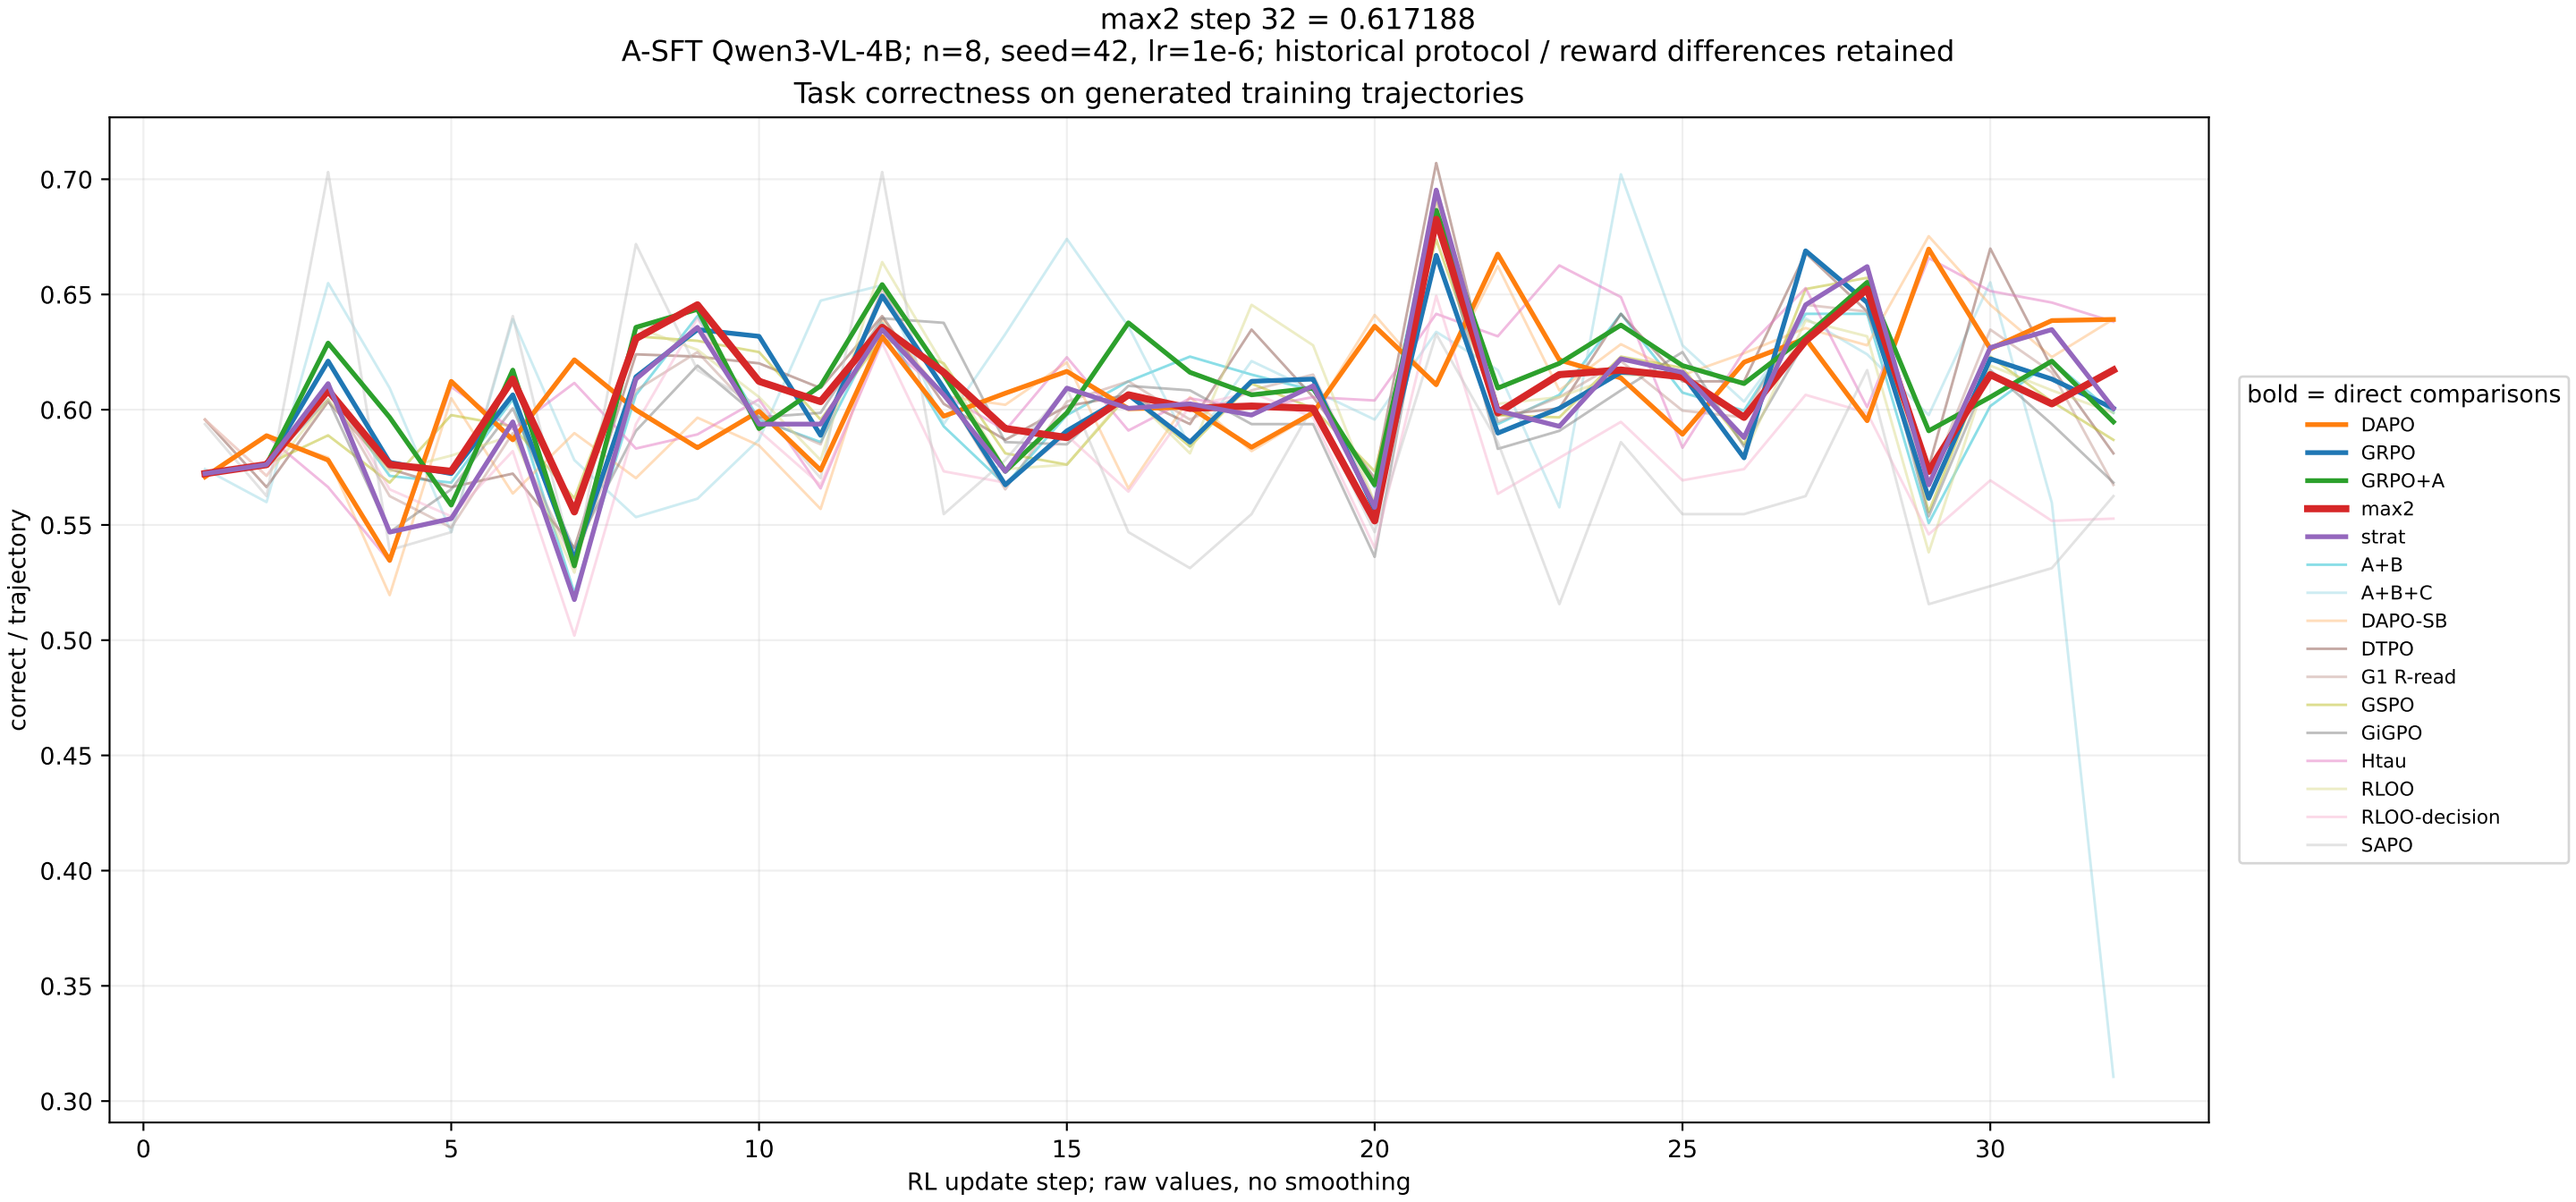
max2 step 32 = 0.617188
A-SFT Qwen3-VL-4B; n=8, seed=42, lr=1e-6; historical protocol / reward differences retained
Task correctness on generated training trajectories
correct / trajectory
0.70
0.65
0.60
0.55
0.50
0.45
0.40
0.35
0.30
0
5
10
15
20
25
30
RL update step; raw values, no smoothing
bold = direct comparisons
DAPO
GRPO
GRPO+A
max2
strat
A+B
A+B+C
DAPO-SB
DTPO
G1 R-read
GSPO
GiGPO
Htau
RLOO
RLOO-decision
SAPO