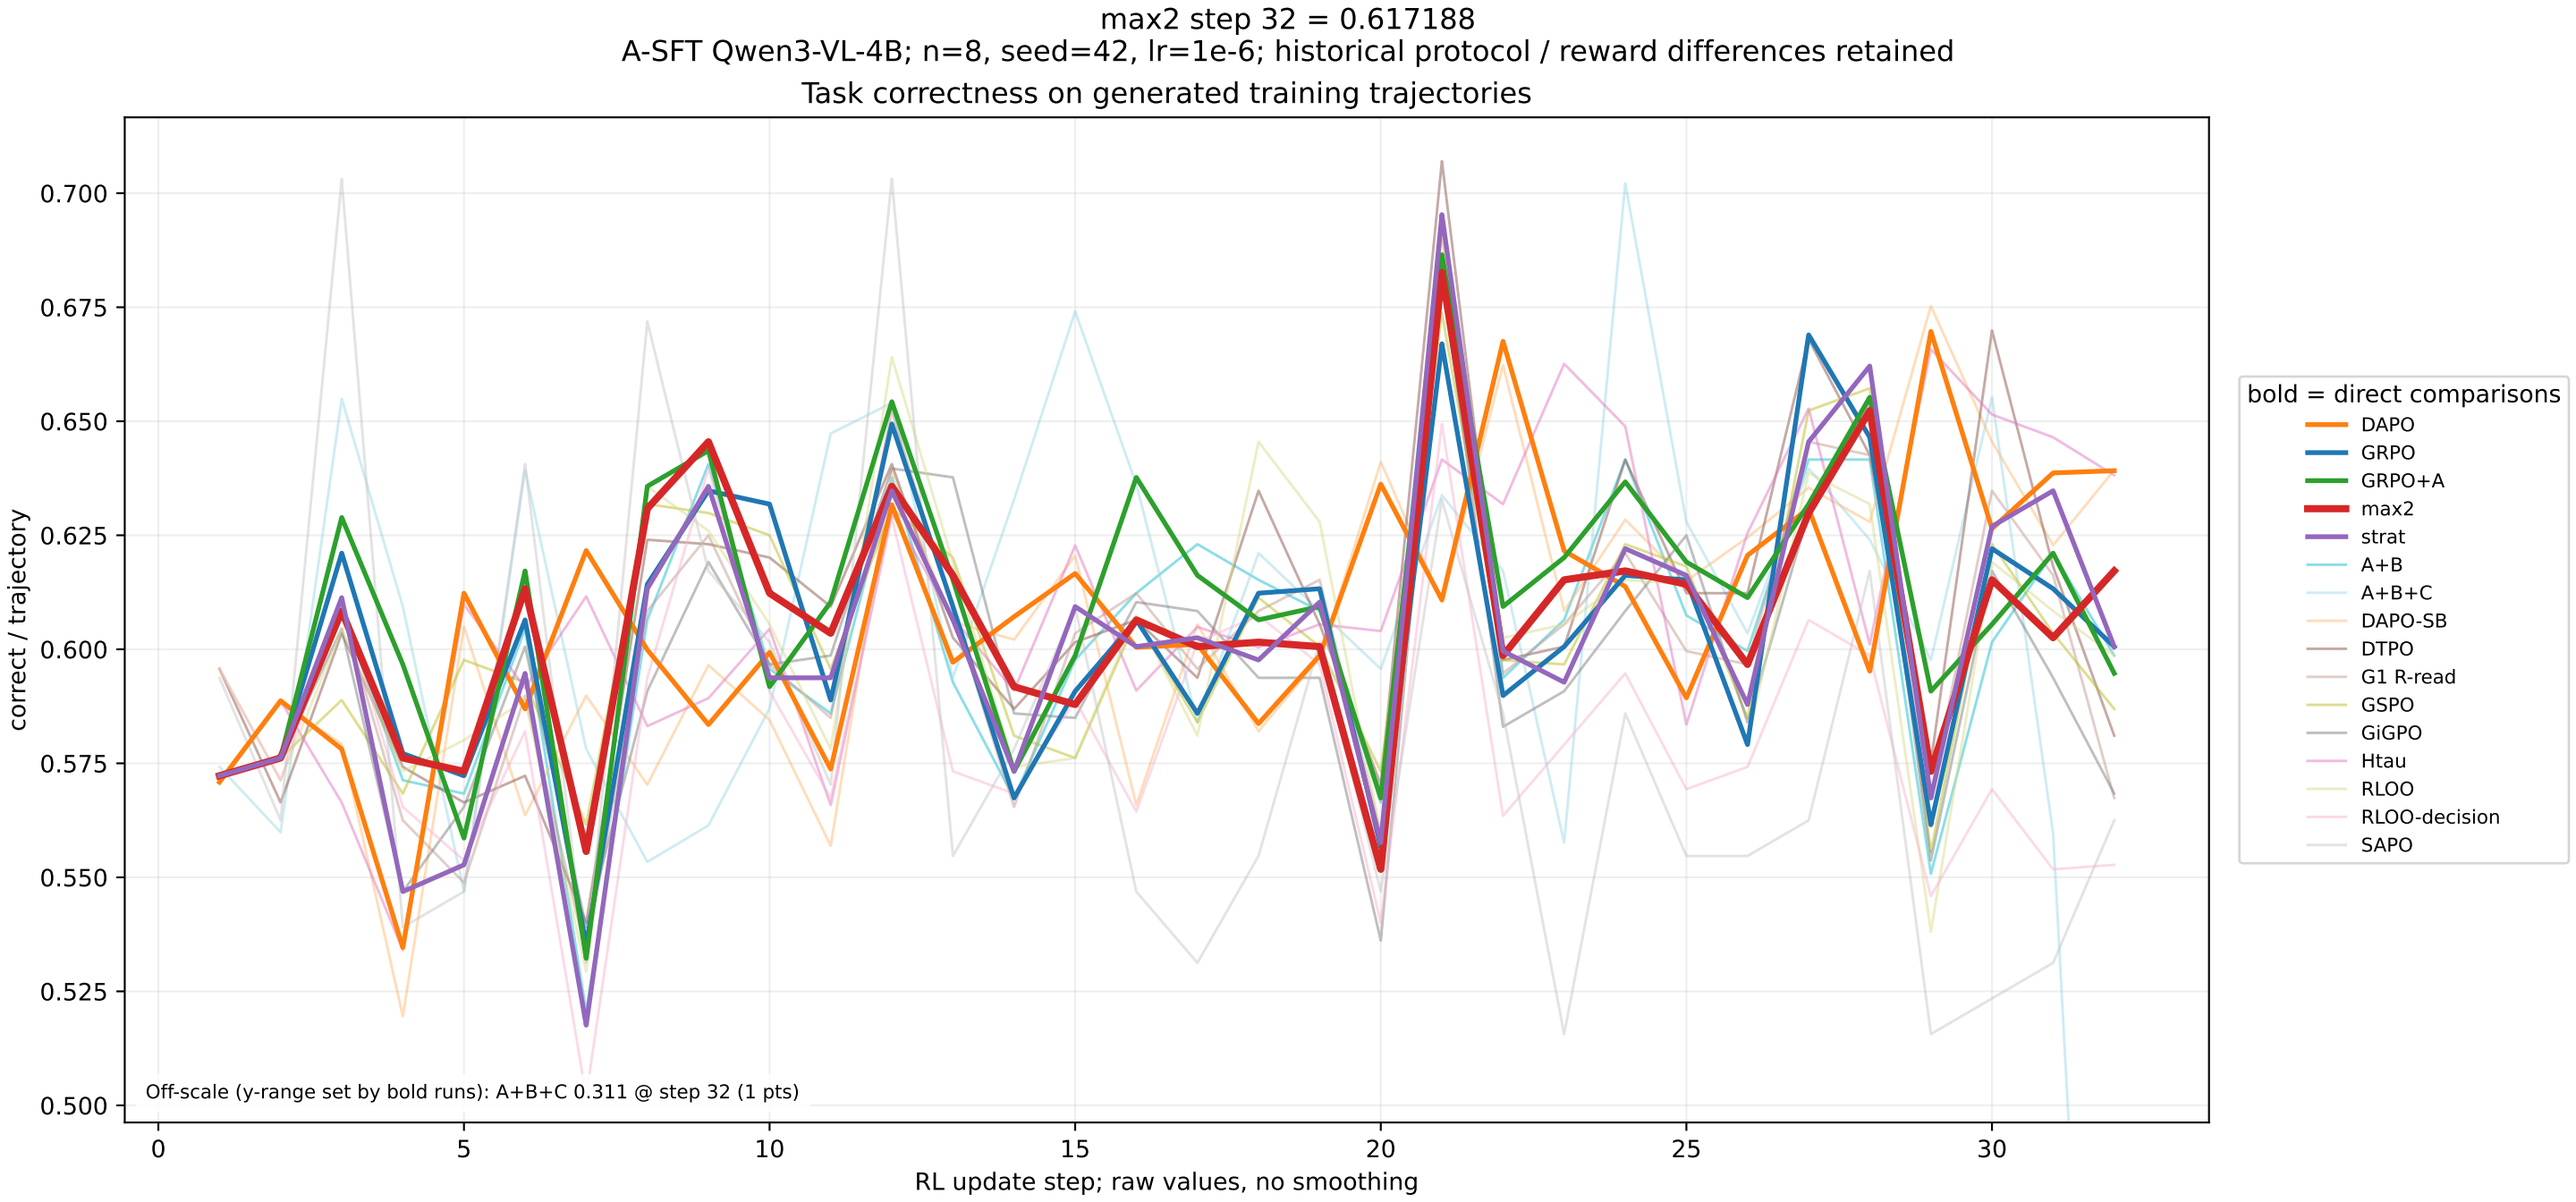
max2 step 32 = 0.617188
A-SFT Qwen3-VL-4B; n=8, seed=42, lr=1e-6; historical protocol / reward differences retained
Task correctness on generated training trajectories
correct / trajectory
RL update step; raw values, no smoothing
0.700
0.675
0.650
0.625
0.600
0.575
0.550
0.525
0.500
0
5
10
15
20
25
30
Off-scale (y-range set by bold runs): A+B+C 0.311 @ step 32 (1 pts)
bold = direct comparisons
DAPO
GRPO
GRPO+A
max2
strat
A+B
A+B+C
DAPO-SB
DTPO
G1 R-read
GSPO
GiGPO
Htau
RLOO
RLOO-decision
SAPO