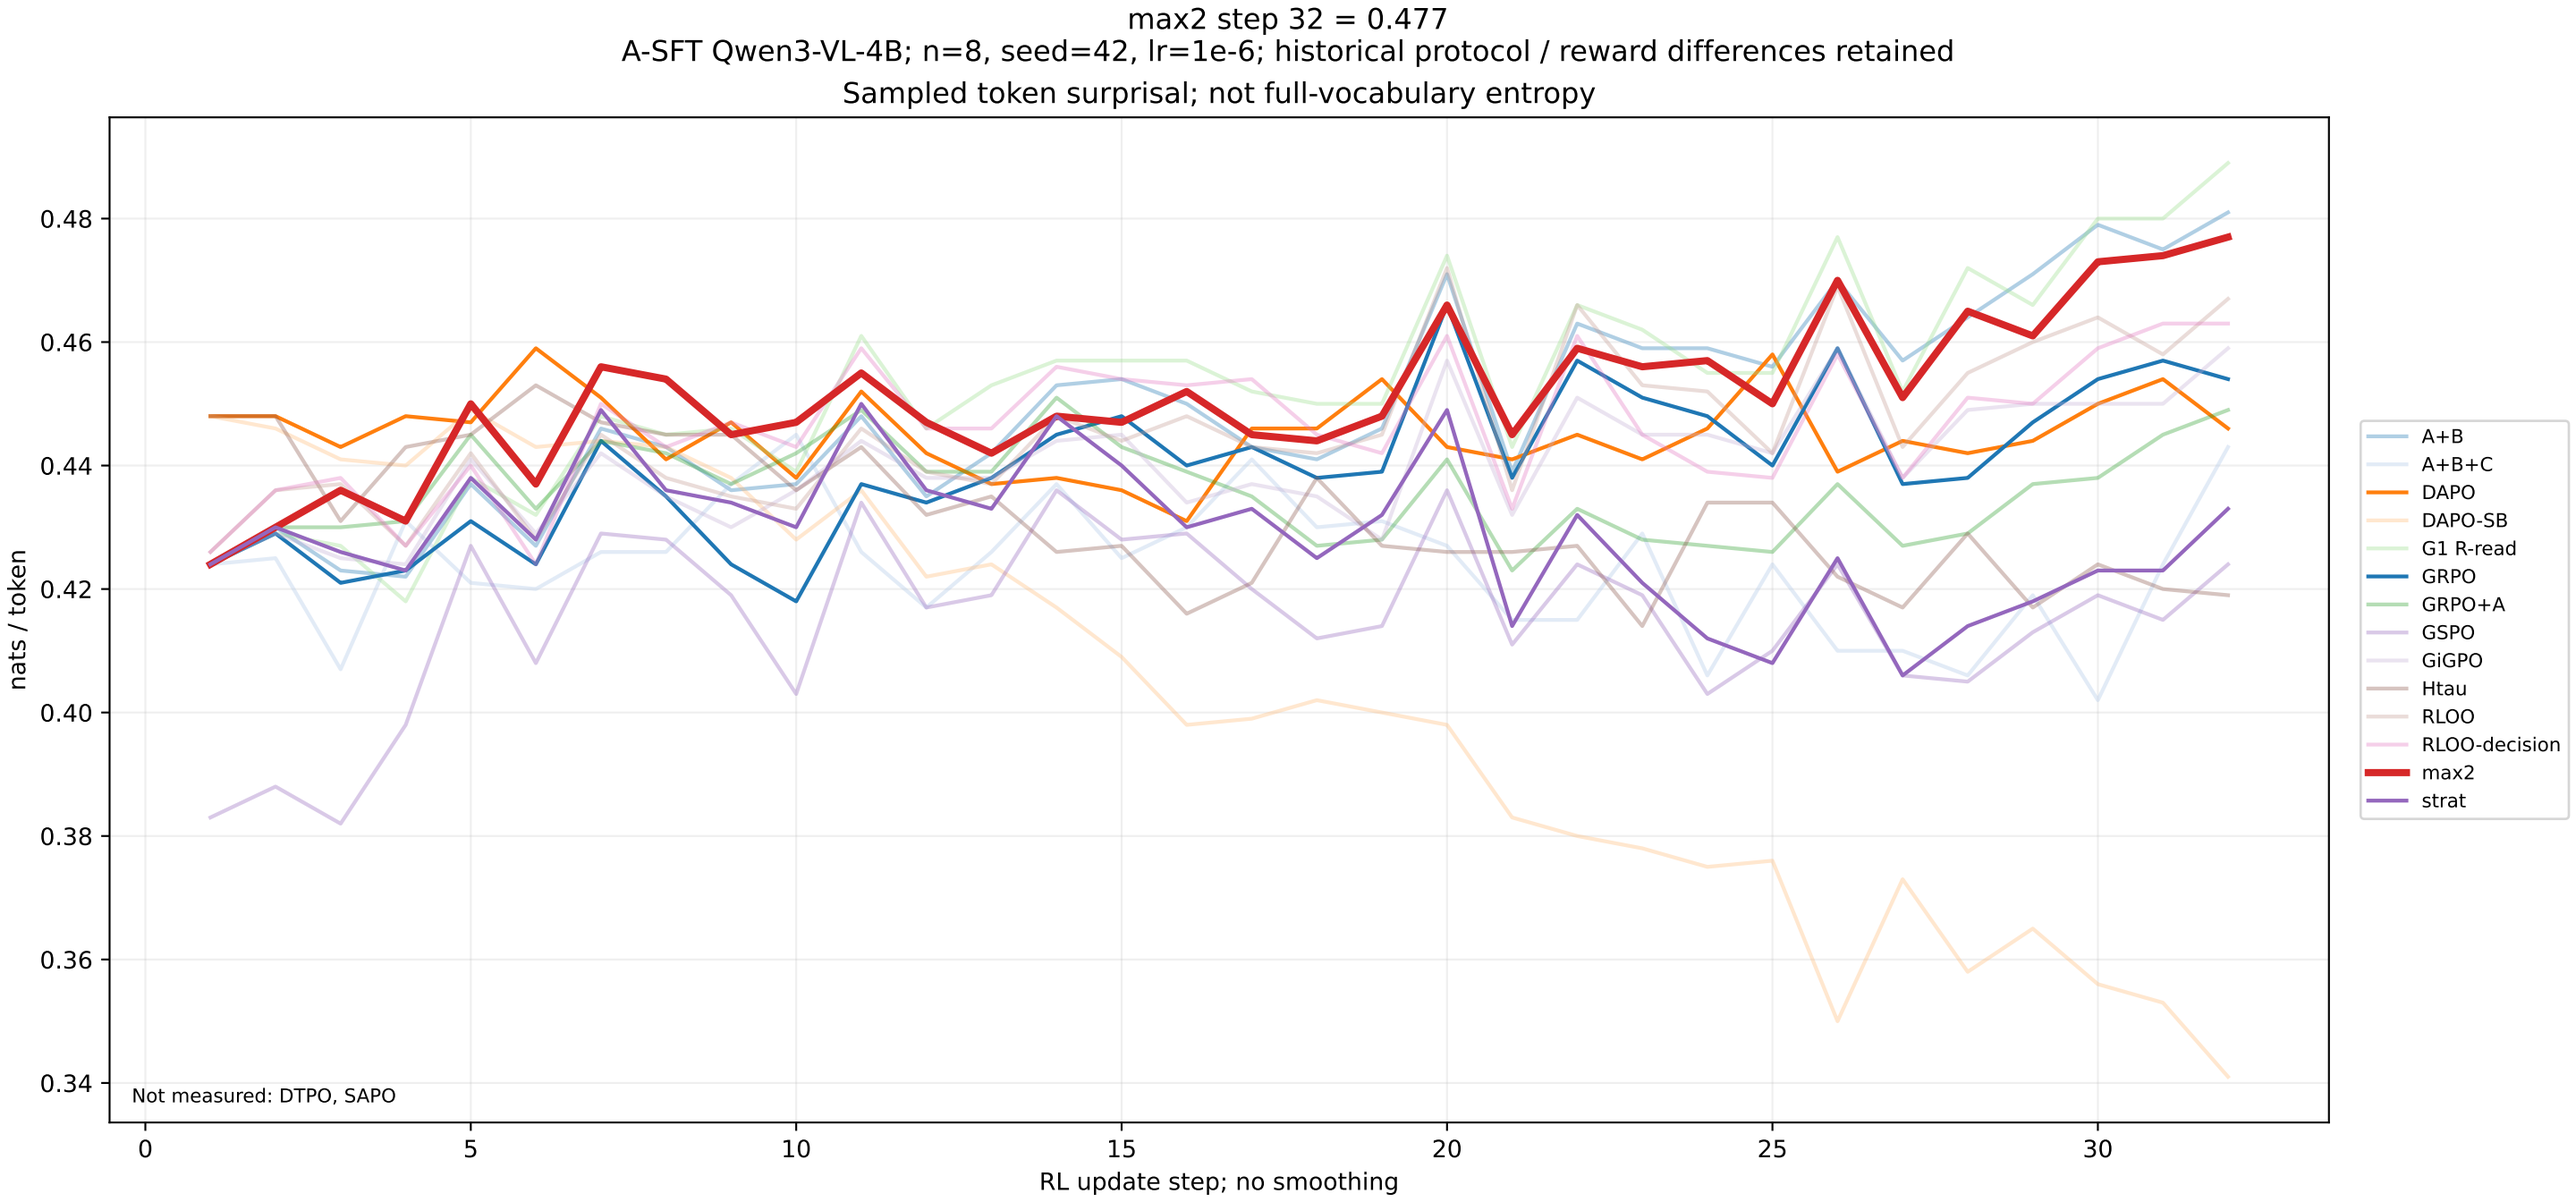
max2 step 32 = 0.477
A-SFT Qwen3-VL-4B; n=8, seed=42, lr=1e-6; historical protocol / reward differences retained
Sampled token surprisal; not full-vocabulary entropy
nats / token
0.48
0.46
0.44
0.42
0.40
0.38
0.36
0.34
0
5
10
15
20
25
30
RL update step; no smoothing
Not measured: DTPO, SAPO
A+B
A+B+C
DAPO
DAPO-SB
G1 R-read
GRPO
GRPO+A
GSPO
GiGPO
Htau
RLOO
RLOO-decision
max2
strat