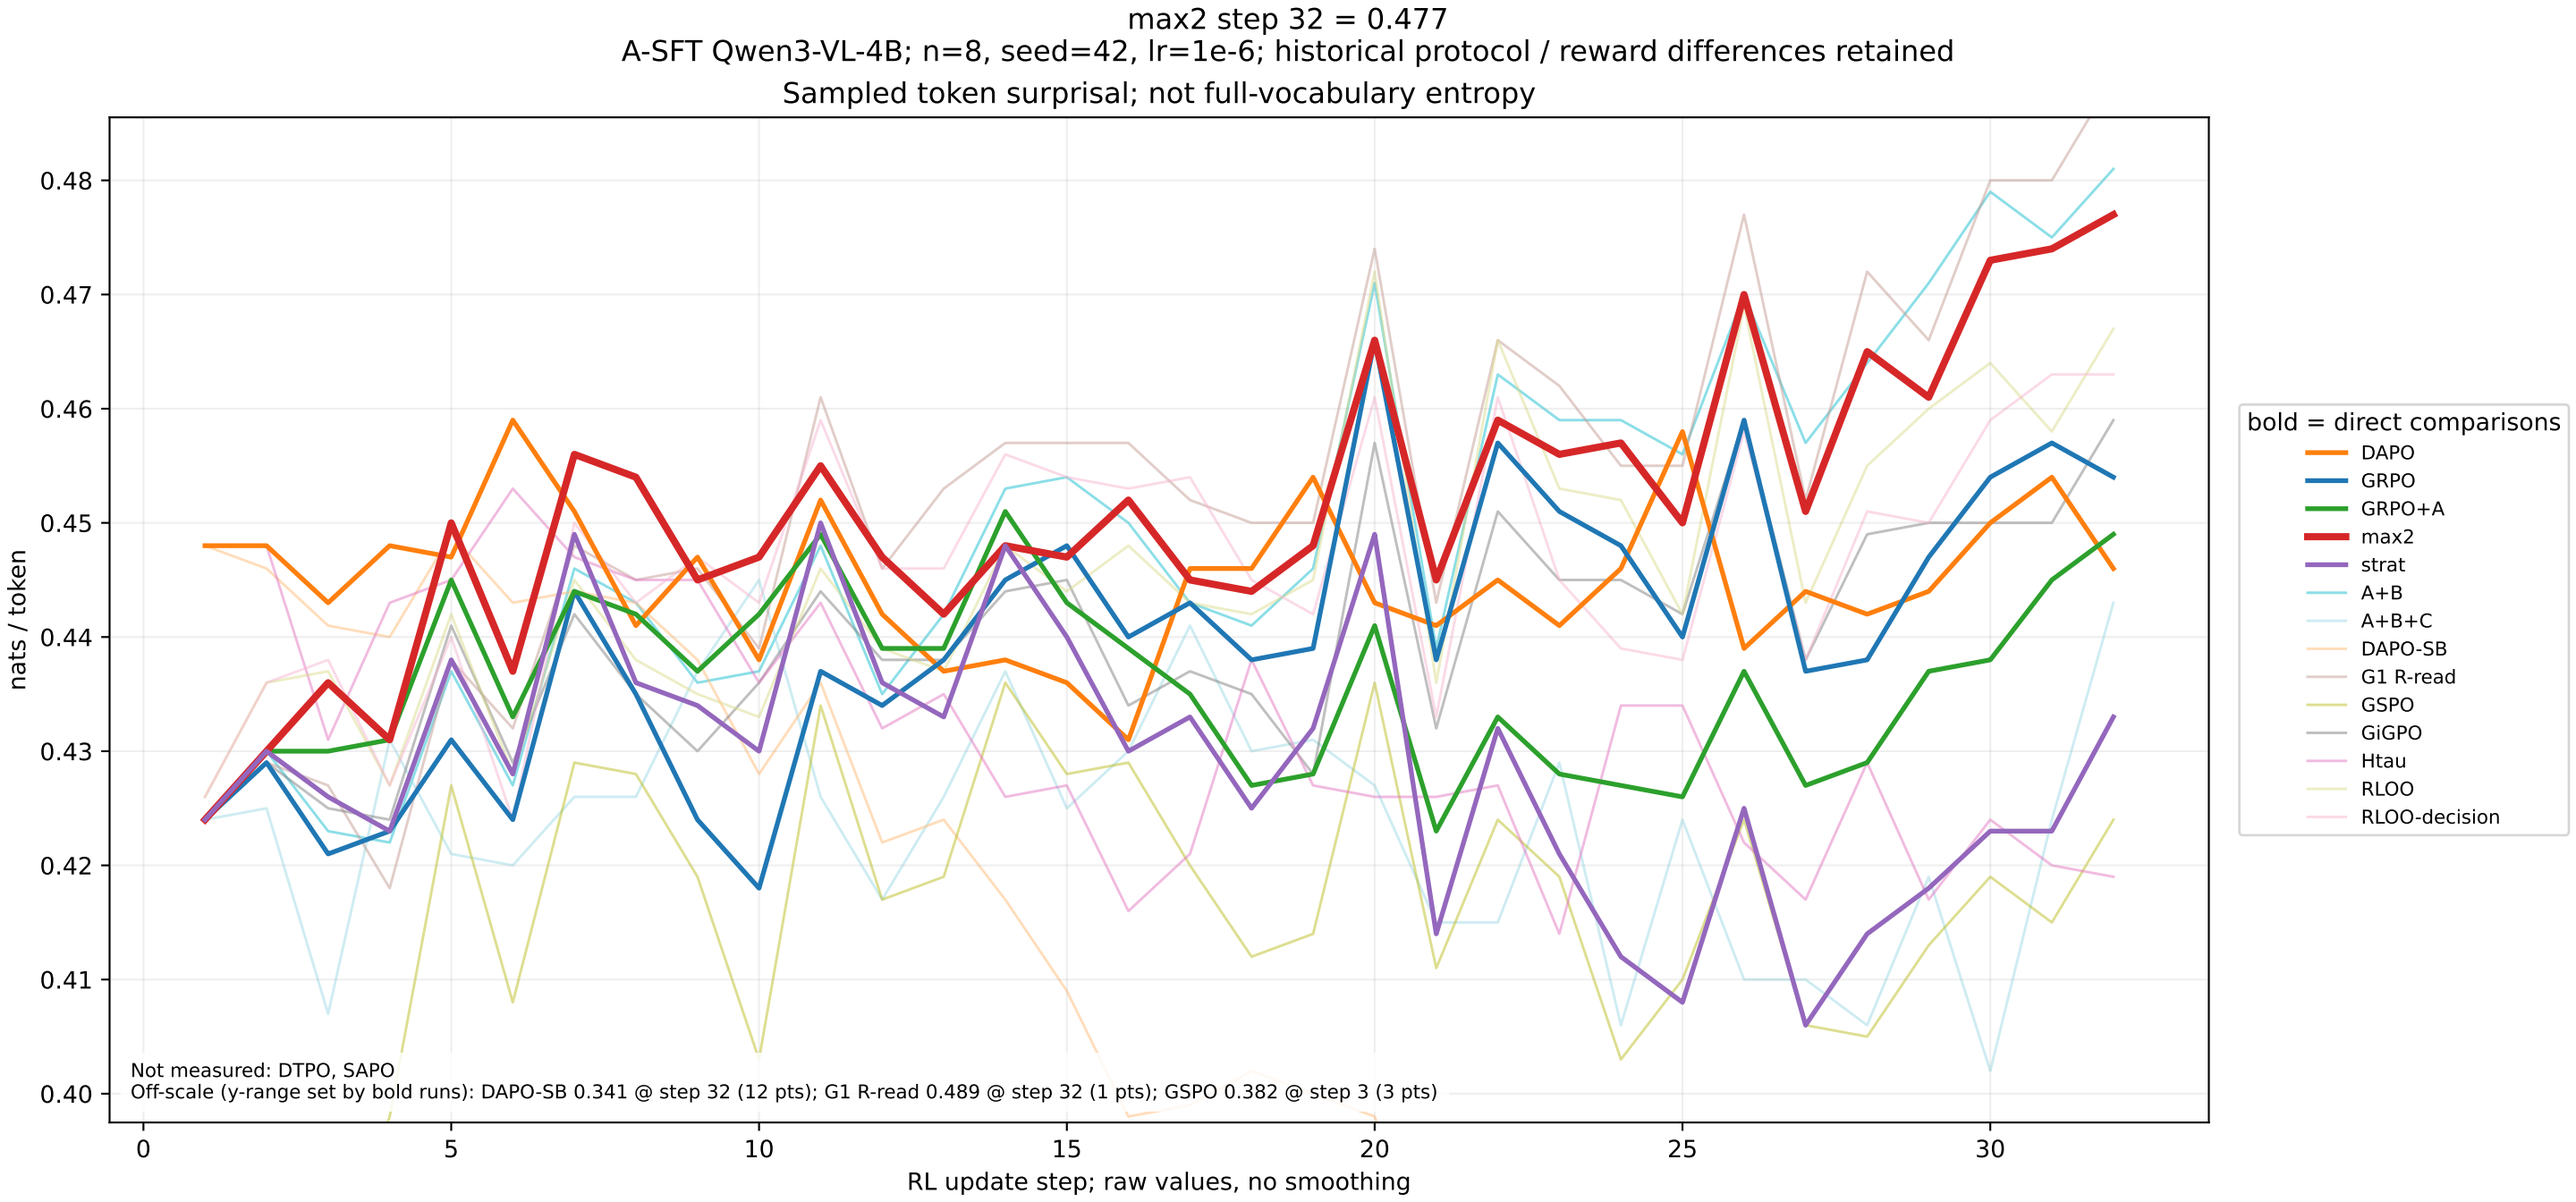
max2 step 32 = 0.477
A-SFT Qwen3-VL-4B; n=8, seed=42, lr=1e-6; historical protocol / reward differences retained
Sampled token surprisal; not full-vocabulary entropy
nats / token
RL update step; raw values, no smoothing
0.40
0.41
0.42
0.43
0.44
0.45
0.46
0.47
0.48
0
5
10
15
20
25
30
Not measured: DTPO, SAPO
Off-scale (y-range set by bold runs): DAPO-SB 0.341 @ step 32 (12 pts); G1 R-read 0.489 @ step 32 (1 pts); GSPO 0.382 @ step 3 (3 pts)
bold = direct comparisons
DAPO
GRPO
GRPO+A
max2
strat
A+B
A+B+C
DAPO-SB
G1 R-read
GSPO
GiGPO
Htau
RLOO
RLOO-decision