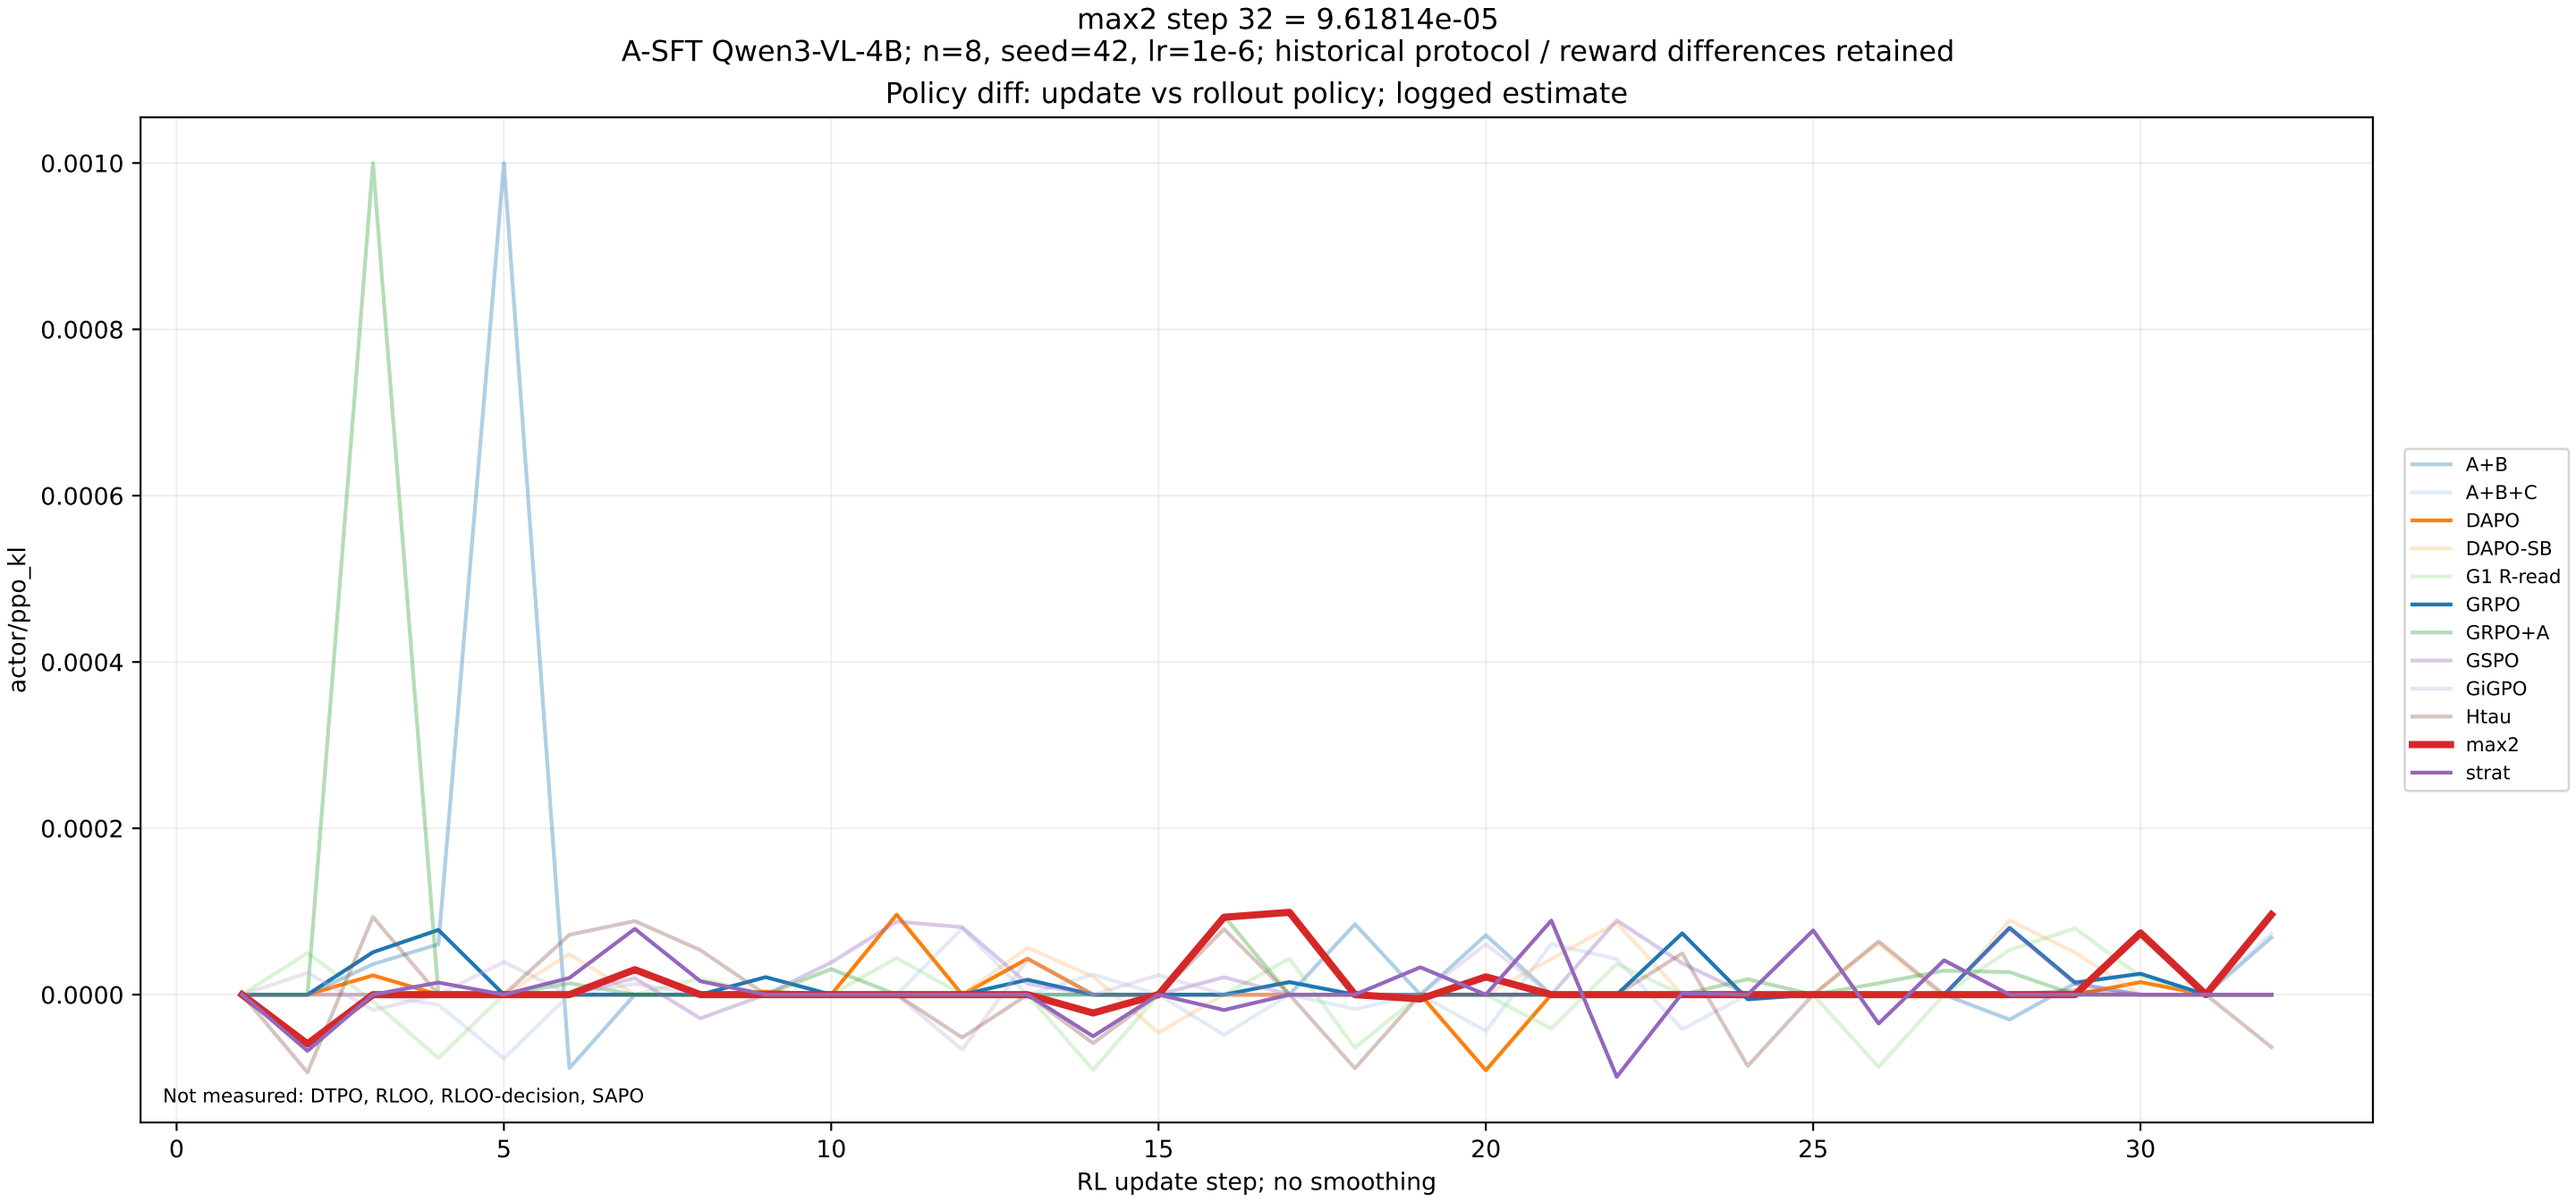
max2 step 32 = 9.61814e-05
A-SFT Qwen3-VL-4B; n=8, seed=42, lr=1e-6; historical protocol / reward differences retained
Policy diff: update vs rollout policy; logged estimate
actor/ppo_kl
0.0010
0.0008
0.0006
0.0004
0.0002
0.0000
0
5
10
15
20
25
30
RL update step; no smoothing
Not measured: DTPO, RLOO, RLOO-decision, SAPO
A+B
A+B+C
DAPO
DAPO-SB
G1 R-read
GRPO
GRPO+A
GSPO
GiGPO
Htau
max2
strat