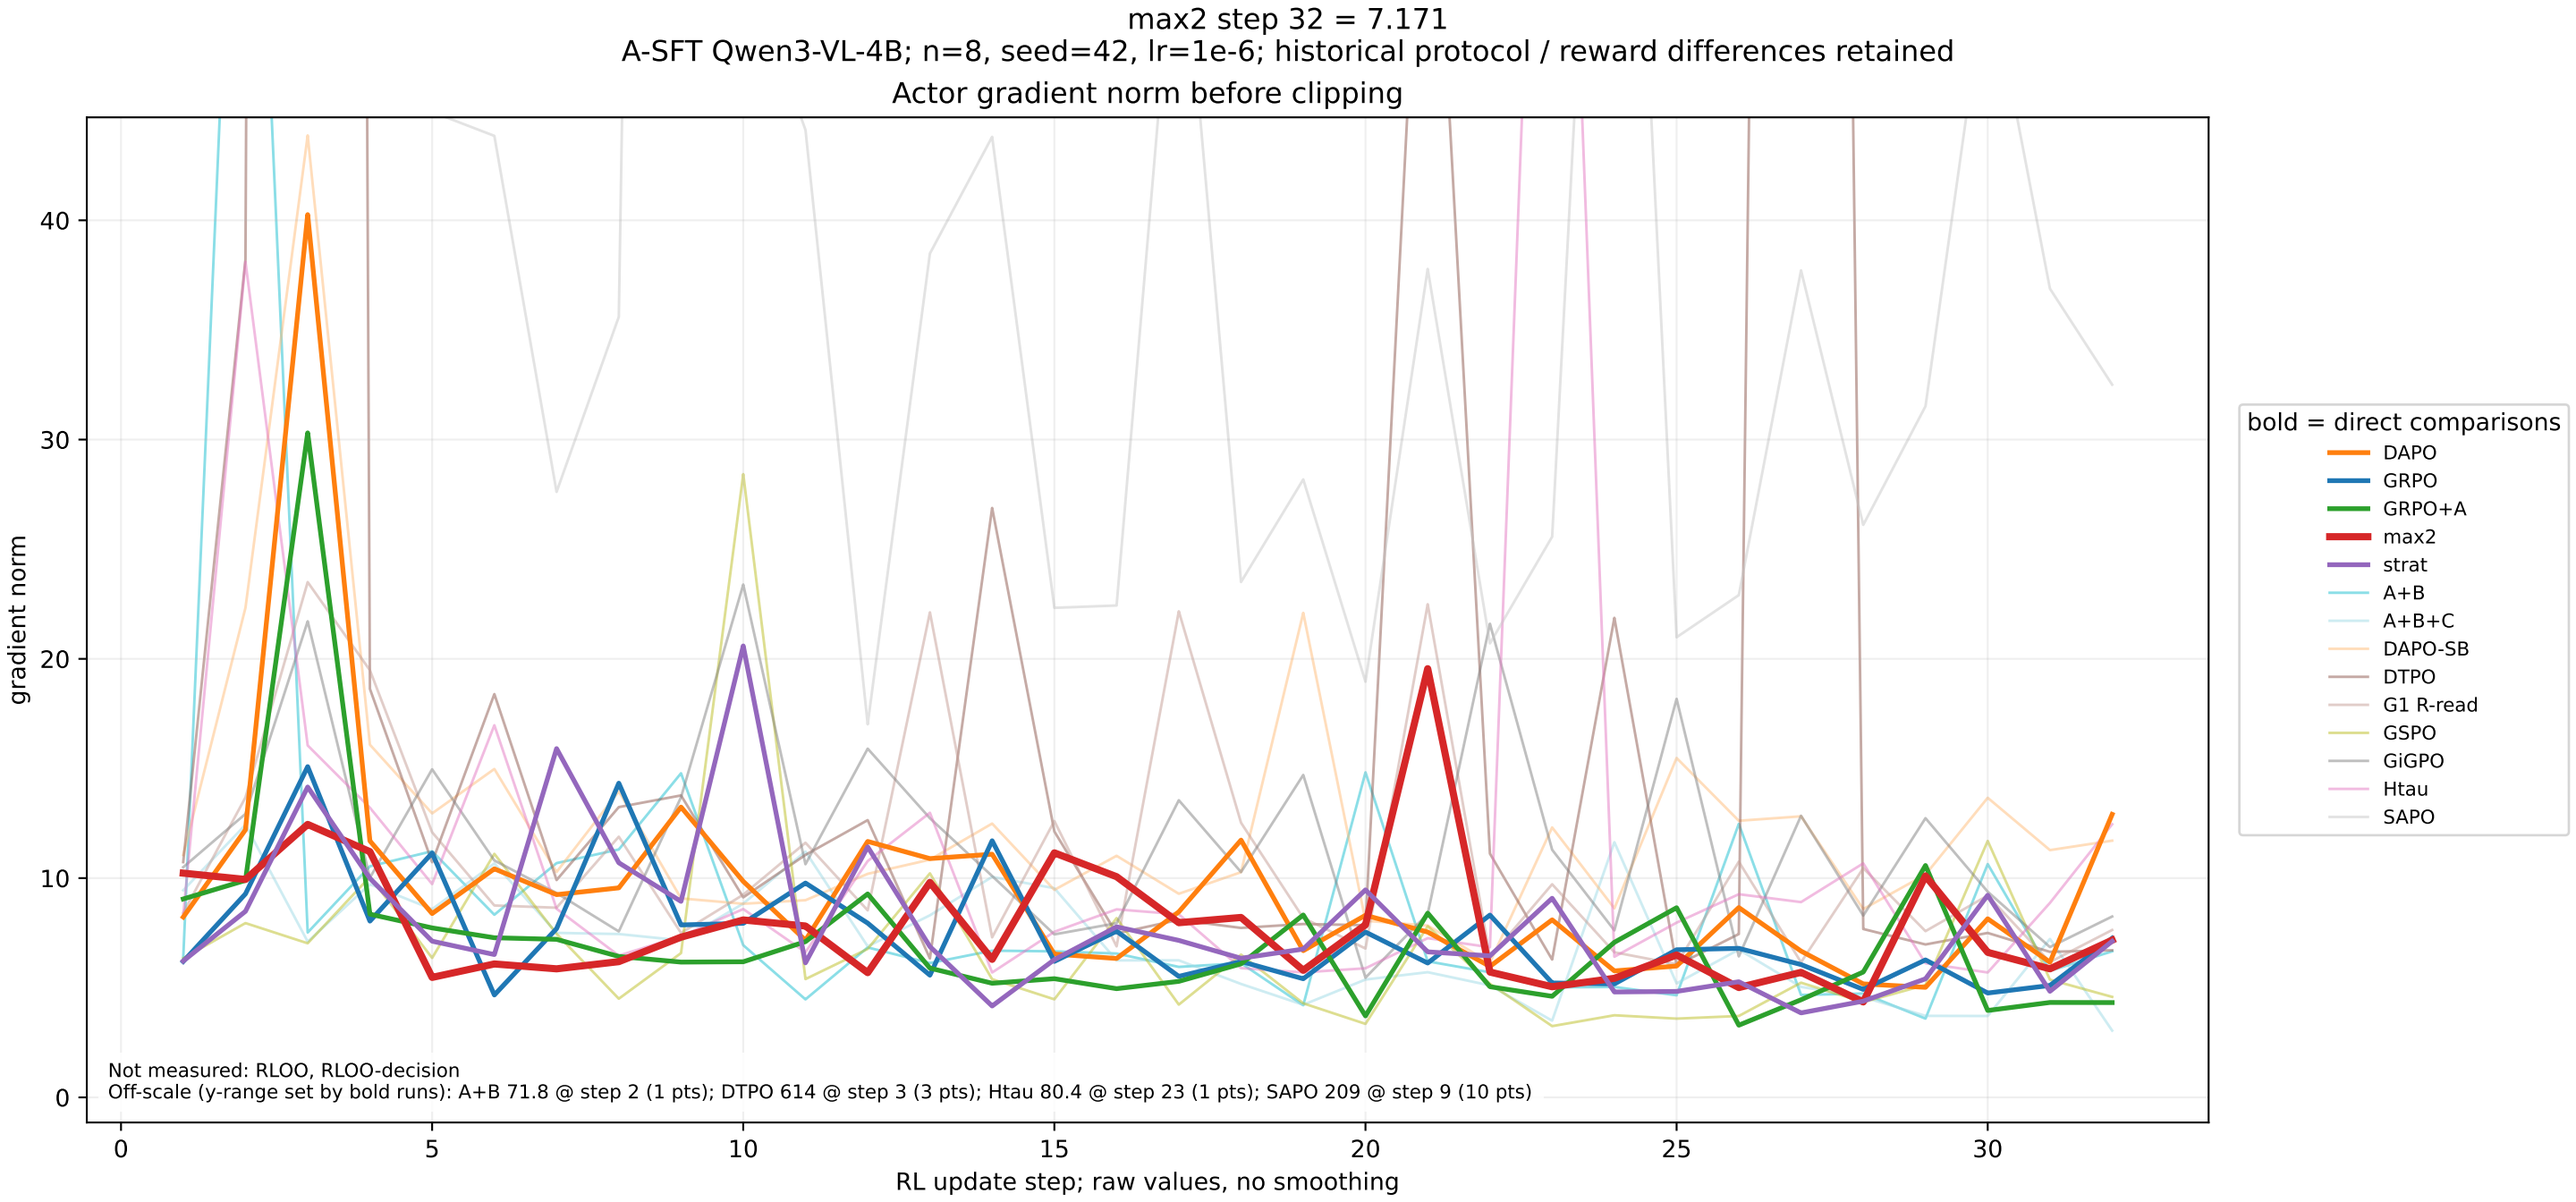
max2 step 32 = 7.171
A-SFT Qwen3-VL-4B; n=8, seed=42, lr=1e-6; historical protocol / reward differences retained
Actor gradient norm before clipping
gradient norm
RL update step; raw values, no smoothing
0
10
20
30
40
0
5
10
15
20
25
30
bold = direct comparisons
DAPO
GRPO
GRPO+A
max2
strat
A+B
A+B+C
DAPO-SB
DTPO
G1 R-read
GSPO
GiGPO
Htau
SAPO
Not measured: RLOO, RLOO-decision
Off-scale (y-range set by bold runs): A+B 71.8 @ step 2 (1 pts); DTPO 614 @ step 3 (3 pts); Htau 80.4 @ step 23 (1 pts); SAPO 209 @ step 9 (10 pts)