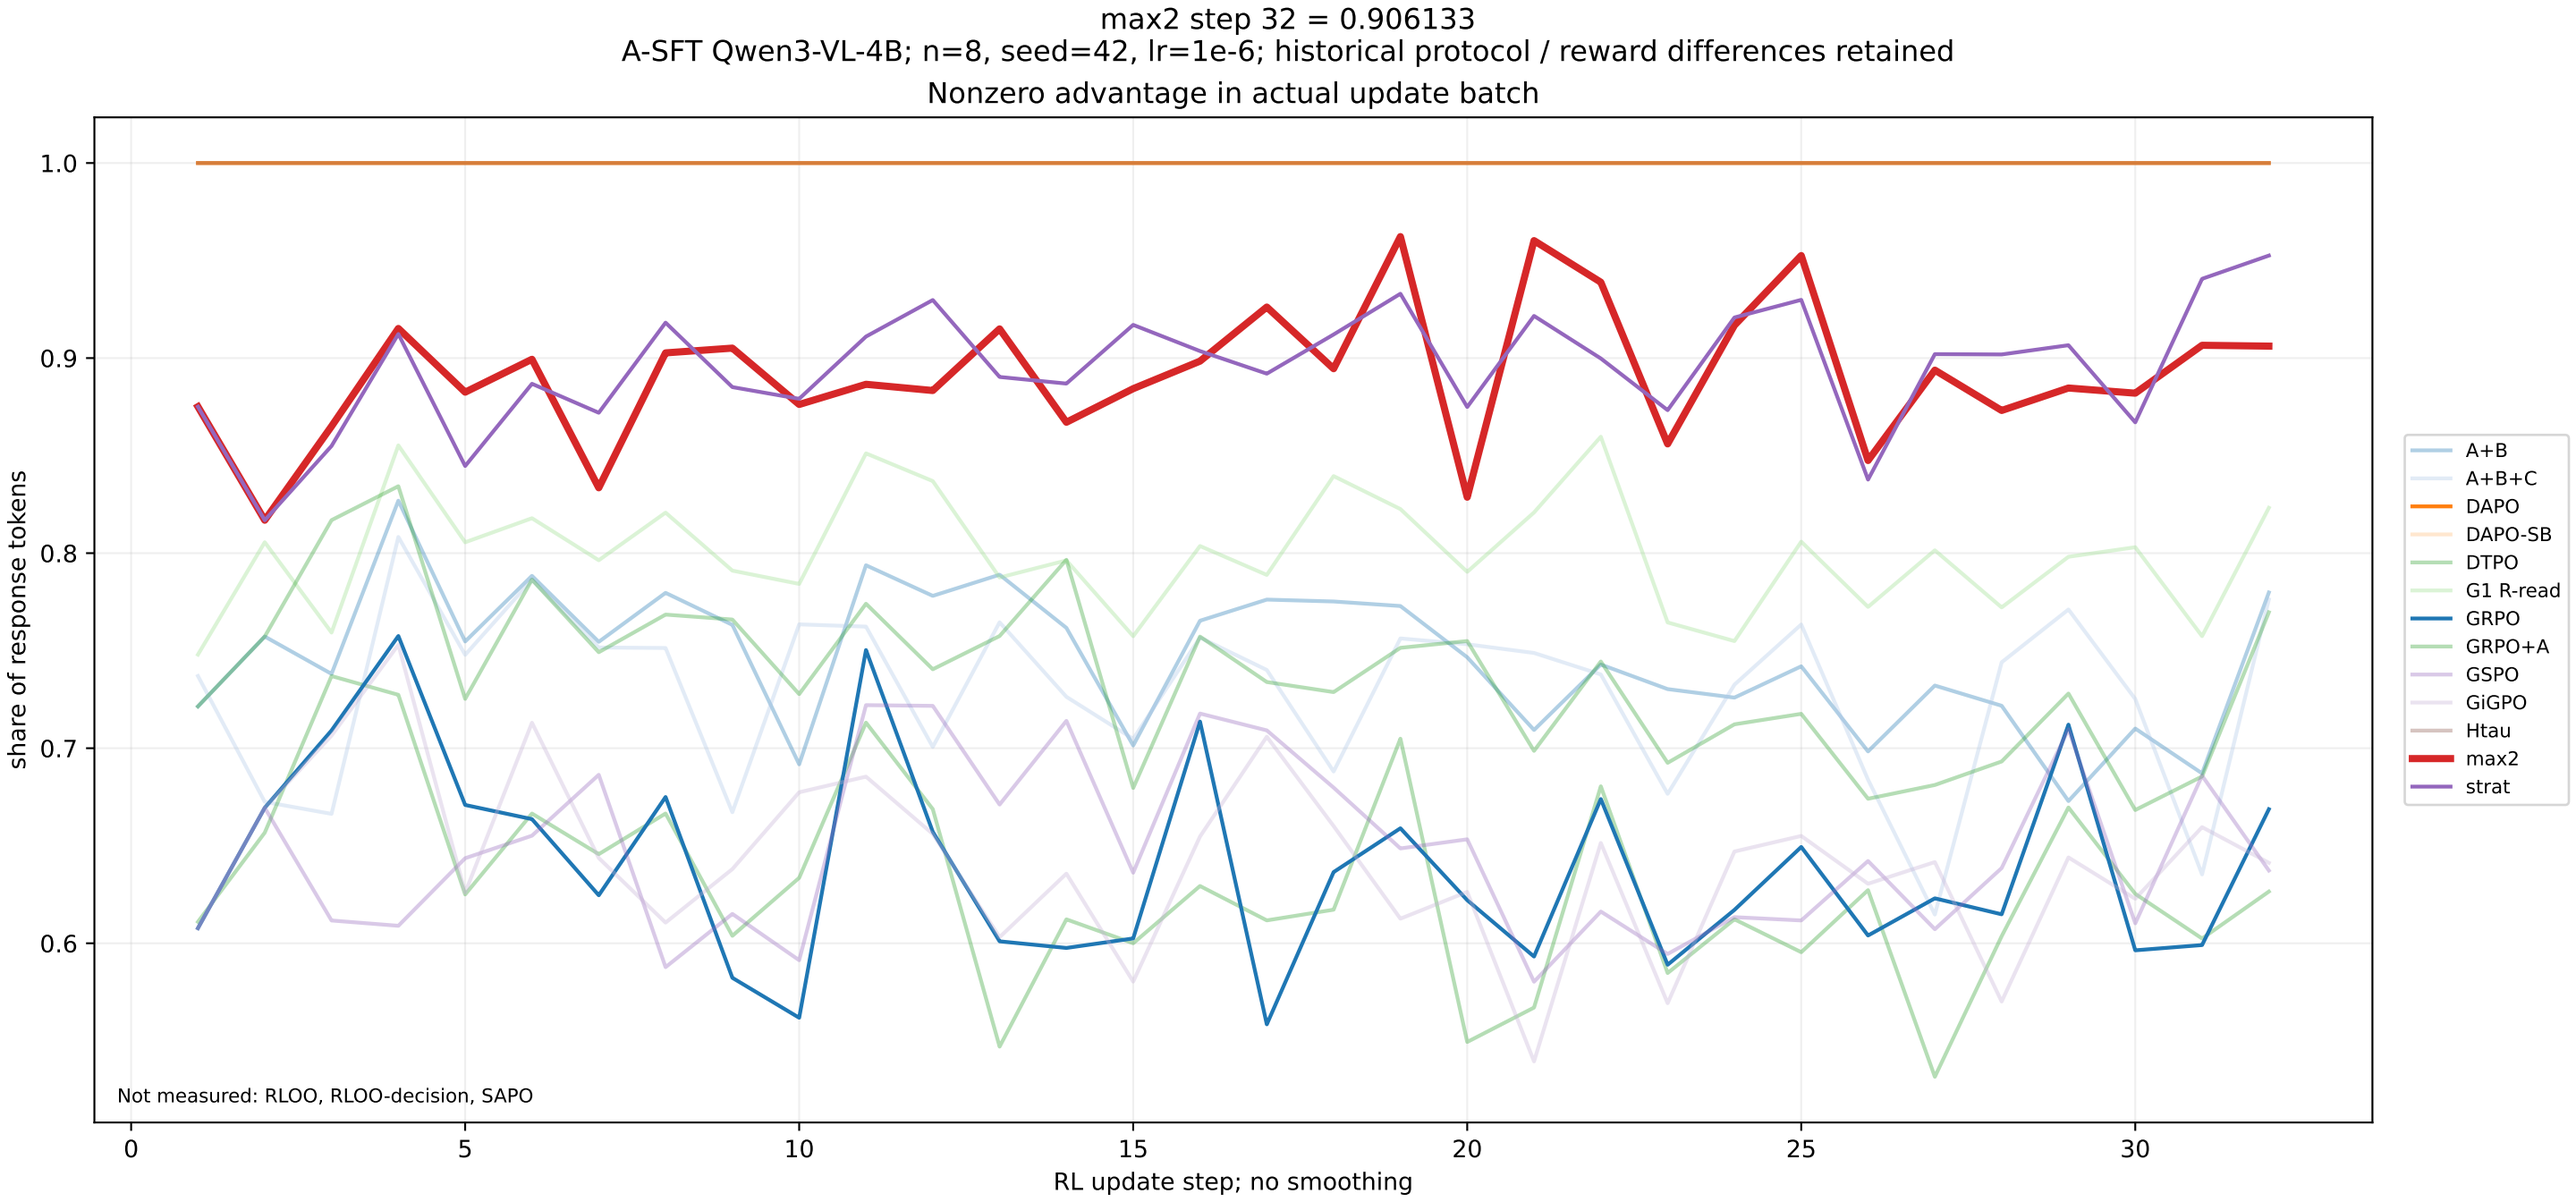
max2 step 32 = 0.906133
A-SFT Qwen3-VL-4B; n=8, seed=42, lr=1e-6; historical protocol / reward differences retained
Nonzero advantage in actual update batch
share of response tokens
RL update step; no smoothing
Not measured: RLOO, RLOO-decision, SAPO
A+B
A+B+C
DAPO
DAPO-SB
DTPO
G1 R-read
GRPO
GRPO+A
GSPO
GiGPO
Htau
max2
strat