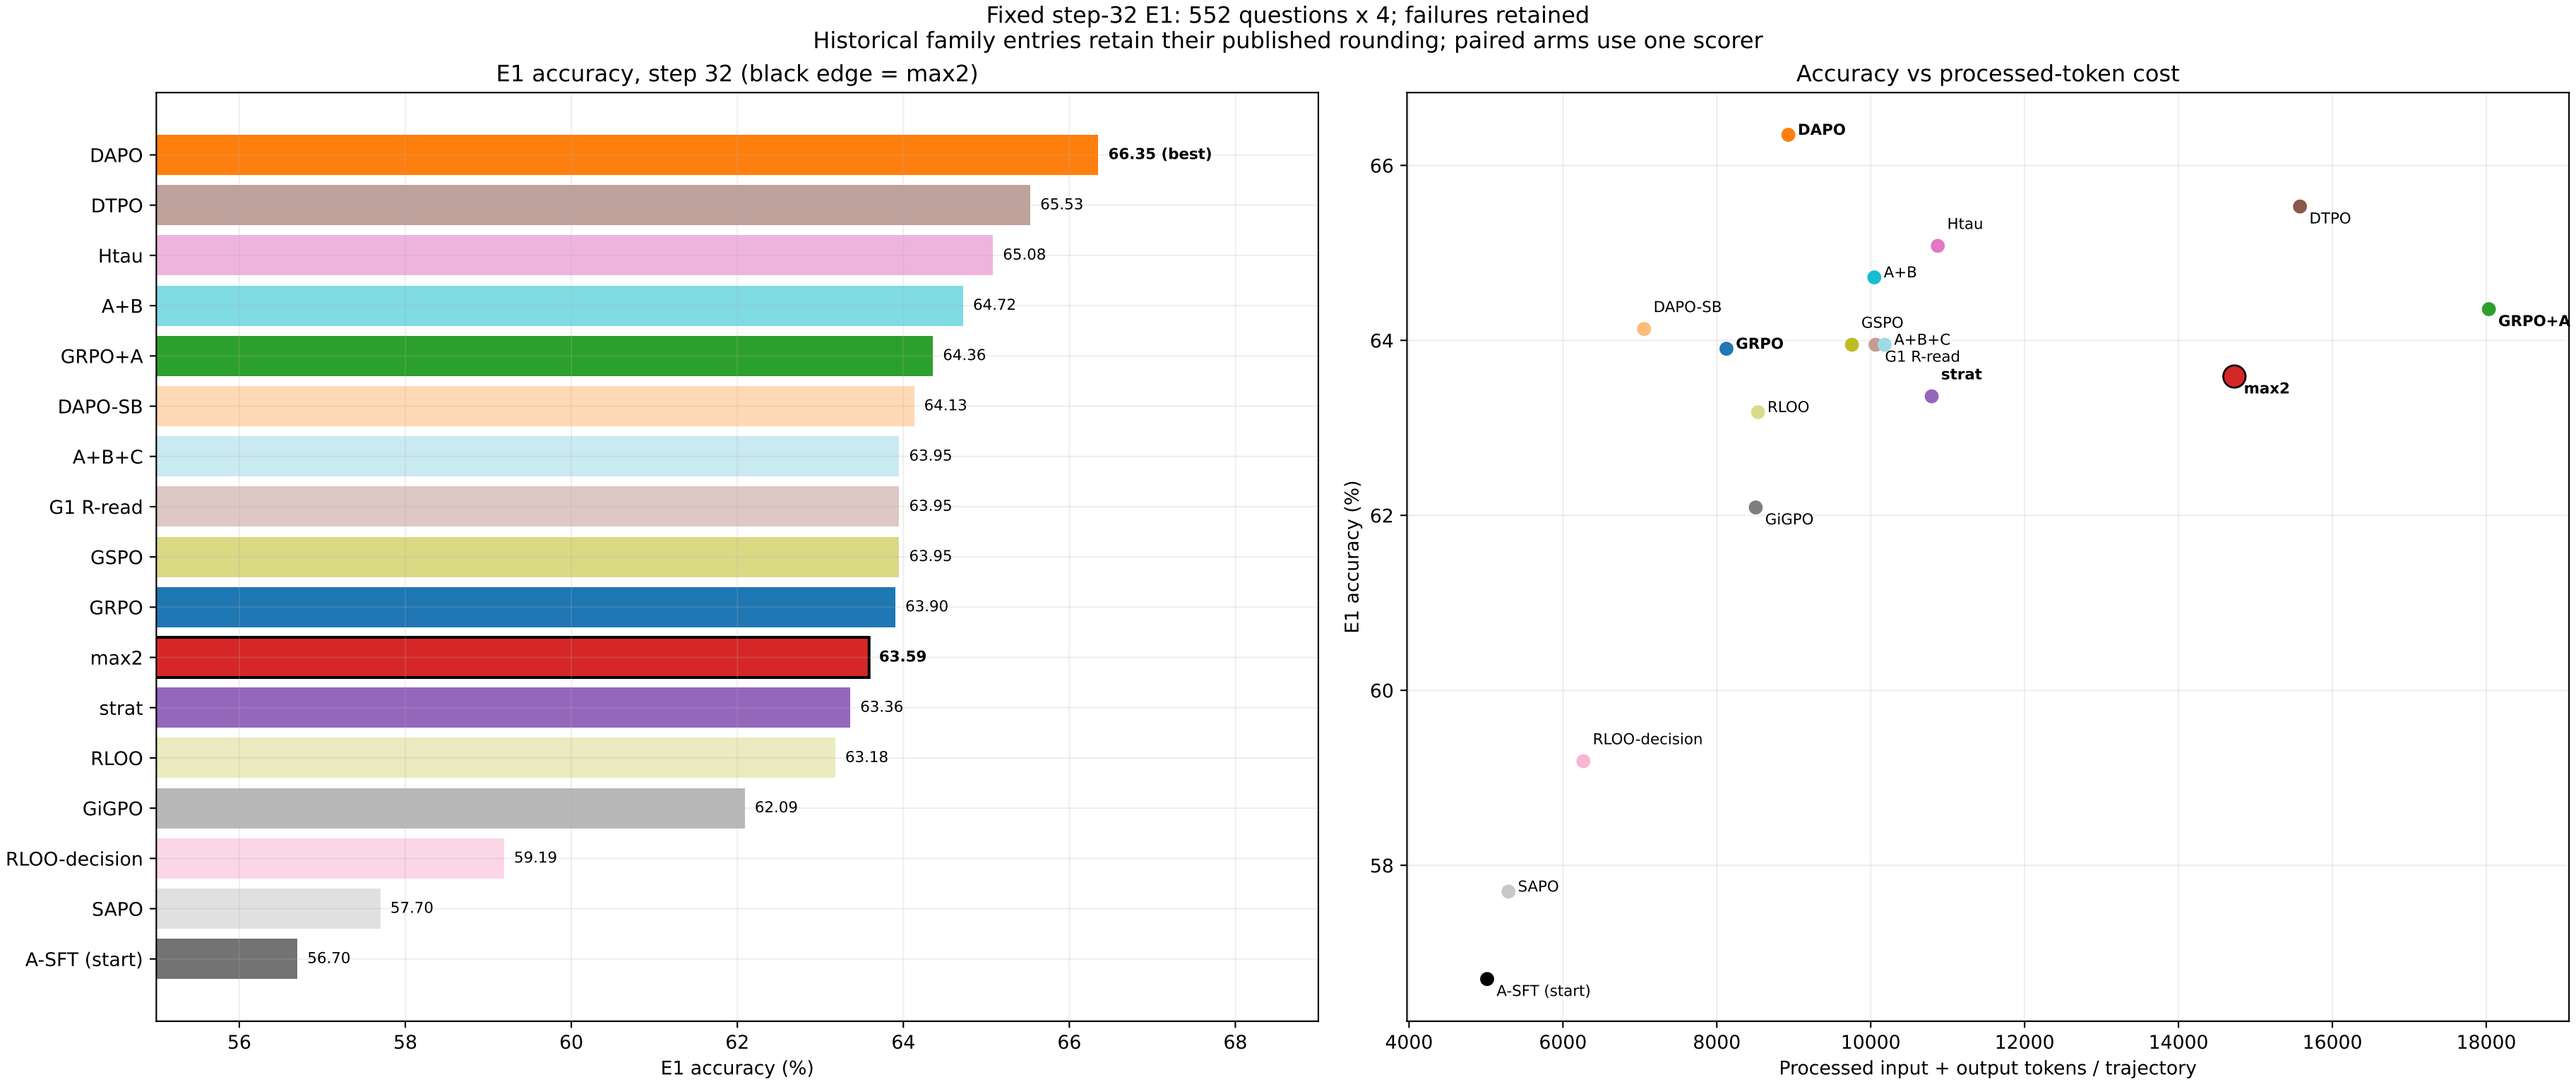
Fixed step-32 E1: 552 questions x 4; failures retained
Historical family entries retain their published rounding; paired arms use one scorer
E1 accuracy, step 32 (black edge = max2)
DAPO
DTPO
Htau
A+B
GRPO+A
DAPO-SB
A+B+C
G1 R-read
GSPO
GRPO
max2
strat
RLOO
GiGPO
RLOO-decision
SAPO
A-SFT (start)
66.35 (best)
65.53
65.08
64.72
64.36
64.13
63.95
63.95
63.95
63.90
63.59
63.36
63.18
62.09
59.19
57.70
56.70
56
58
60
62
64
66
68
E1 accuracy (%)
Accuracy vs processed-token cost
E1 accuracy (%)
Processed input + output tokens / trajectory
4000
6000
8000
10000
12000
14000
16000
18000
58
60
62
64
66