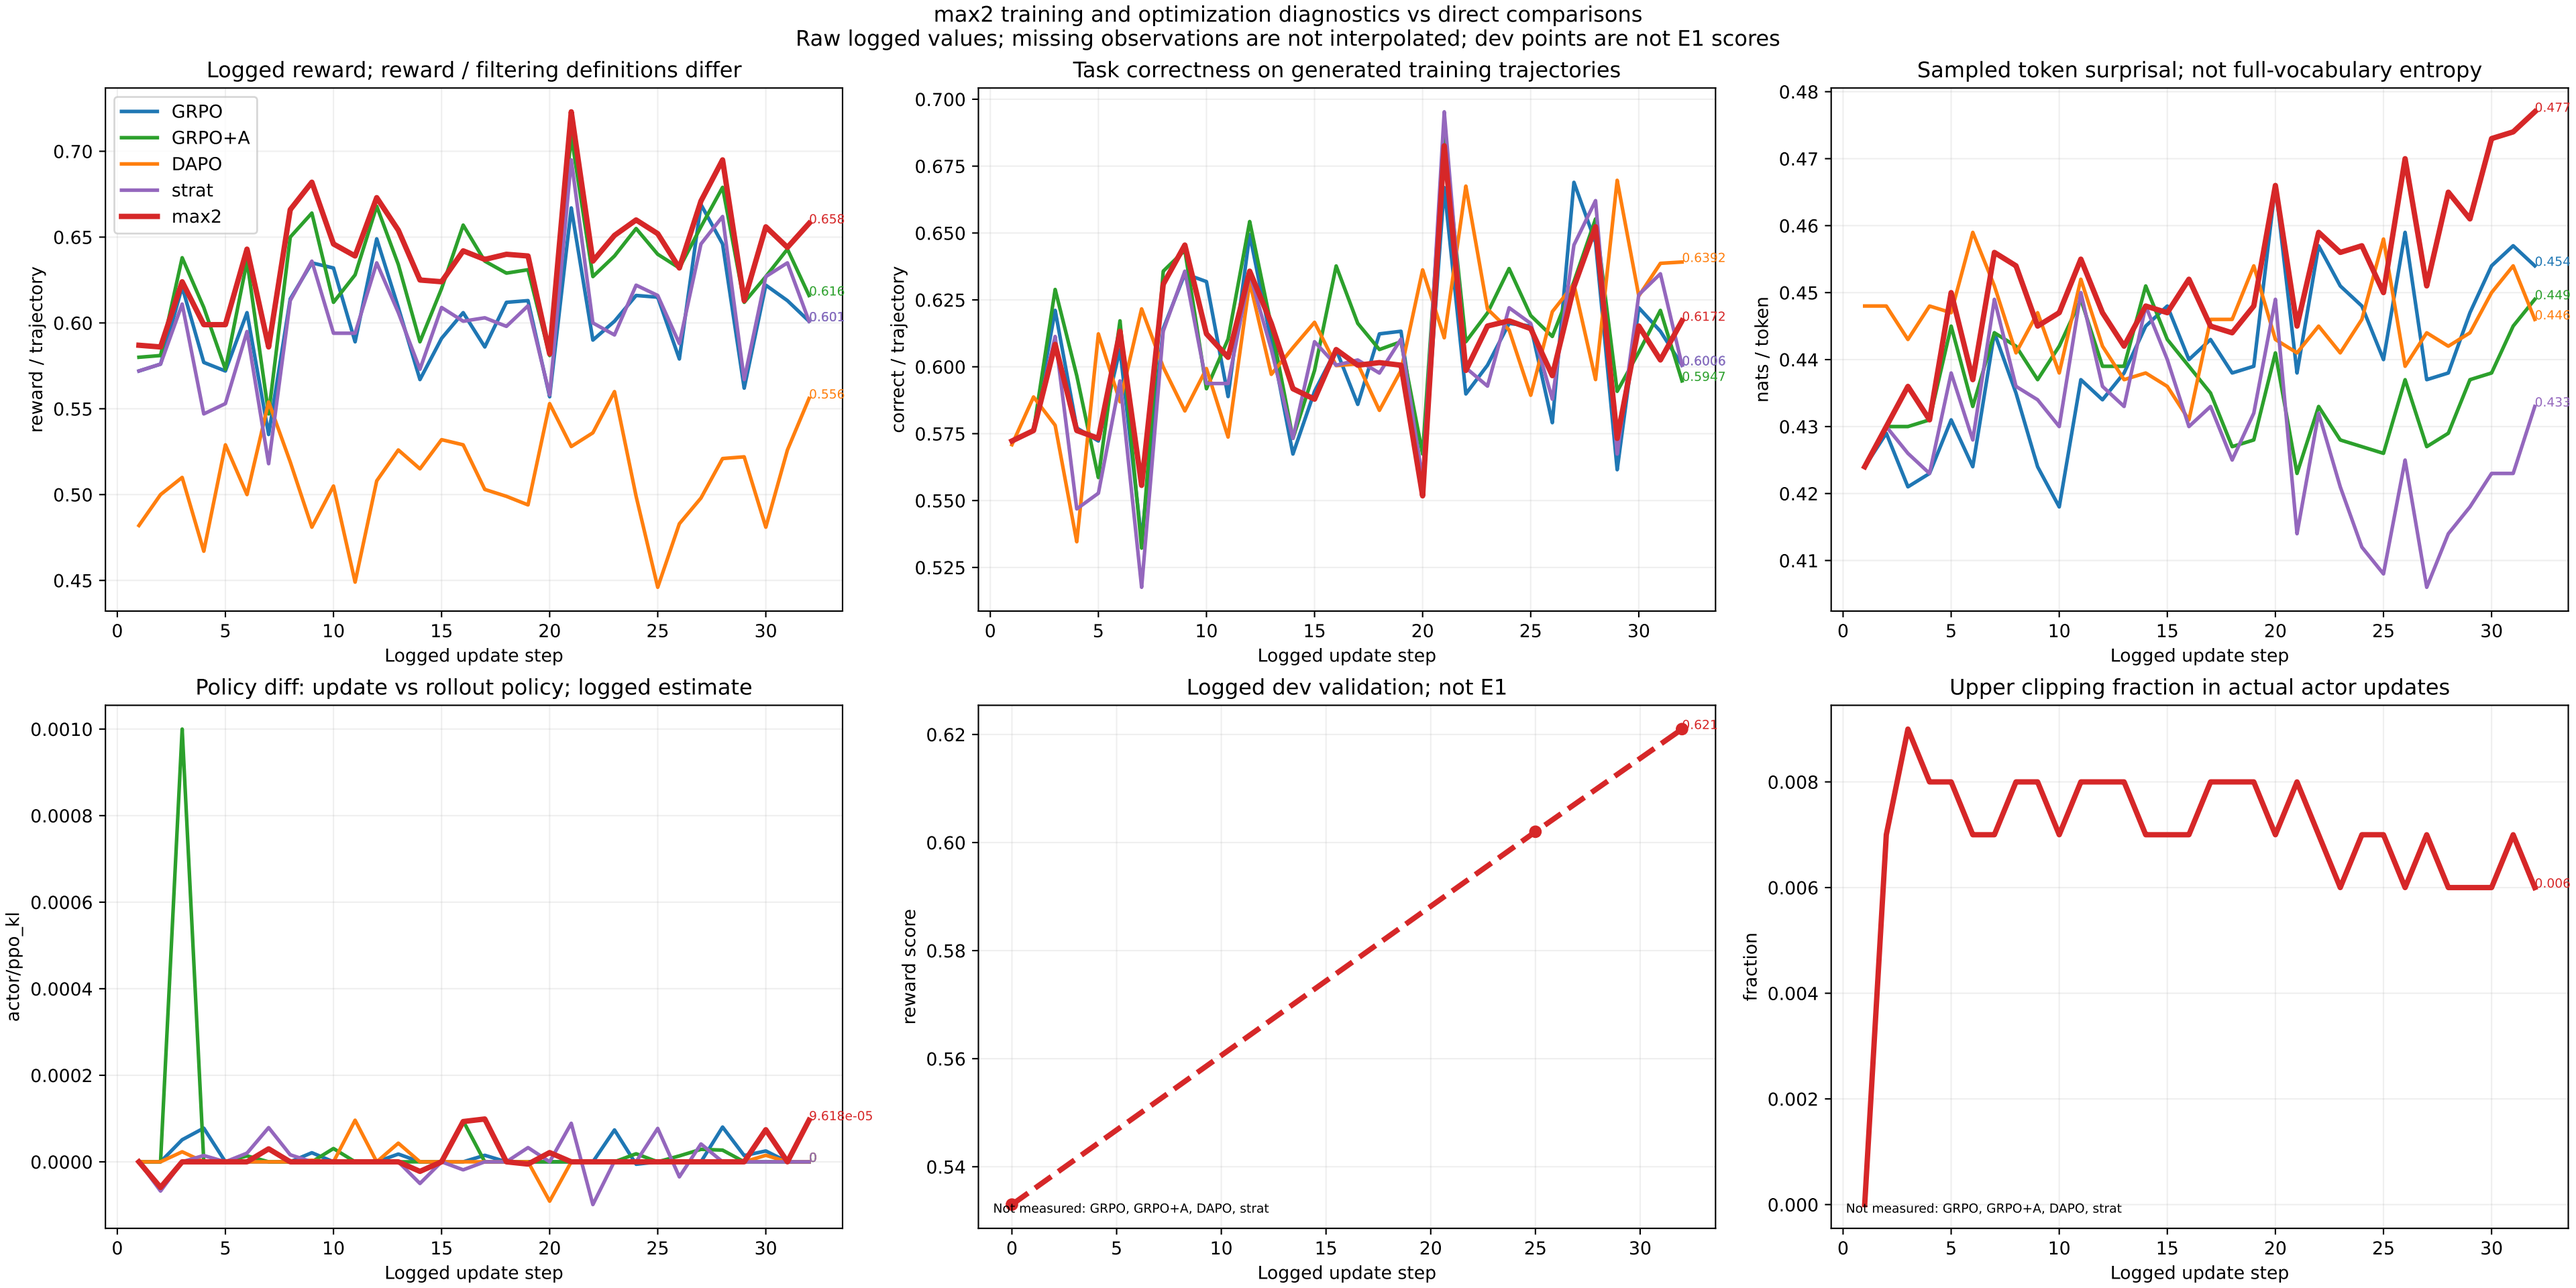
max2 training and optimization diagnostics vs direct comparisons
Raw logged values; missing observations are not interpolated; dev points are not E1 scores
Logged reward; reward / filtering definitions differ
reward / trajectory
Logged update step
GRPO
GRPO+A
DAPO
strat
max2
0.658
0.616
0.601
0.556
Task correctness on generated training trajectories
correct / trajectory
Logged update step
0.6392
0.6172
0.6006
0.5947
Sampled token surprisal; not full-vocabulary entropy
nats / token
Logged update step
0.477
0.454
0.449
0.446
0.433
Policy diff: update vs rollout policy; logged estimate
actor/ppo_kl
Logged update step
9.618e-05
0
Logged dev validation; not E1
reward score
Logged update step
0.621
Not measured: GRPO, GRPO+A, DAPO, strat
Upper clipping fraction in actual actor updates
fraction
Logged update step
0.006
Not measured: GRPO, GRPO+A, DAPO, strat