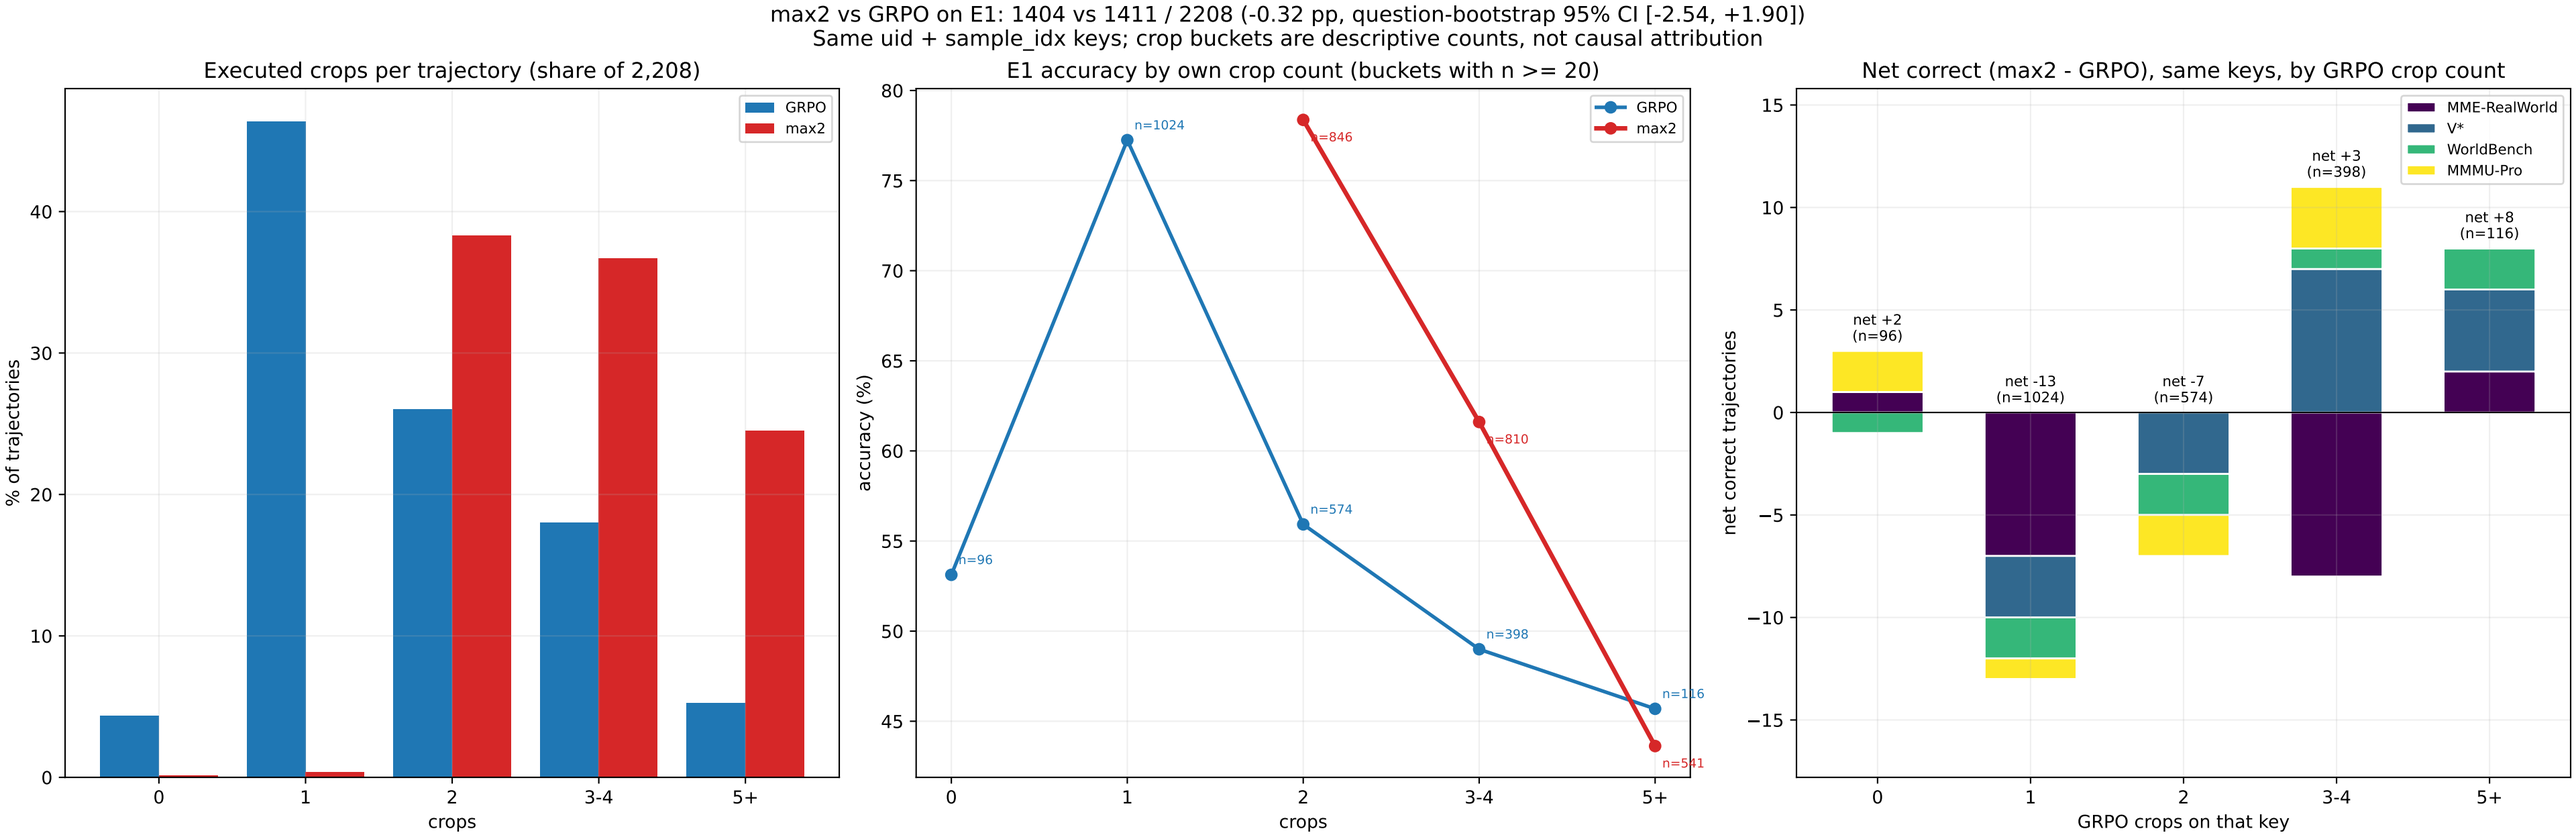
max2 vs GRPO on E1: 1404 vs 1411 / 2208 (-0.32 pp, question-bootstrap 95% CI [-2.54, +1.90])
Same uid + sample_idx keys; crop buckets are descriptive counts, not causal attribution
Executed crops per trajectory (share of 2,208)
GRPO
max2
% of trajectories
crops
0
1
2
3-4
5+
E1 accuracy by own crop count (buckets with n >= 20)
accuracy (%)
n=96
n=1024
n=574
n=398
n=116
n=846
n=810
n=541
Net correct (max2 - GRPO), same keys, by GRPO crop count
MME-RealWorld
V*
WorldBench
MMMU-Pro
net +2
(n=96)
net -13
(n=1024)
net -7
(n=574)
net +3
(n=398)
net +8
(n=116)
net correct trajectories
GRPO crops on that key