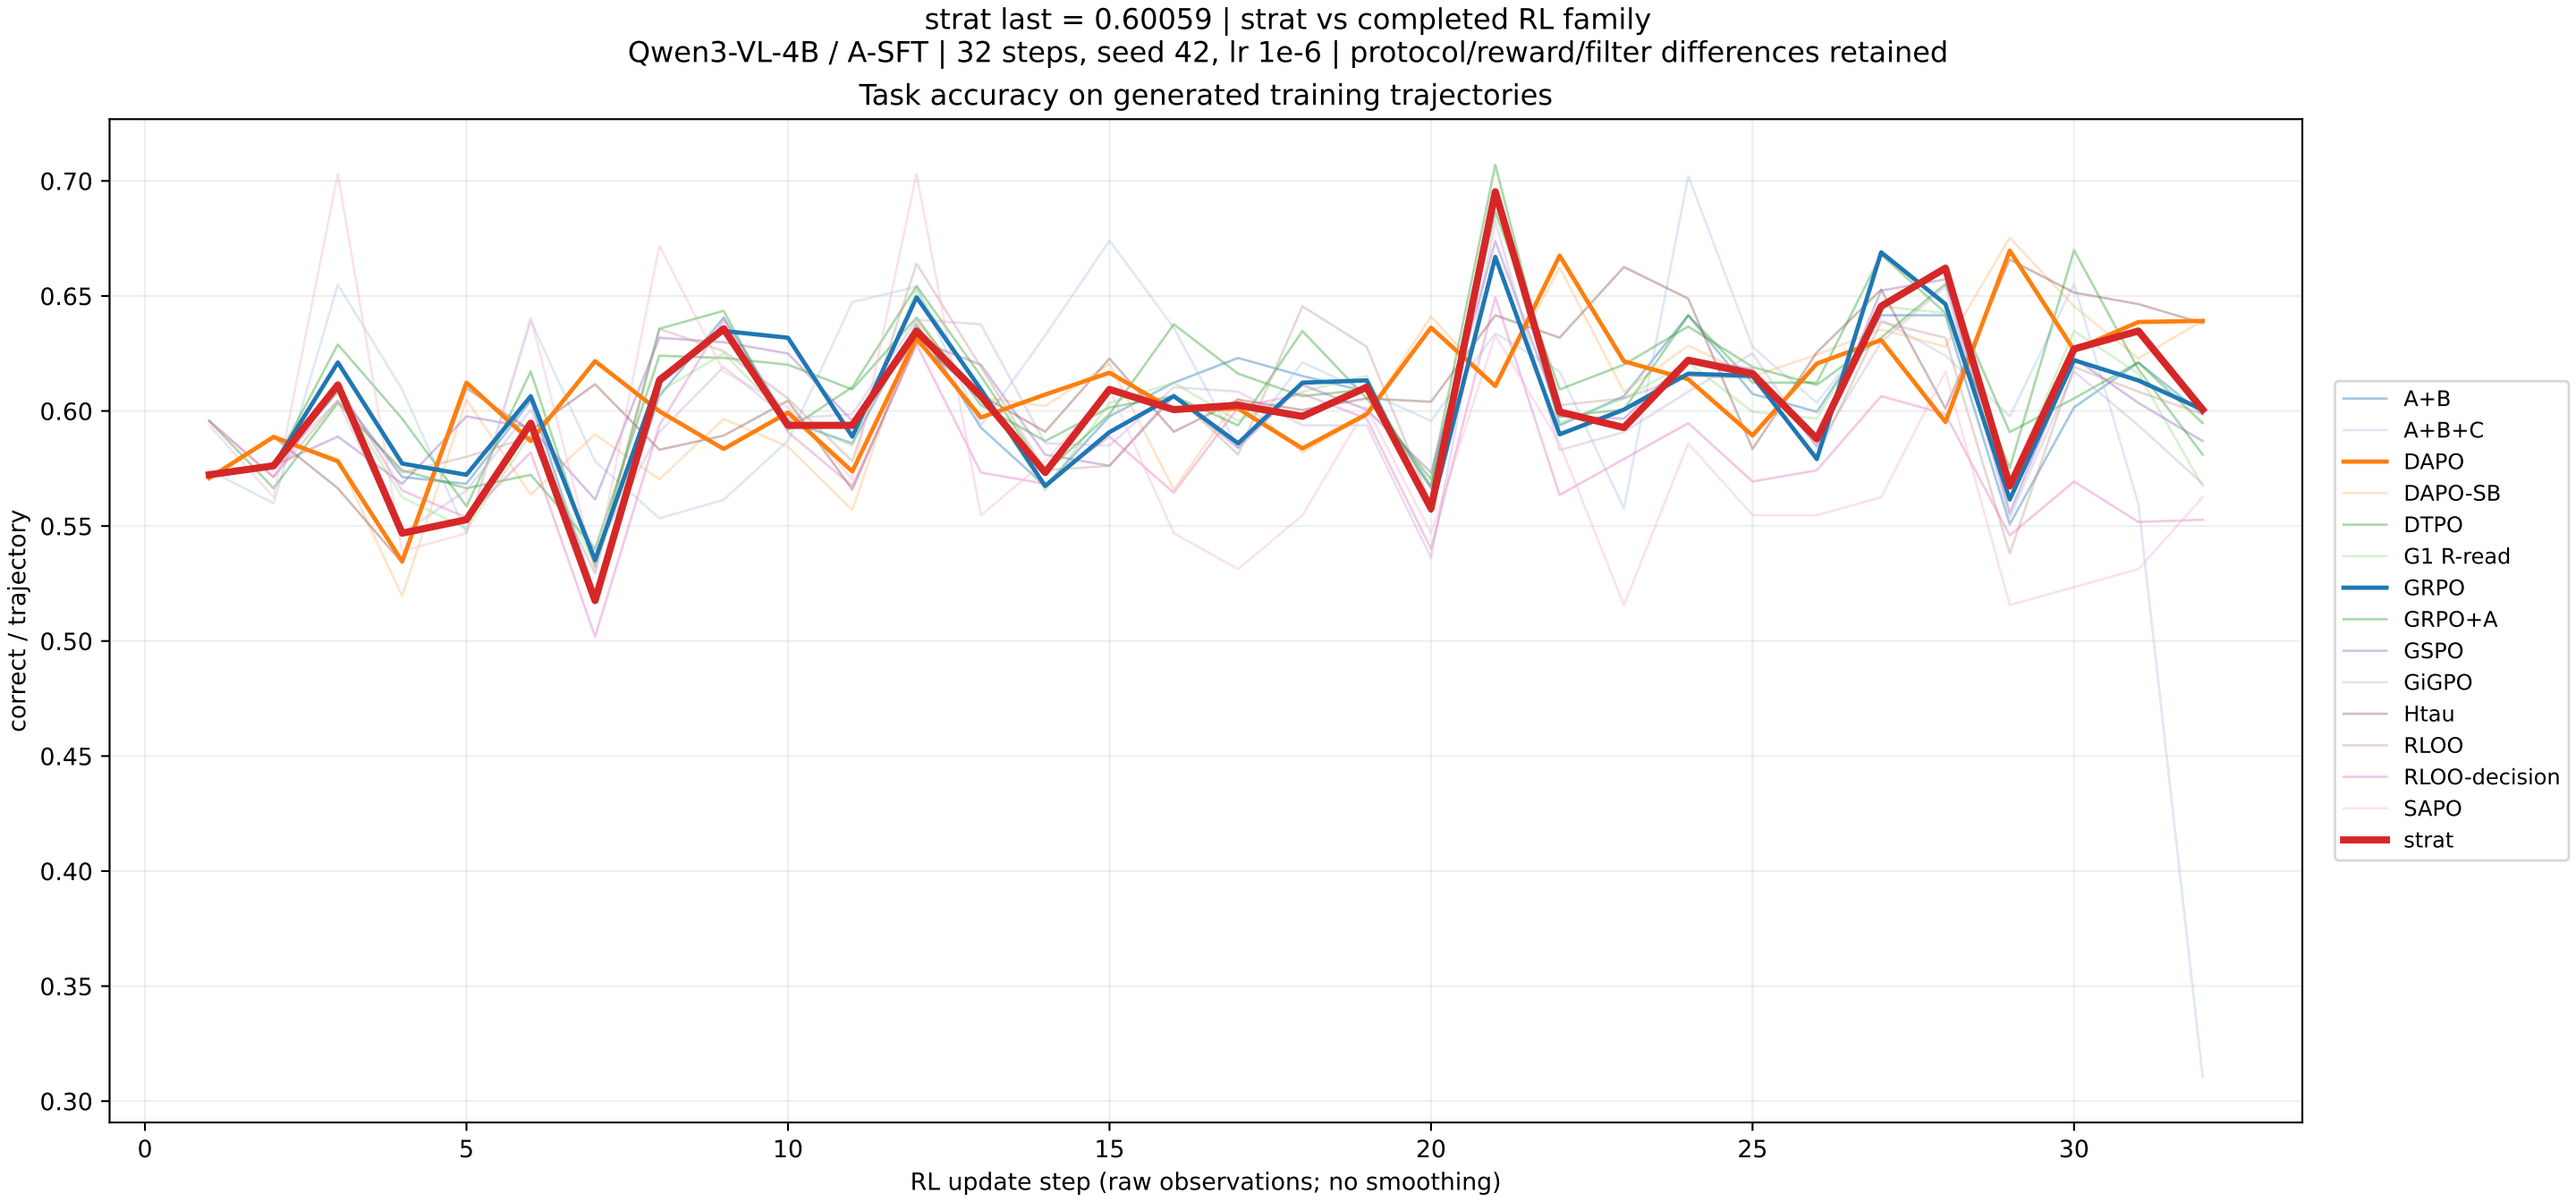
strat last = 0.60059 | strat vs completed RL family
Qwen3-VL-4B / A-SFT | 32 steps, seed 42, lr 1e-6 | protocol/reward/filter differences retained
Task accuracy on generated training trajectories
correct / trajectory
RL update step (raw observations; no smoothing)
0.70
0.65
0.60
0.55
0.50
0.45
0.40
0.35
0.30
0
5
10
15
20
25
30
A+B
A+B+C
DAPO
DAPO-SB
DTPO
G1 R-read
GRPO
GRPO+A
GSPO
GiGPO
Htau
RLOO
RLOO-decision
SAPO
strat