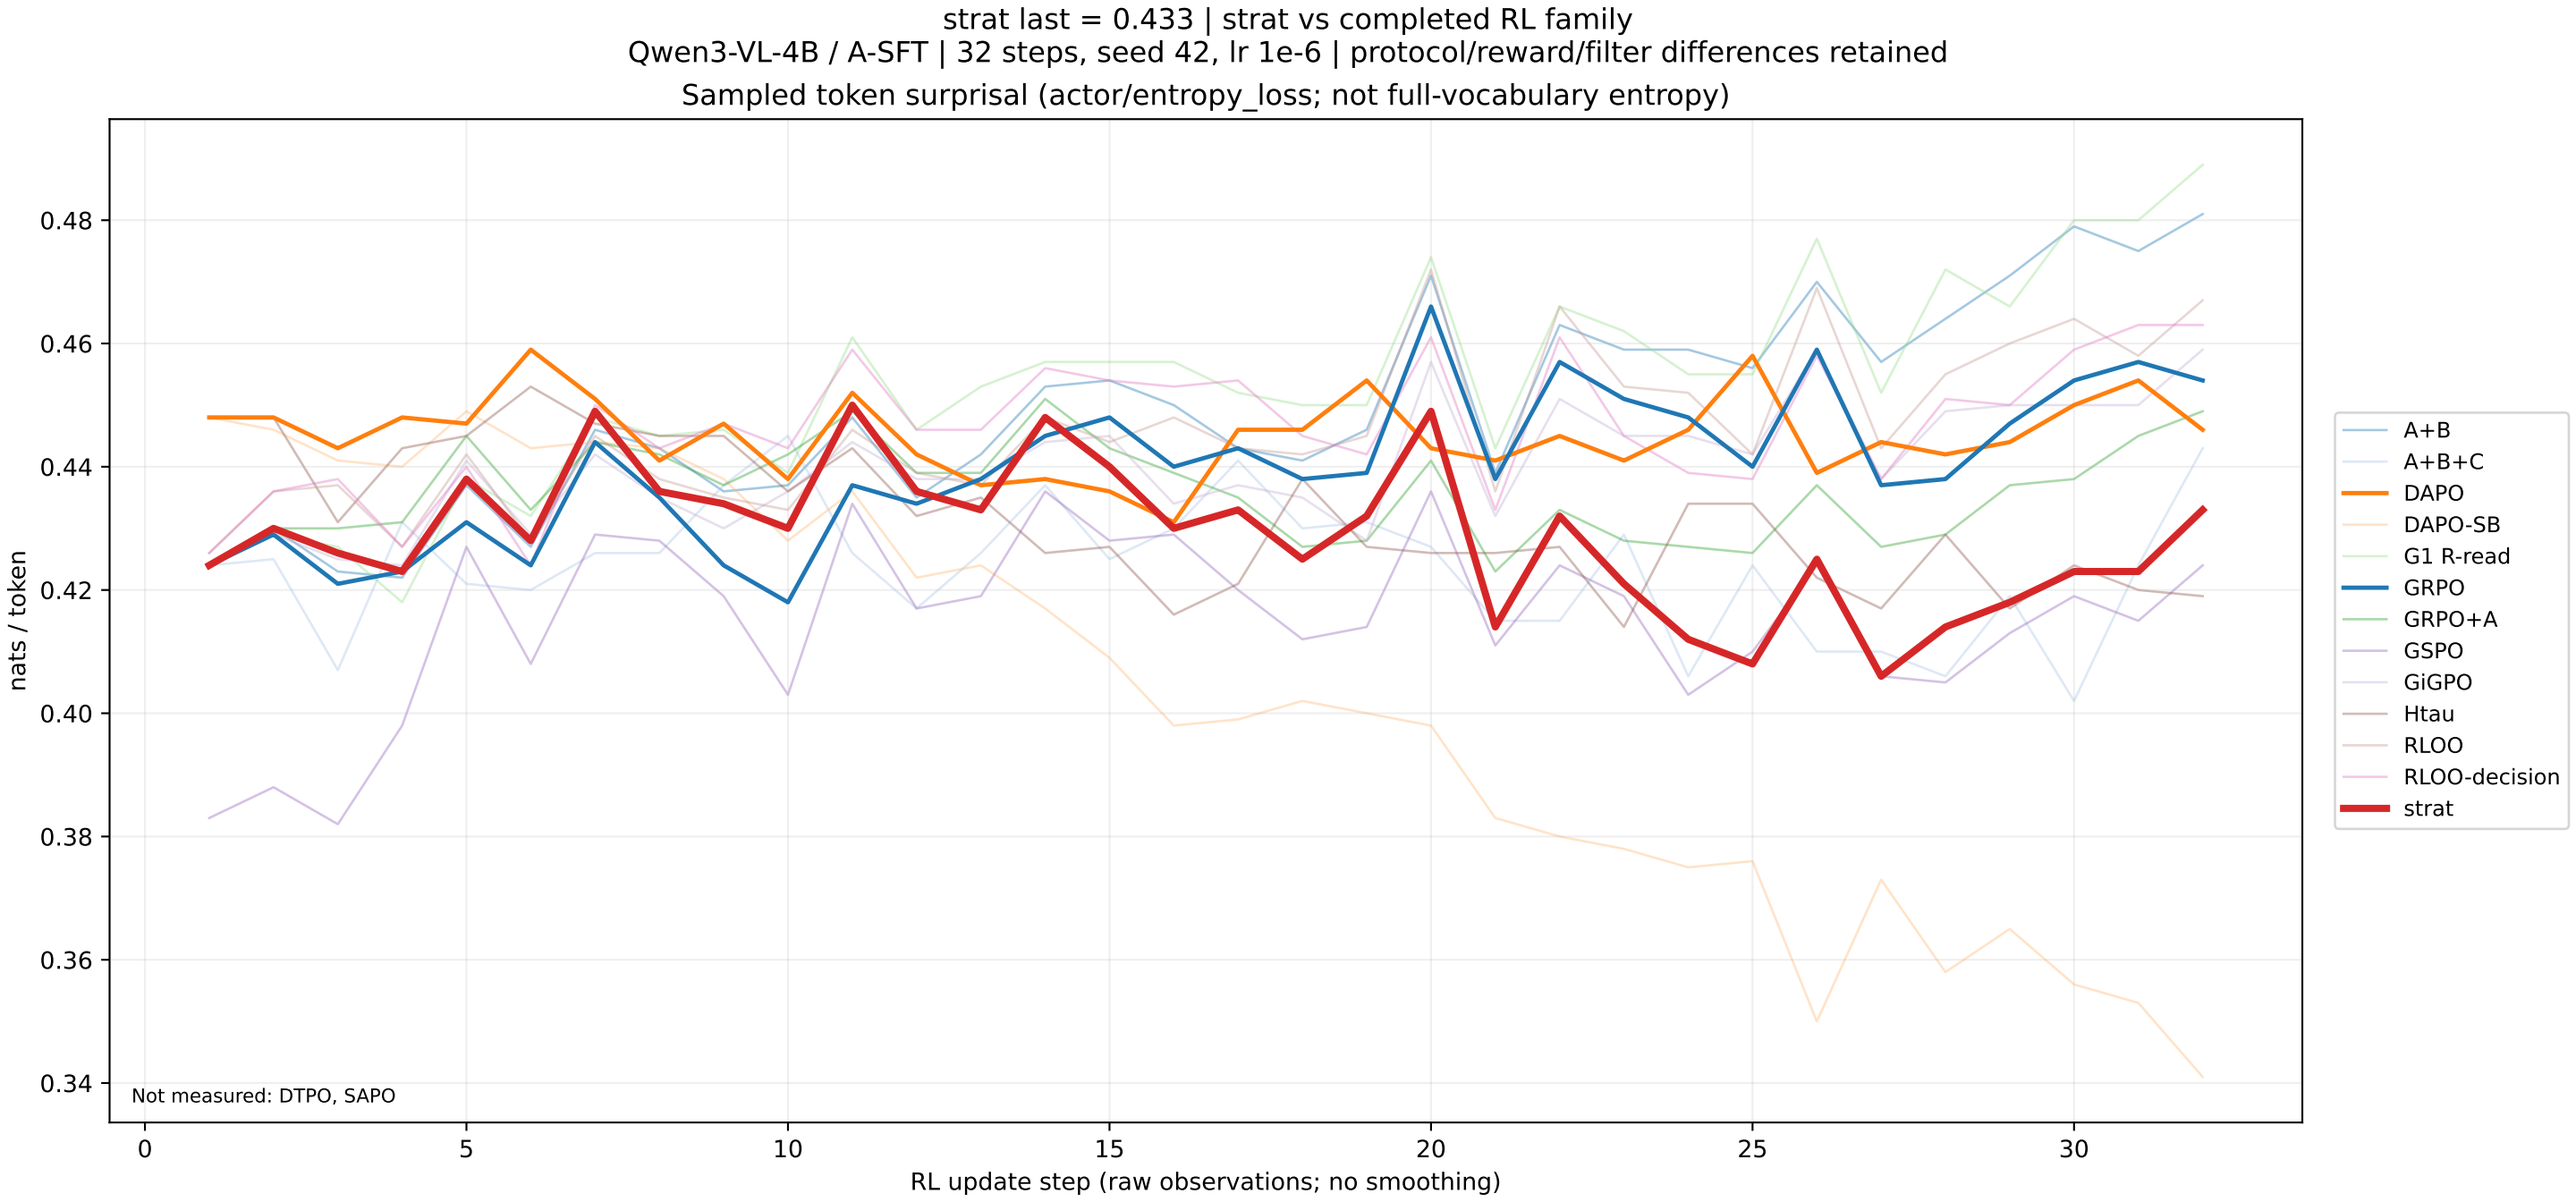
strat last = 0.433 | strat vs completed RL family
Qwen3-VL-4B / A-SFT | 32 steps, seed 42, lr 1e-6 | protocol/reward/filter differences retained
Sampled token surprisal (actor/entropy_loss; not full-vocabulary entropy)
nats / token
0.48
0.46
0.44
0.42
0.40
0.38
0.36
0.34
Not measured: DTPO, SAPO
0
5
10
15
20
25
30
RL update step (raw observations; no smoothing)
A+B
A+B+C
DAPO
DAPO-SB
G1 R-read
GRPO
GRPO+A
GSPO
GiGPO
Htau
RLOO
RLOO-decision
strat