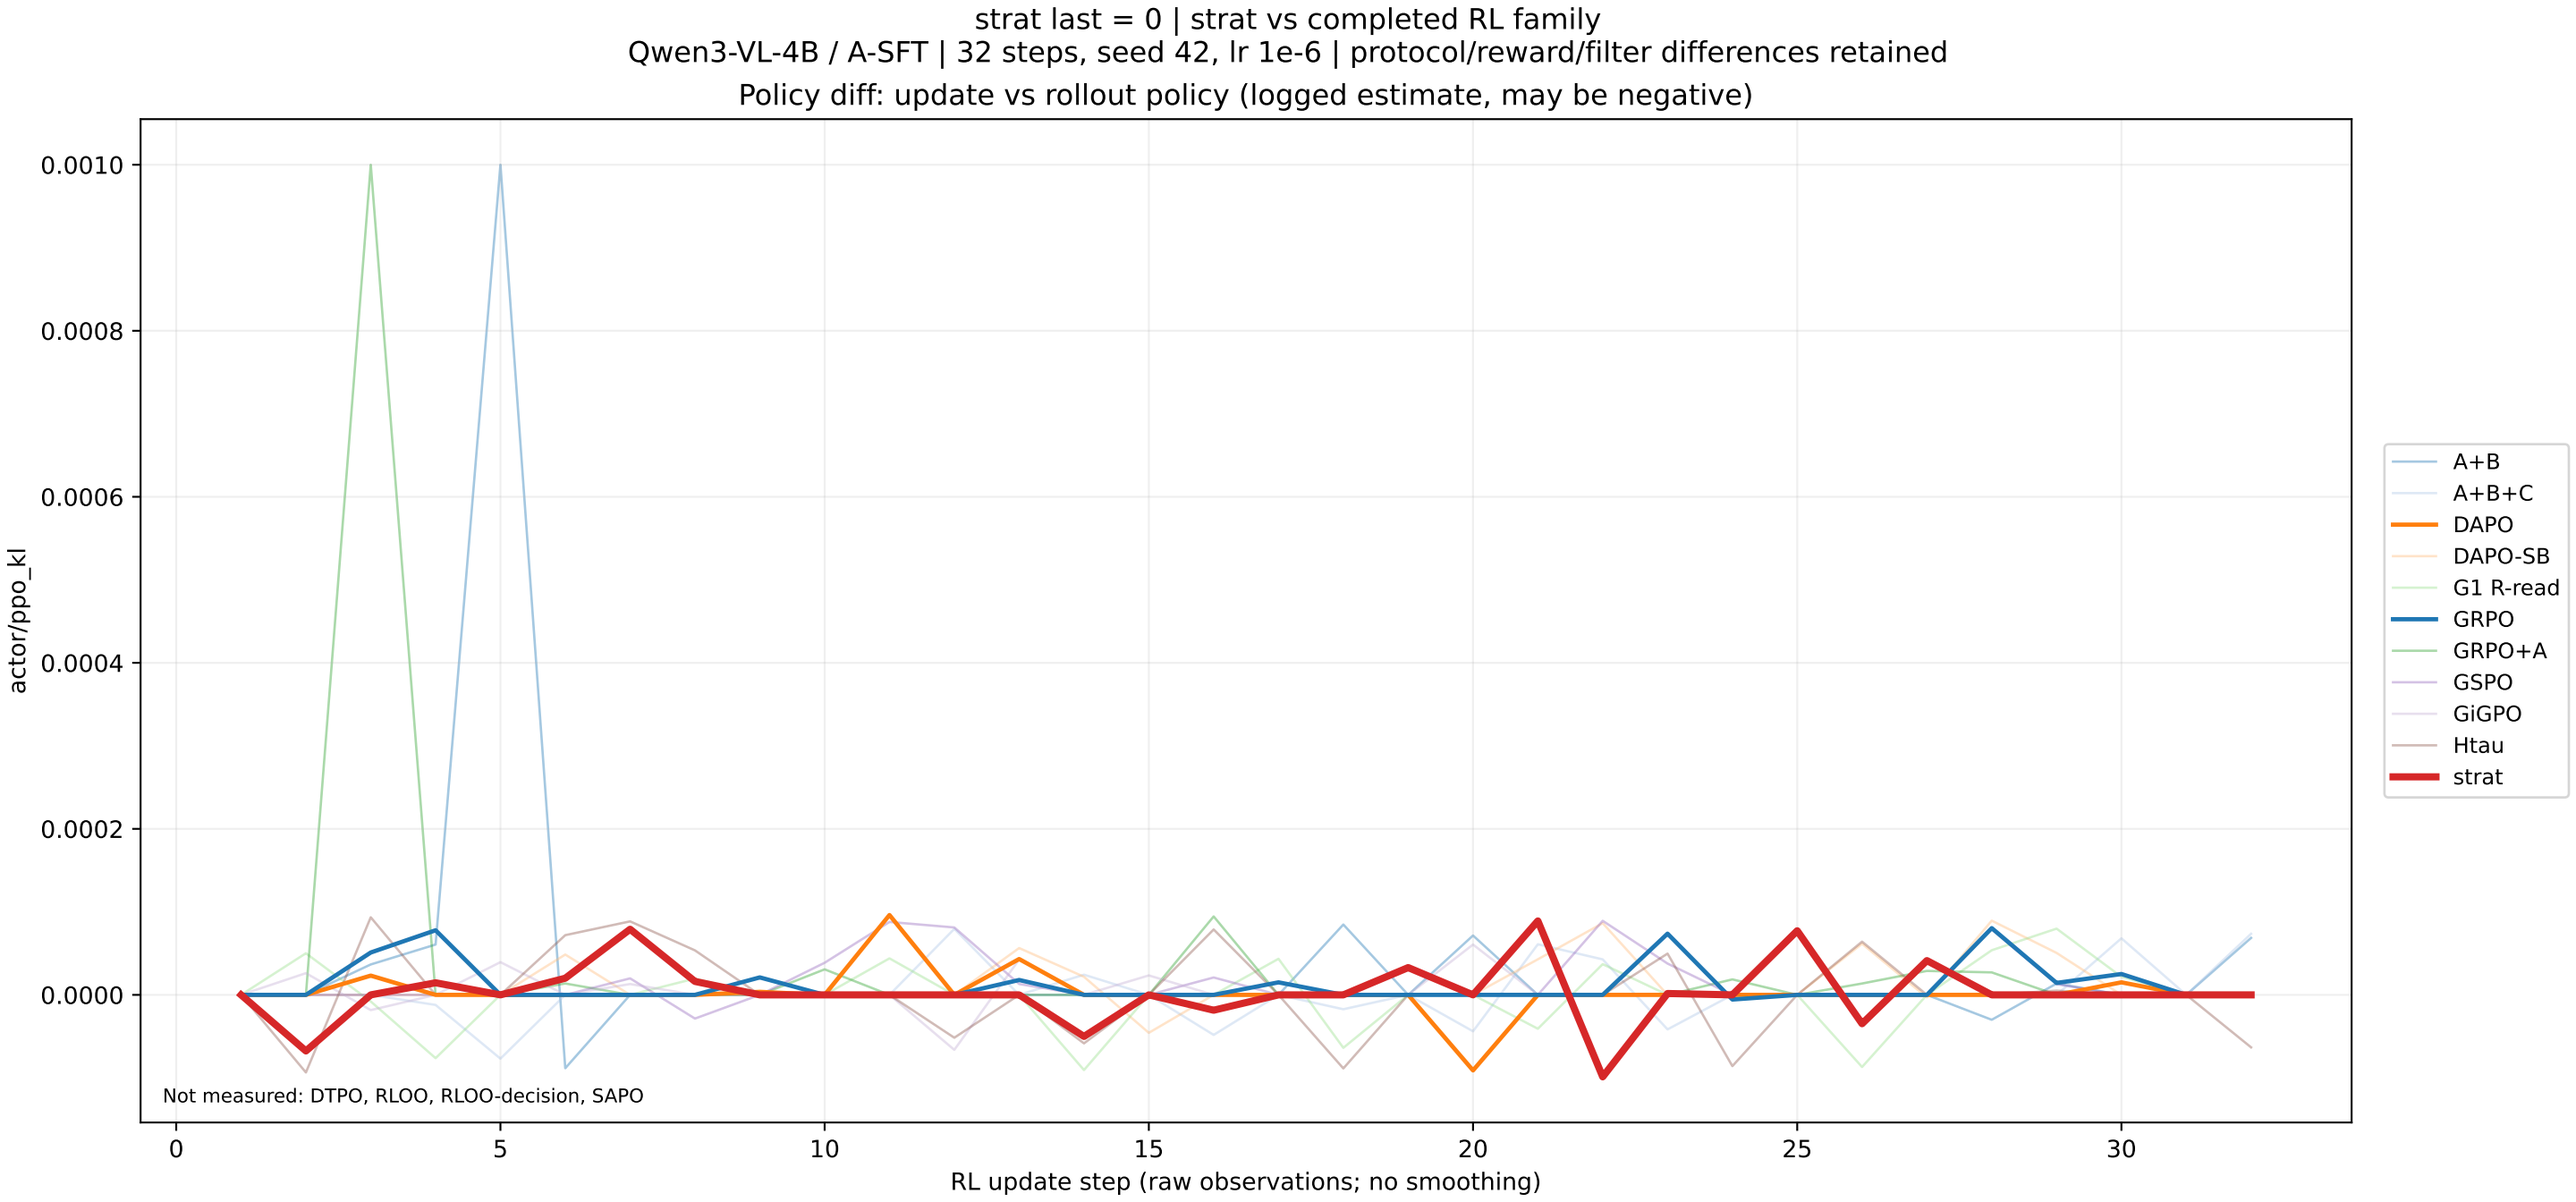
strat last = 0 | strat vs completed RL family
Qwen3-VL-4B / A-SFT | 32 steps, seed 42, lr 1e-6 | protocol/reward/filter differences retained
Policy diff: update vs rollout policy (logged estimate, may be negative)
actor/ppo_kl
0.0010
0.0008
0.0006
0.0004
0.0002
0.0000
0
5
10
15
20
25
30
RL update step (raw observations; no smoothing)
Not measured: DTPO, RLOO, RLOO-decision, SAPO
A+B
A+B+C
DAPO
DAPO-SB
G1 R-read
GRPO
GRPO+A
GSPO
GiGPO
Htau
strat