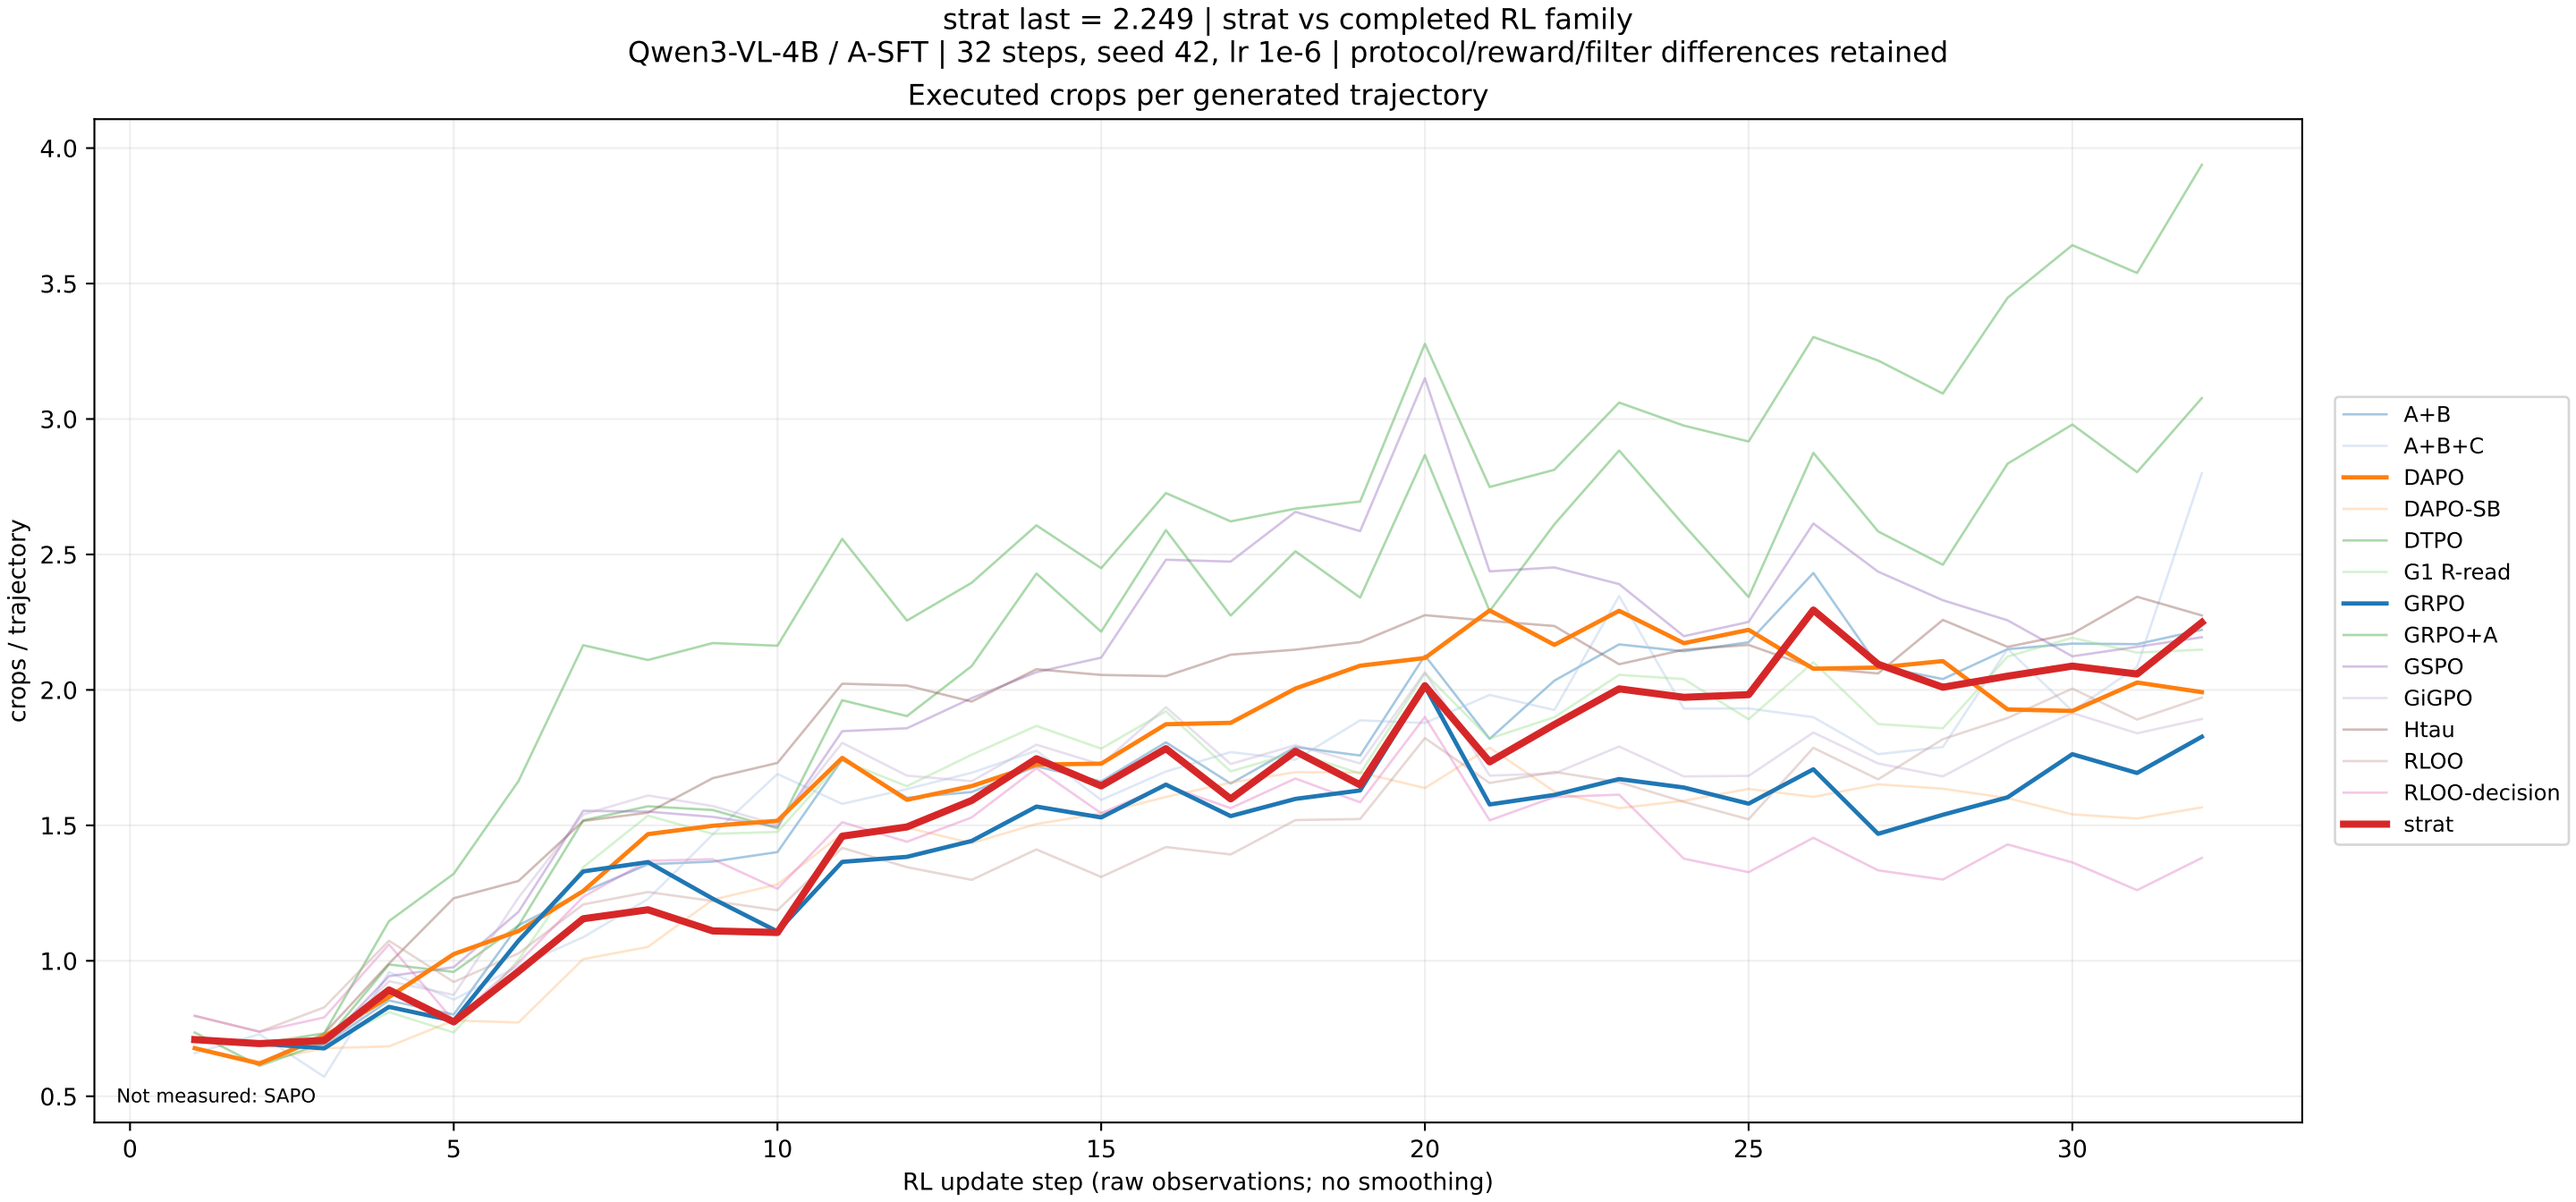
strat last = 2.249 | strat vs completed RL family
Qwen3-VL-4B / A-SFT | 32 steps, seed 42, lr 1e-6 | protocol/reward/filter differences retained
Executed crops per generated trajectory
crops / trajectory
RL update step (raw observations; no smoothing)
Not measured: SAPO
A+B
A+B+C
DAPO
DAPO-SB
DTPO
G1 R-read
GRPO
GRPO+A
GSPO
GiGPO
Htau
RLOO
RLOO-decision
strat