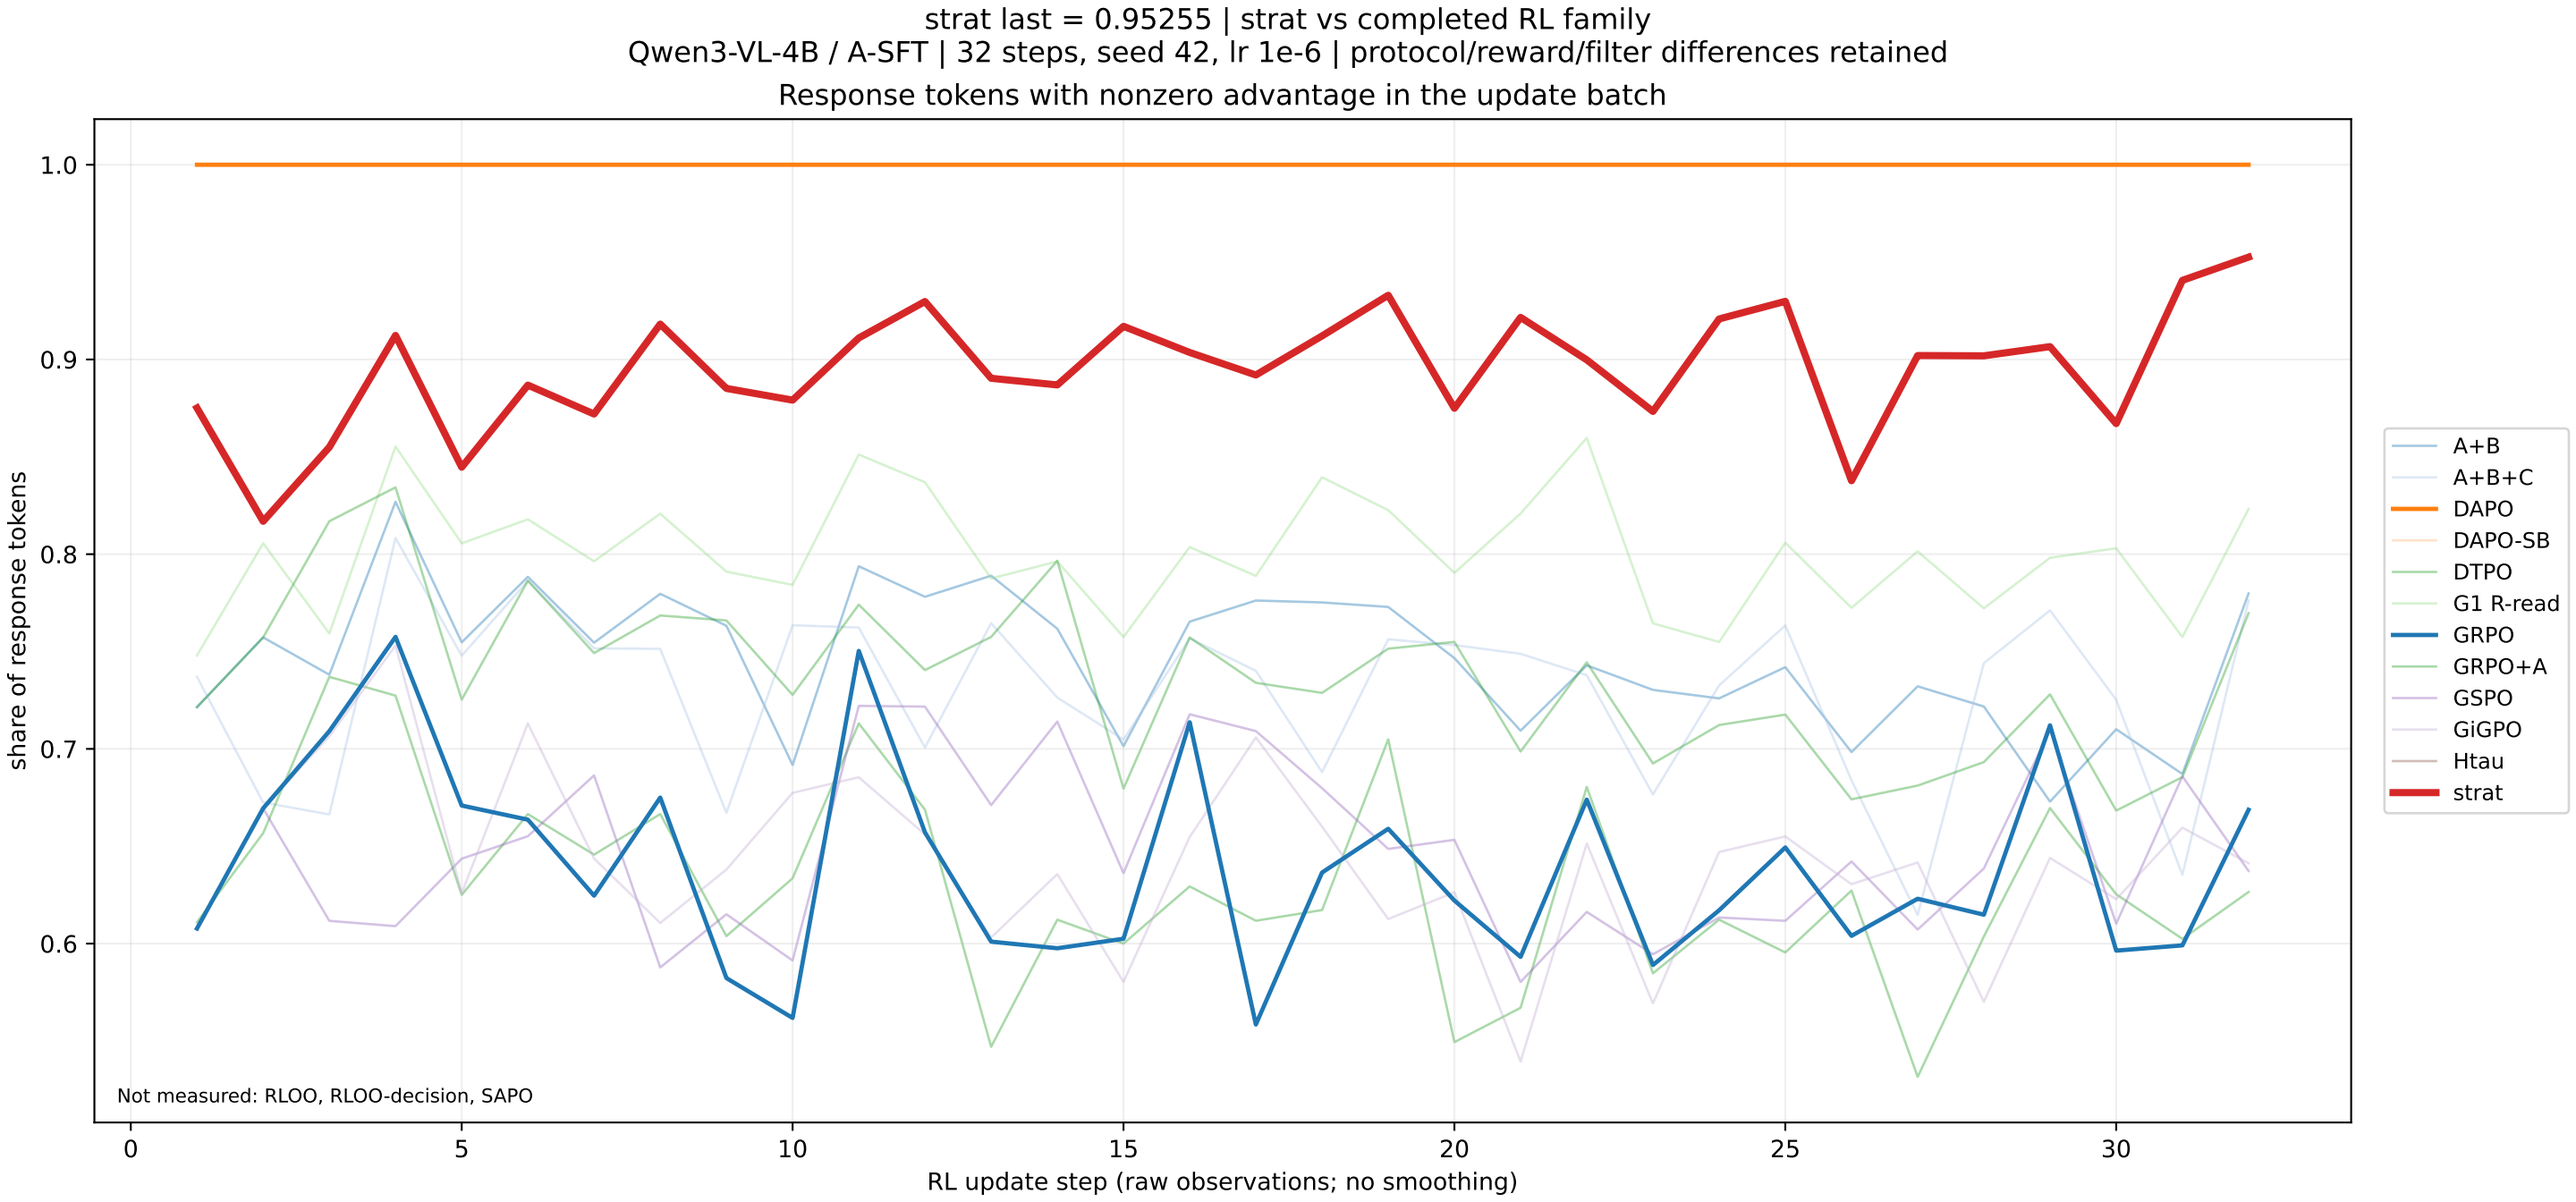
strat last = 0.95255 | strat vs completed RL family
Qwen3-VL-4B / A-SFT | 32 steps, seed 42, lr 1e-6 | protocol/reward/filter differences retained
Response tokens with nonzero advantage in the update batch
share of response tokens
RL update step (raw observations; no smoothing)
Not measured: RLOO, RLOO-decision, SAPO
A+B
A+B+C
DAPO
DAPO-SB
DTPO
G1 R-read
GRPO
GRPO+A
GSPO
GiGPO
Htau
strat