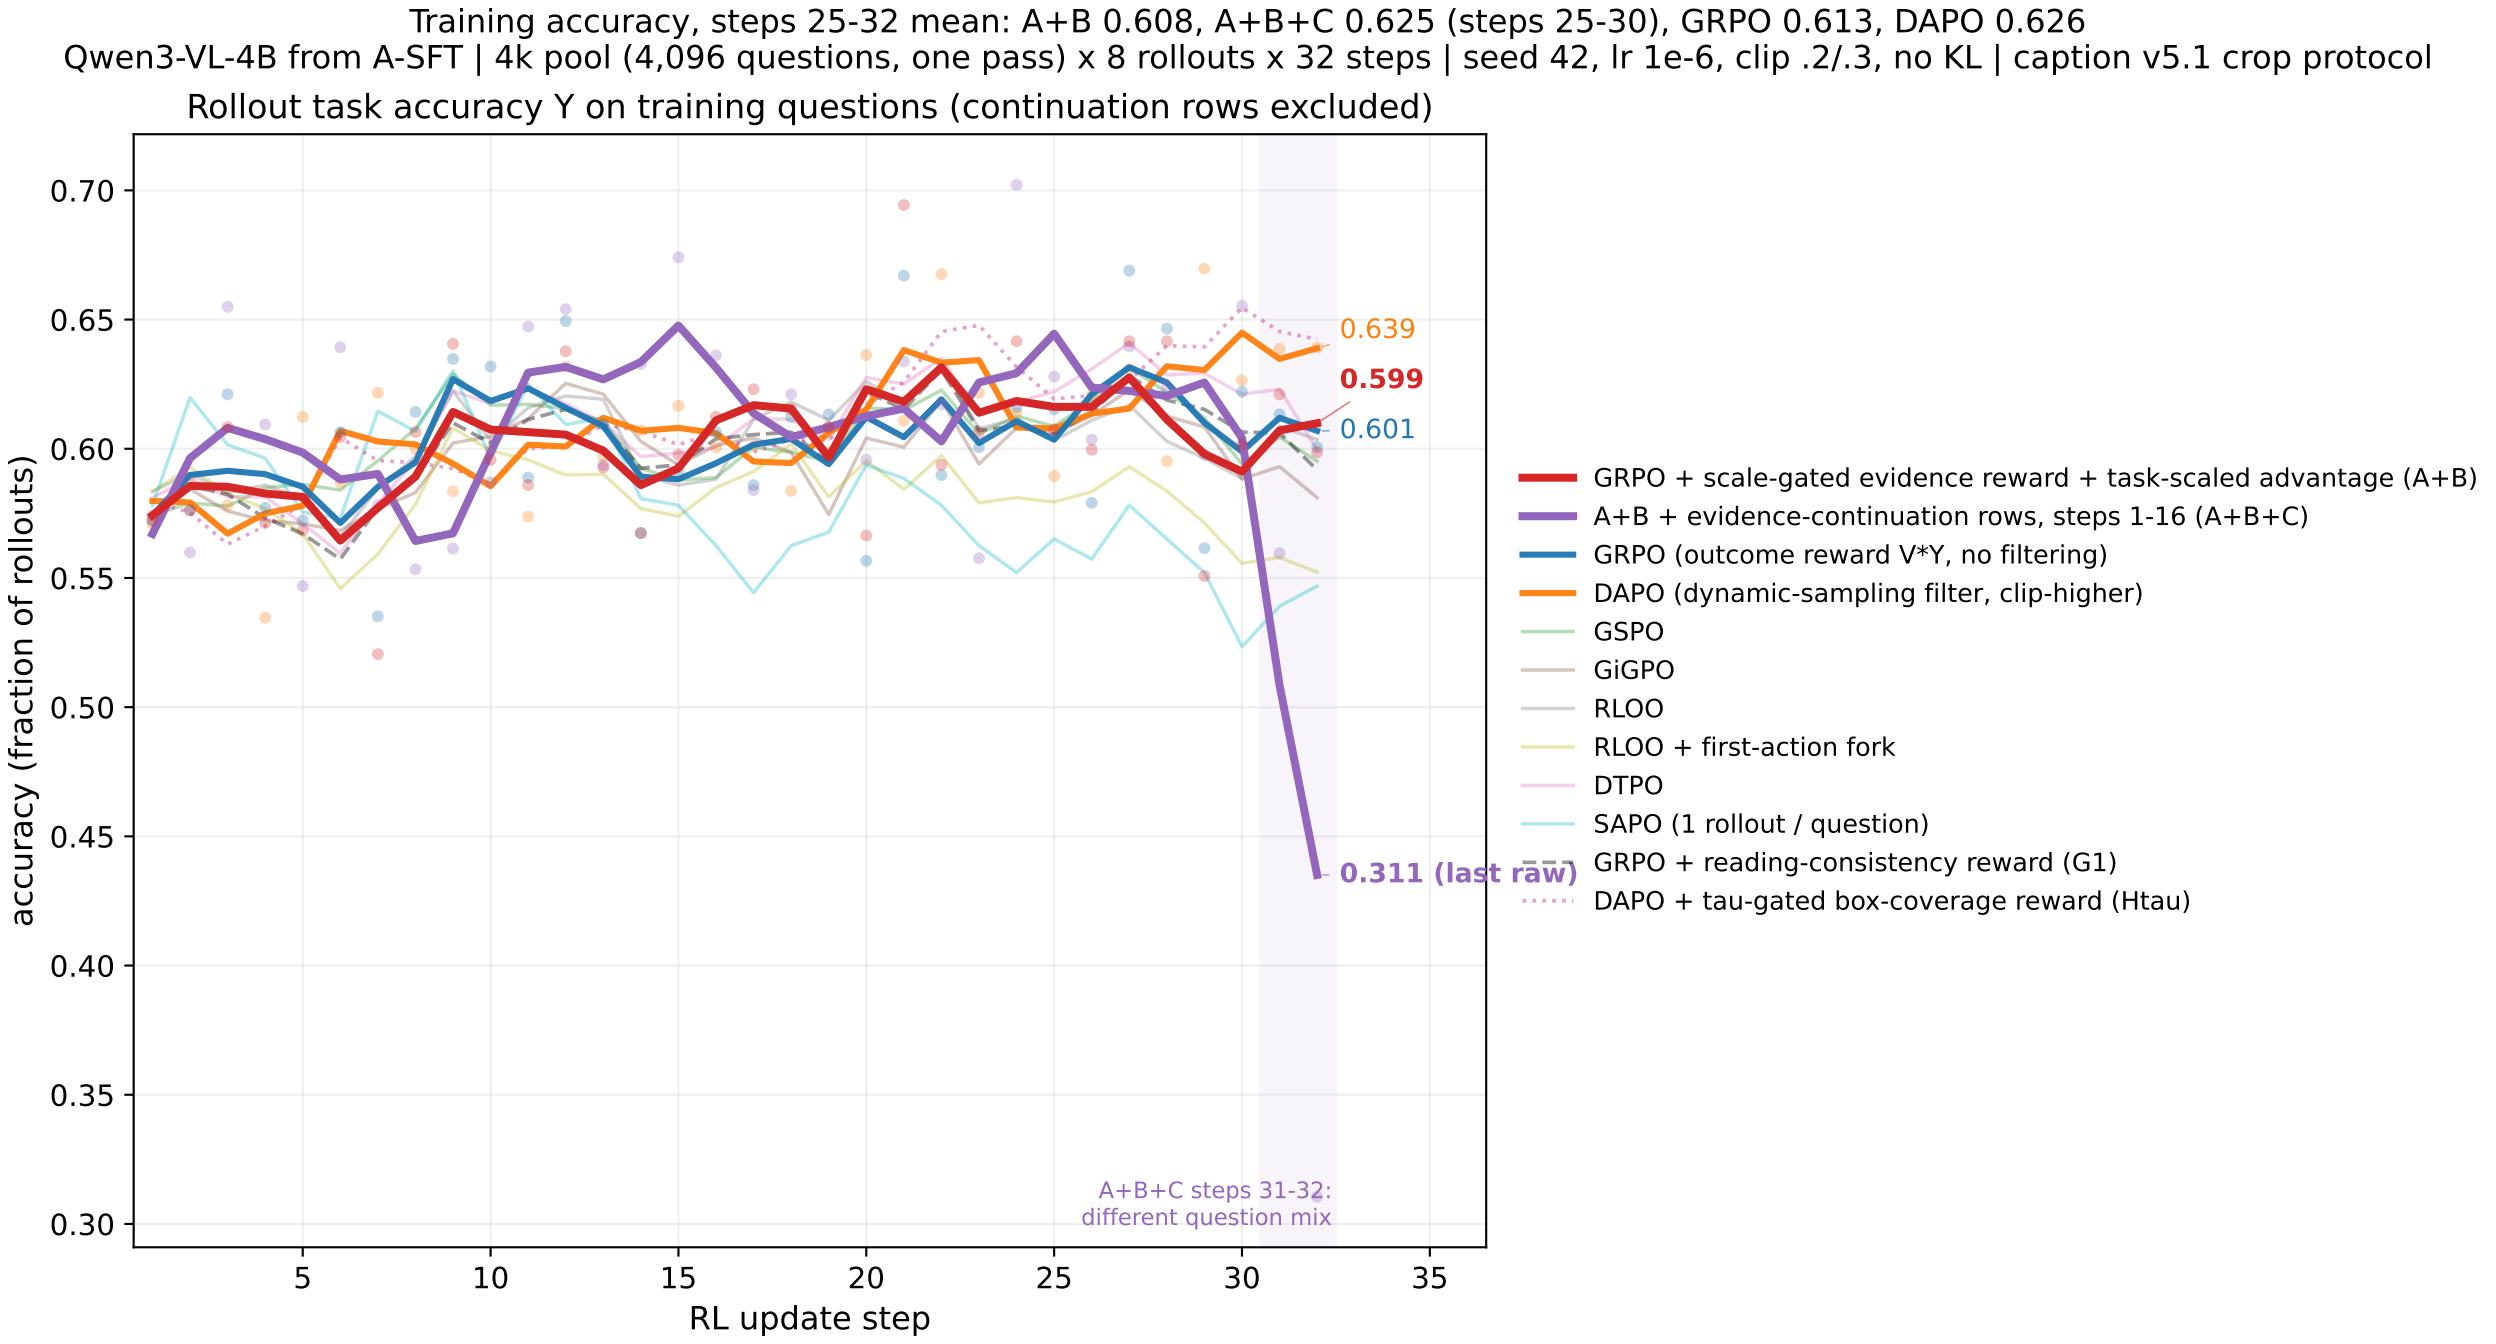
Training accuracy, steps 25-32 mean: A+B 0.608, A+B+C 0.625 (steps 25-30), GRPO 0.613, DAPO 0.626
Qwen3-VL-4B from A-SFT | 4k pool (4,096 questions, one pass) x 8 rollouts x 32 steps | seed 42, lr 1e-6, clip .2/.3, no KL | caption v5.1 crop protocol
Rollout task accuracy Y on training questions (continuation rows excluded)
accuracy (fraction of rollouts)
RL update step
0.70
0.65
0.60
0.55
0.50
0.45
0.40
0.35
0.30
5
10
15
20
25
30
35
0.639
0.599
0.601
0.311 (last raw)
A+B+C steps 31-32:
different question mix
GRPO + scale-gated evidence reward + task-scaled advantage (A+B)
A+B + evidence-continuation rows, steps 1-16 (A+B+C)
GRPO (outcome reward V*Y, no filtering)
DAPO (dynamic-sampling filter, clip-higher)
GSPO
GiGPO
RLOO
RLOO + first-action fork
DTPO
SAPO (1 rollout / question)
GRPO + reading-consistency reward (G1)
DAPO + tau-gated box-coverage reward (Htau)