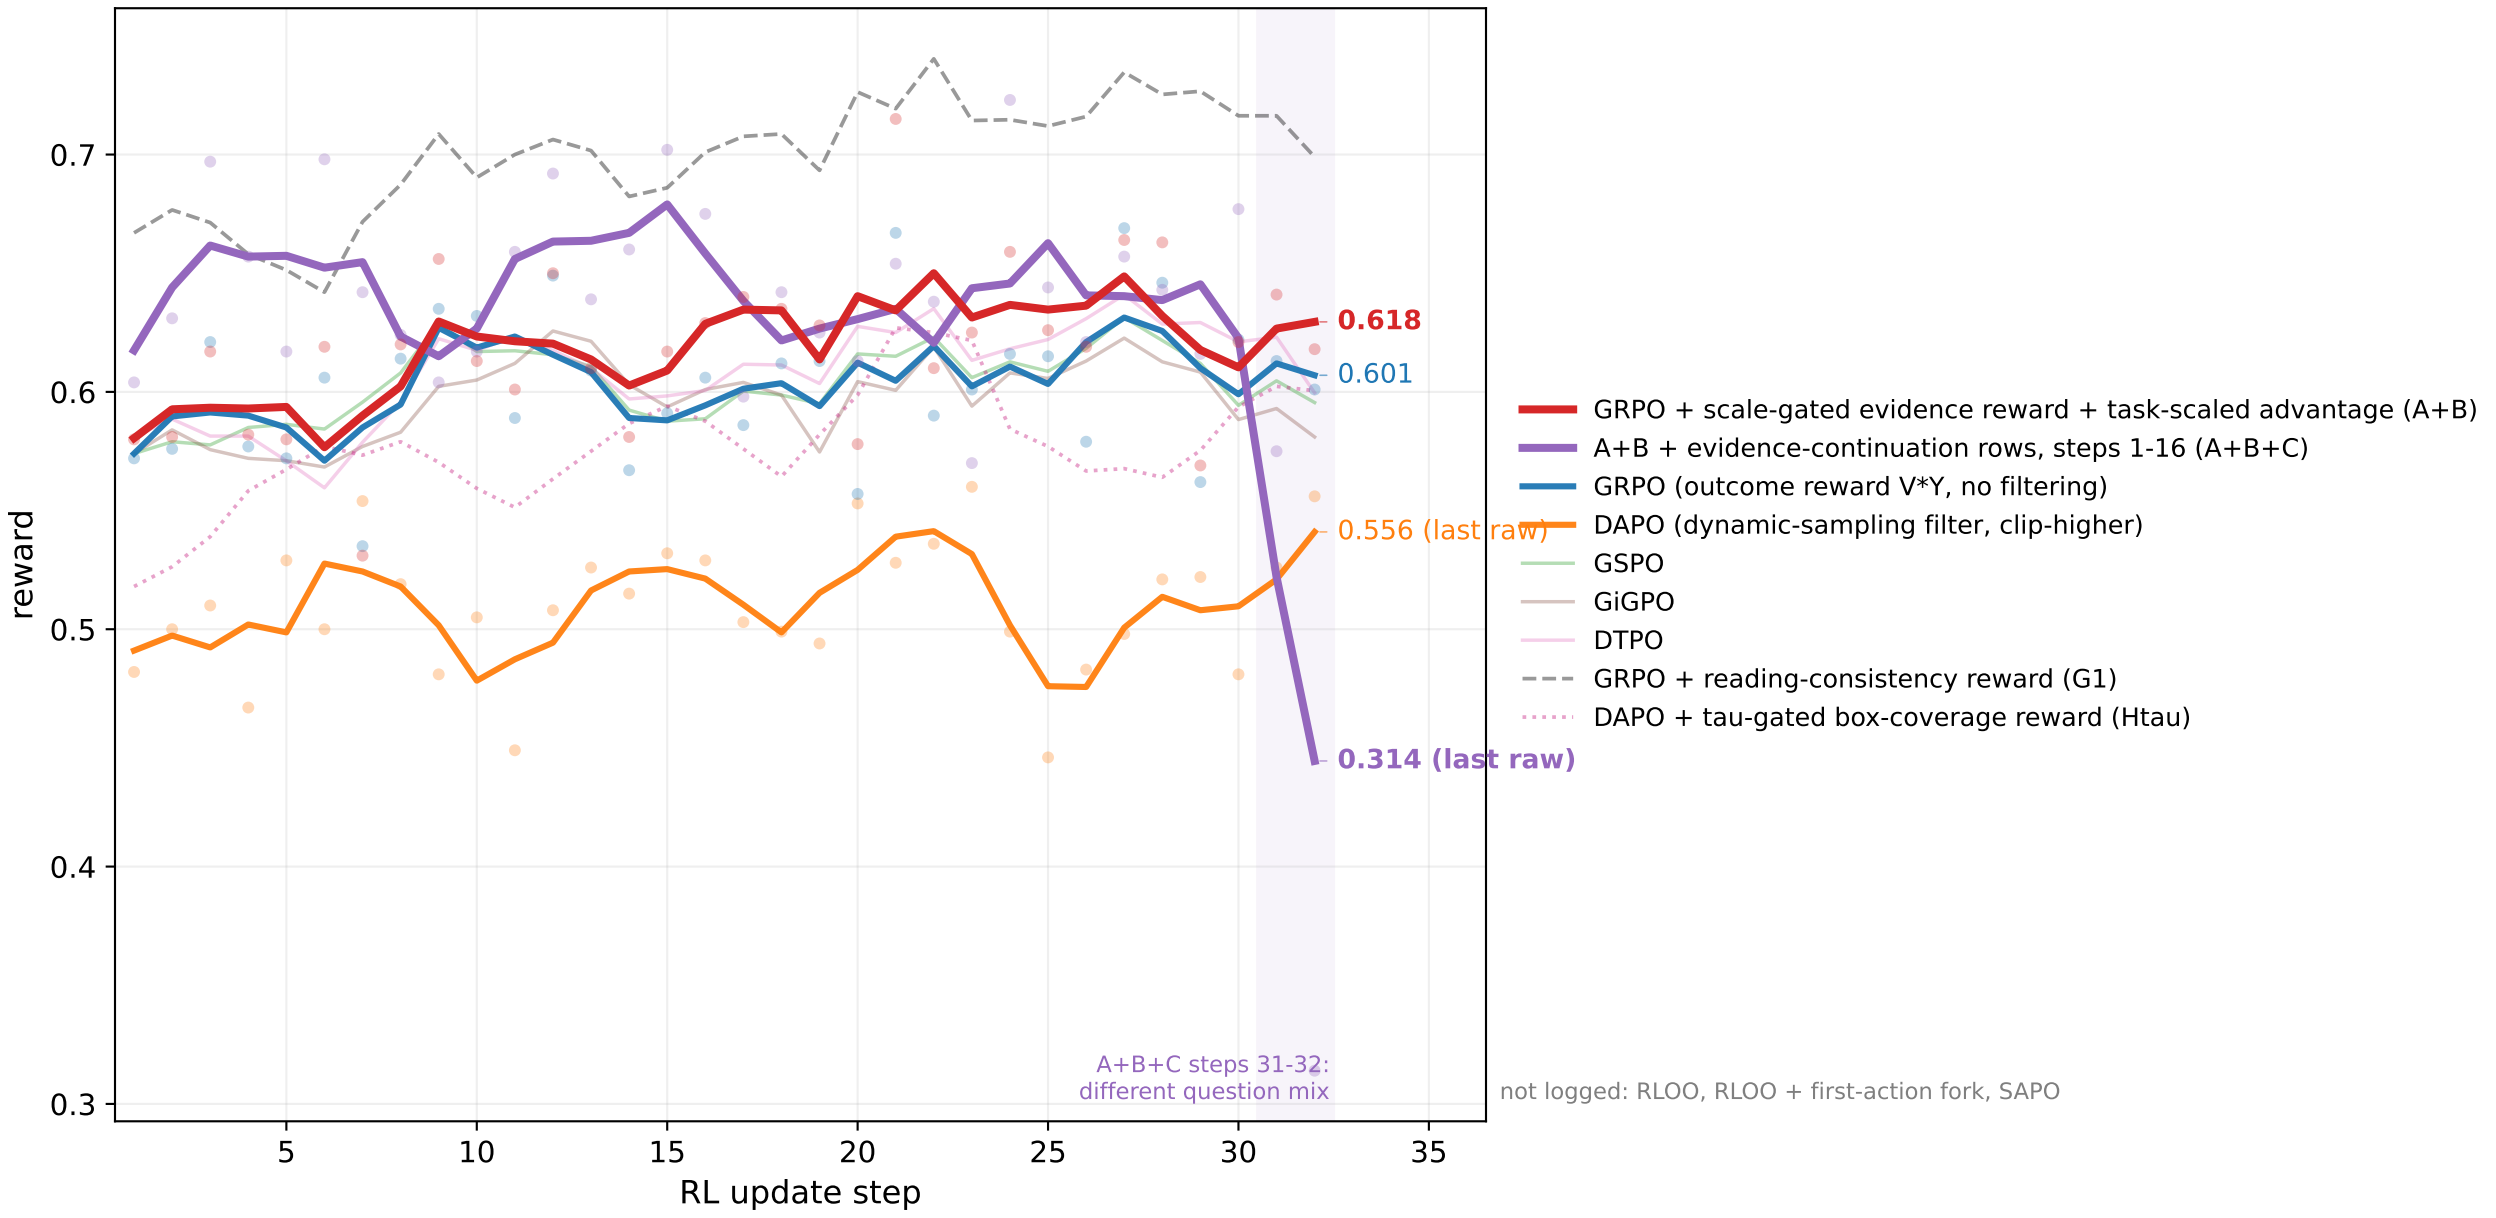Logged reward follows accuracy; A+B ends 0.618 vs GRPO 0.601; A+B+C step 32 = 0.314 (question mix, see shading)
Qwen3-VL-4B from A-SFT | 4k pool (4,096 questions, one pass) x 8 rollouts x 32 steps | seed 42, lr 1e-6, clip .2/.3, no KL | caption v5.1 crop protocol