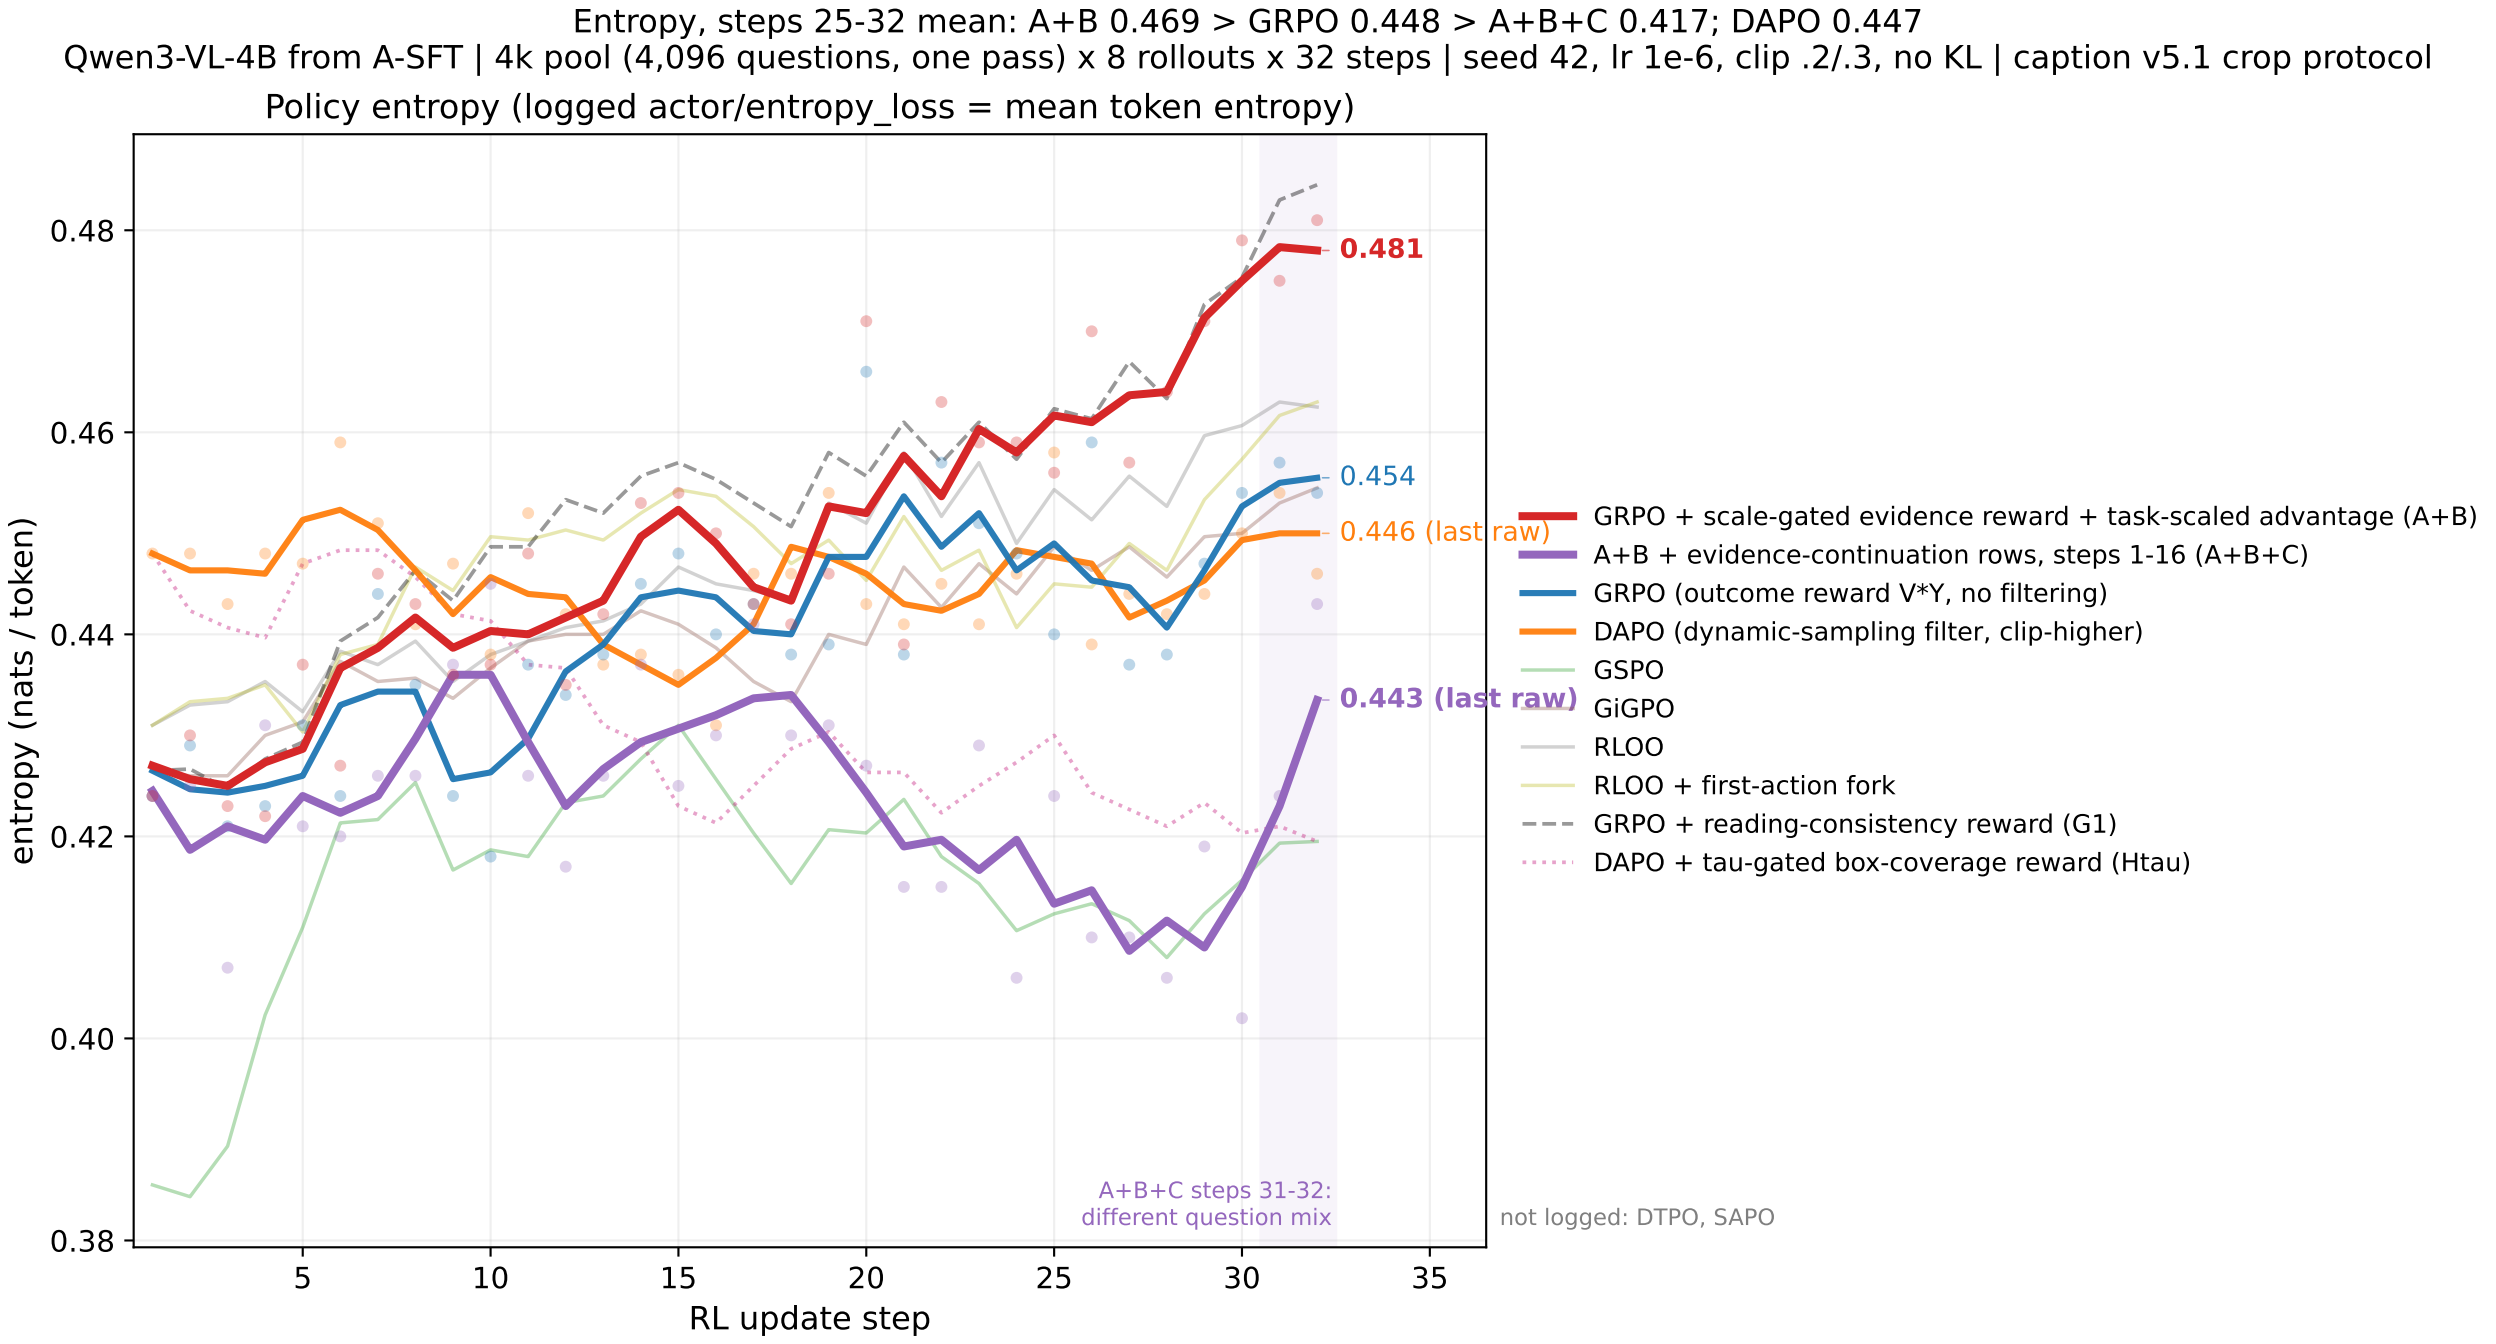
Entropy, steps 25-32 mean: A+B 0.469 > GRPO 0.448 > A+B+C 0.417; DAPO 0.447
Qwen3-VL-4B from A-SFT | 4k pool (4,096 questions, one pass) x 8 rollouts x 32 steps | seed 42, lr 1e-6, clip .2/.3, no KL | caption v5.1 crop protocol
Policy entropy (logged actor/entropy_loss = mean token entropy)
entropy (nats / token)
RL update step
0.481
0.454
0.446 (last raw)
0.443 (last raw)
A+B+C steps 31-32:
different question mix
GRPO + scale-gated evidence reward + task-scaled advantage (A+B)
A+B + evidence-continuation rows, steps 1-16 (A+B+C)
GRPO (outcome reward V*Y, no filtering)
DAPO (dynamic-sampling filter, clip-higher)
GSPO
GiGPO
RLOO
RLOO + first-action fork
GRPO + reading-consistency reward (G1)
DAPO + tau-gated box-coverage reward (Htau)
not logged: DTPO, SAPO