Policy diff stays within +-1e-4 for all logged runs except one A+B step (1.0e-3 at step 5)

Qwen3-VL-4B from A-SFT | 4k pool (4,096 questions, one pass) x 8 rollouts x 32 steps | seed 42, lr 1e-6, clip .2/.3, no KL | caption v5.1 crop protocol
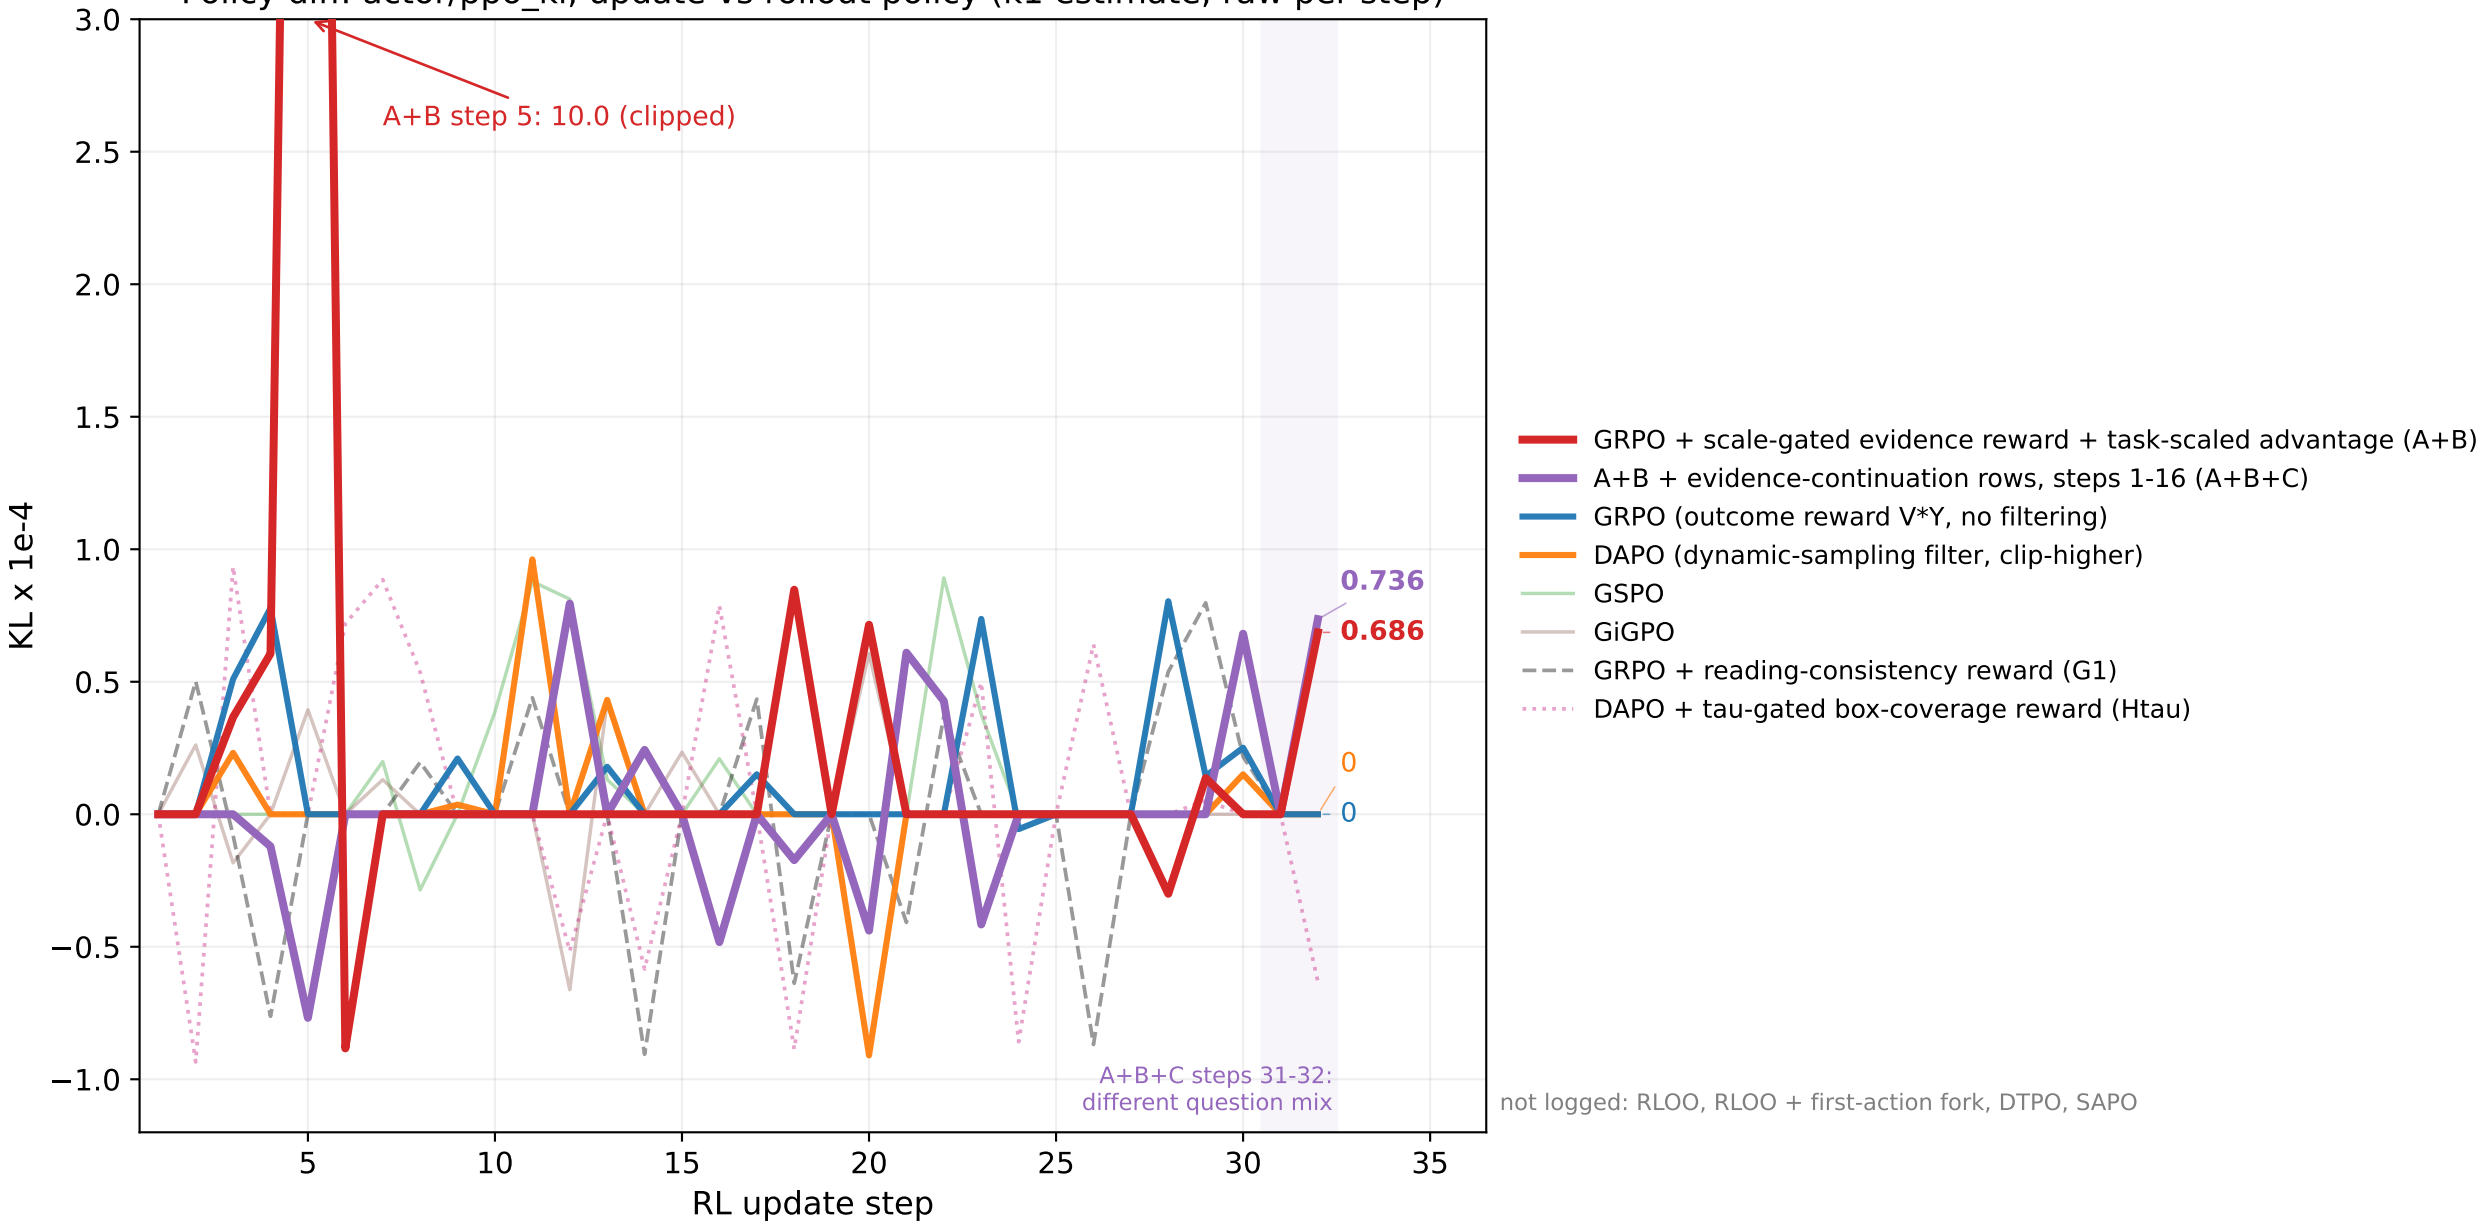
Policy diff: actor/ppo_kl, update vs rollout policy (k1 estimate, raw per step)