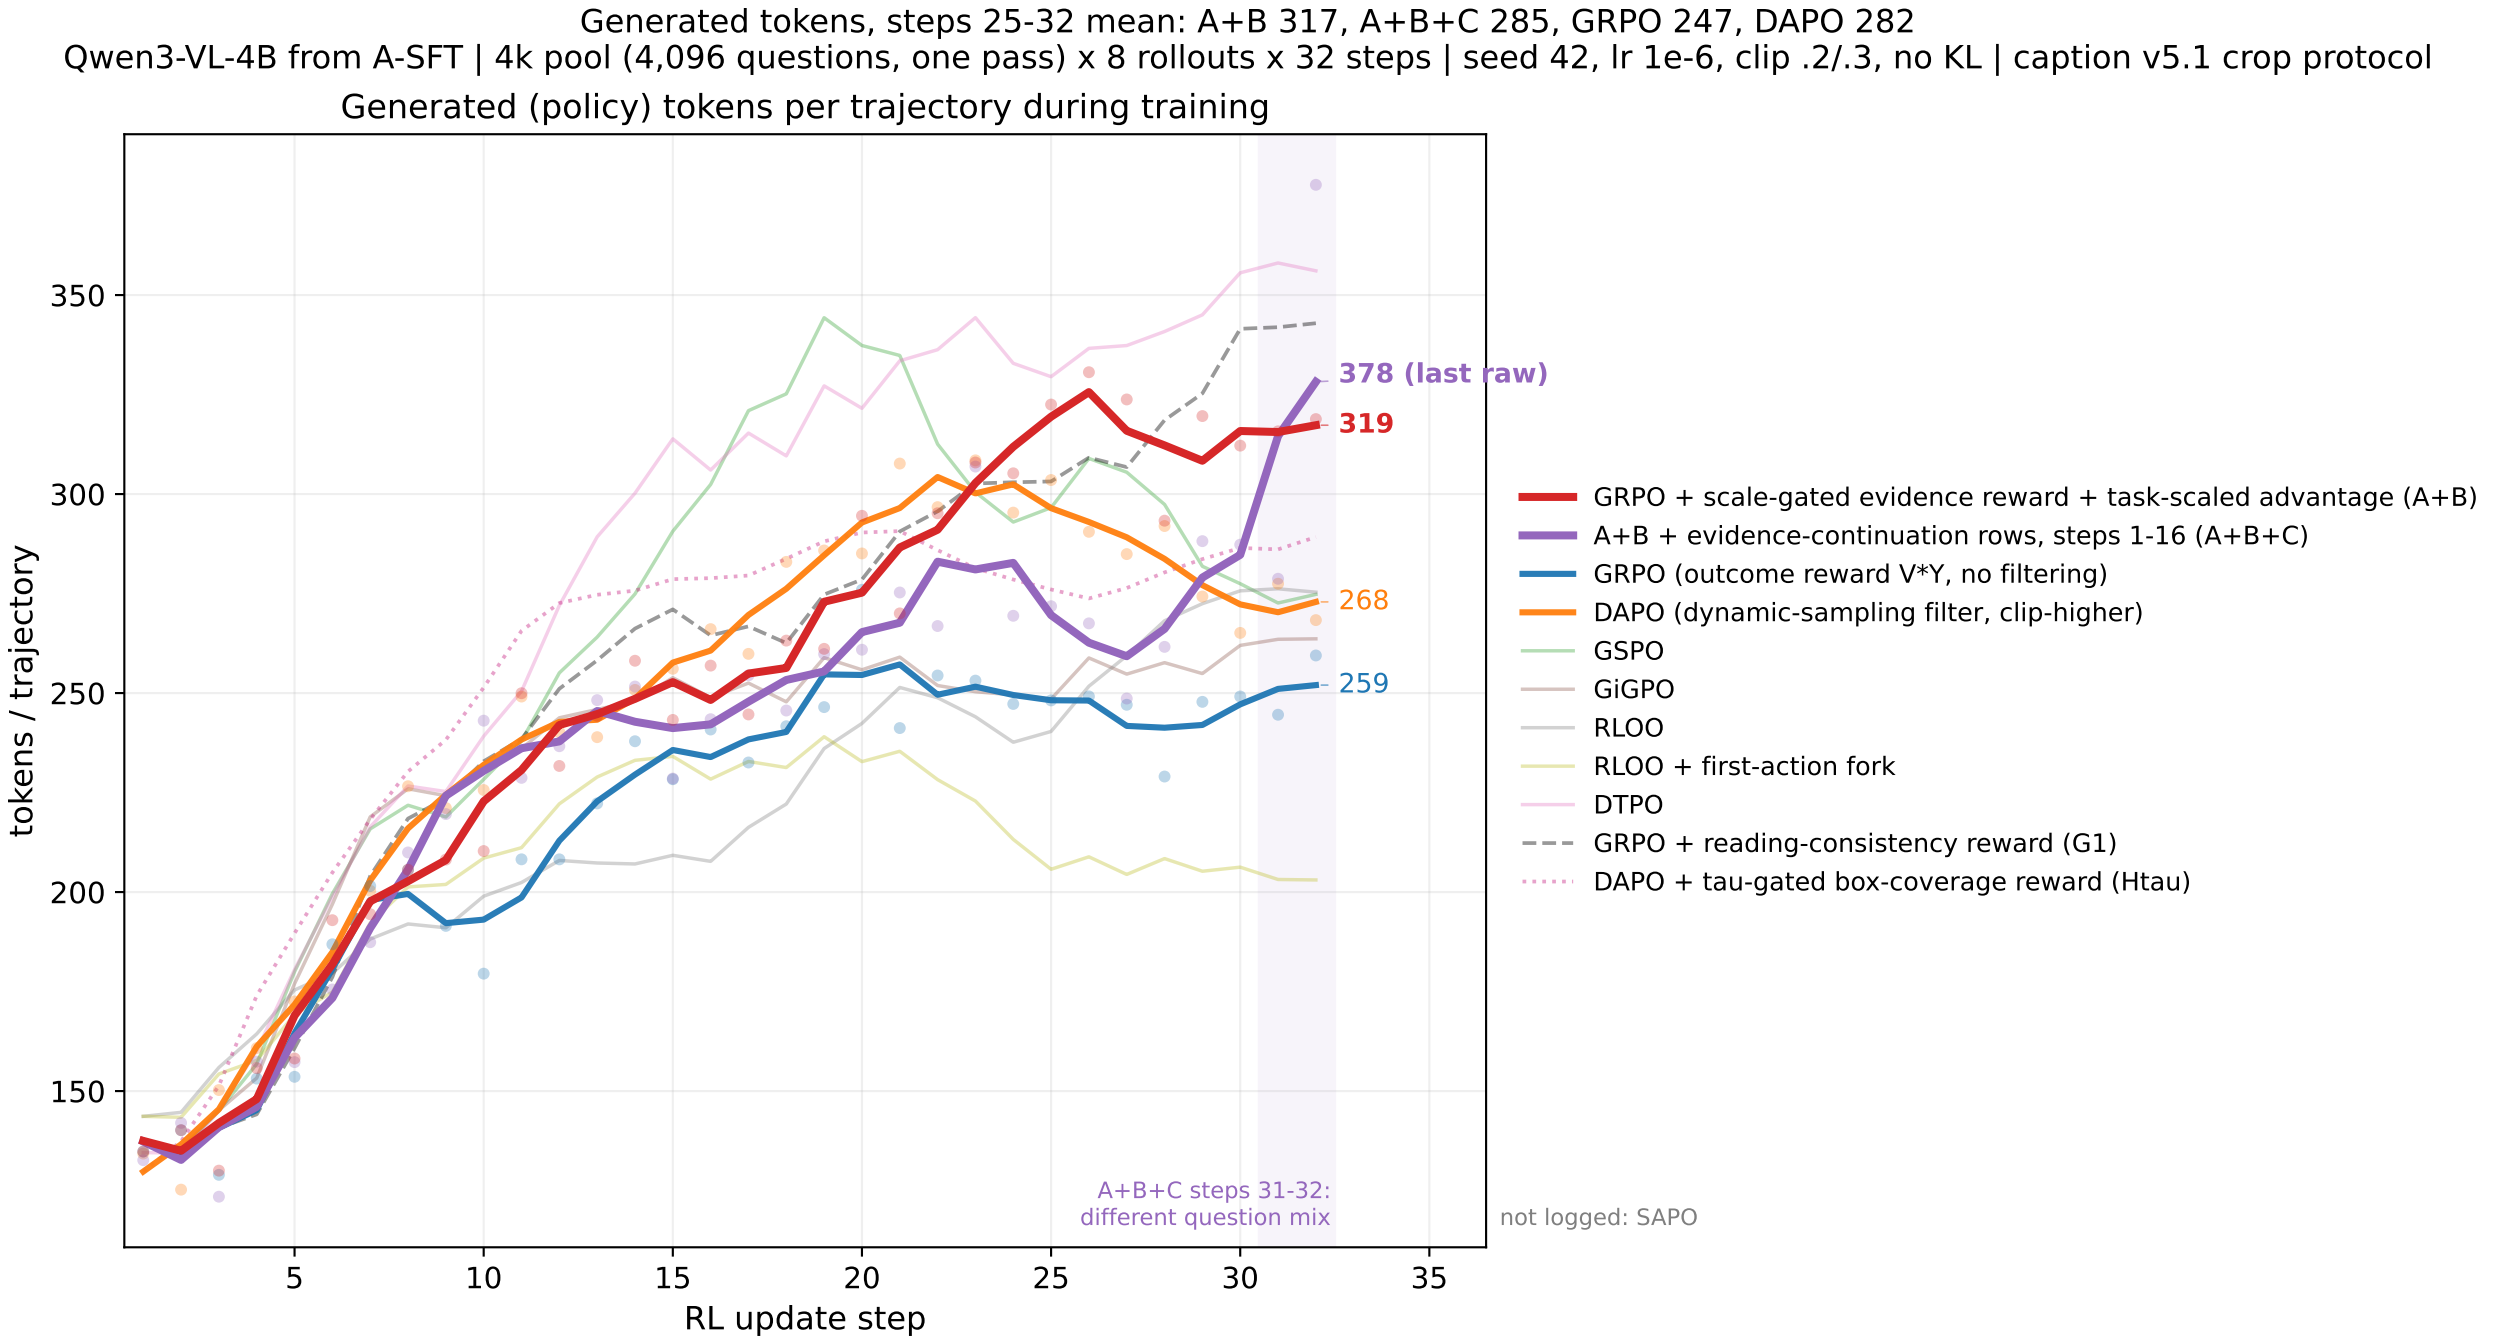
Generated tokens, steps 25-32 mean: A+B 317, A+B+C 285, GRPO 247, DAPO 282
Qwen3-VL-4B from A-SFT | 4k pool (4,096 questions, one pass) x 8 rollouts x 32 steps | seed 42, lr 1e-6, clip .2/.3, no KL | caption v5.1 crop protocol
Generated (policy) tokens per trajectory during training
tokens / trajectory
RL update step
378 (last raw)
319
268
259
A+B+C steps 31-32:
different question mix
GRPO + scale-gated evidence reward + task-scaled advantage (A+B)
A+B + evidence-continuation rows, steps 1-16 (A+B+C)
GRPO (outcome reward V*Y, no filtering)
DAPO (dynamic-sampling filter, clip-higher)
GSPO
GiGPO
RLOO
RLOO + first-action fork
DTPO
GRPO + reading-consistency reward (G1)
DAPO + tau-gated box-coverage reward (Htau)
not logged: SAPO