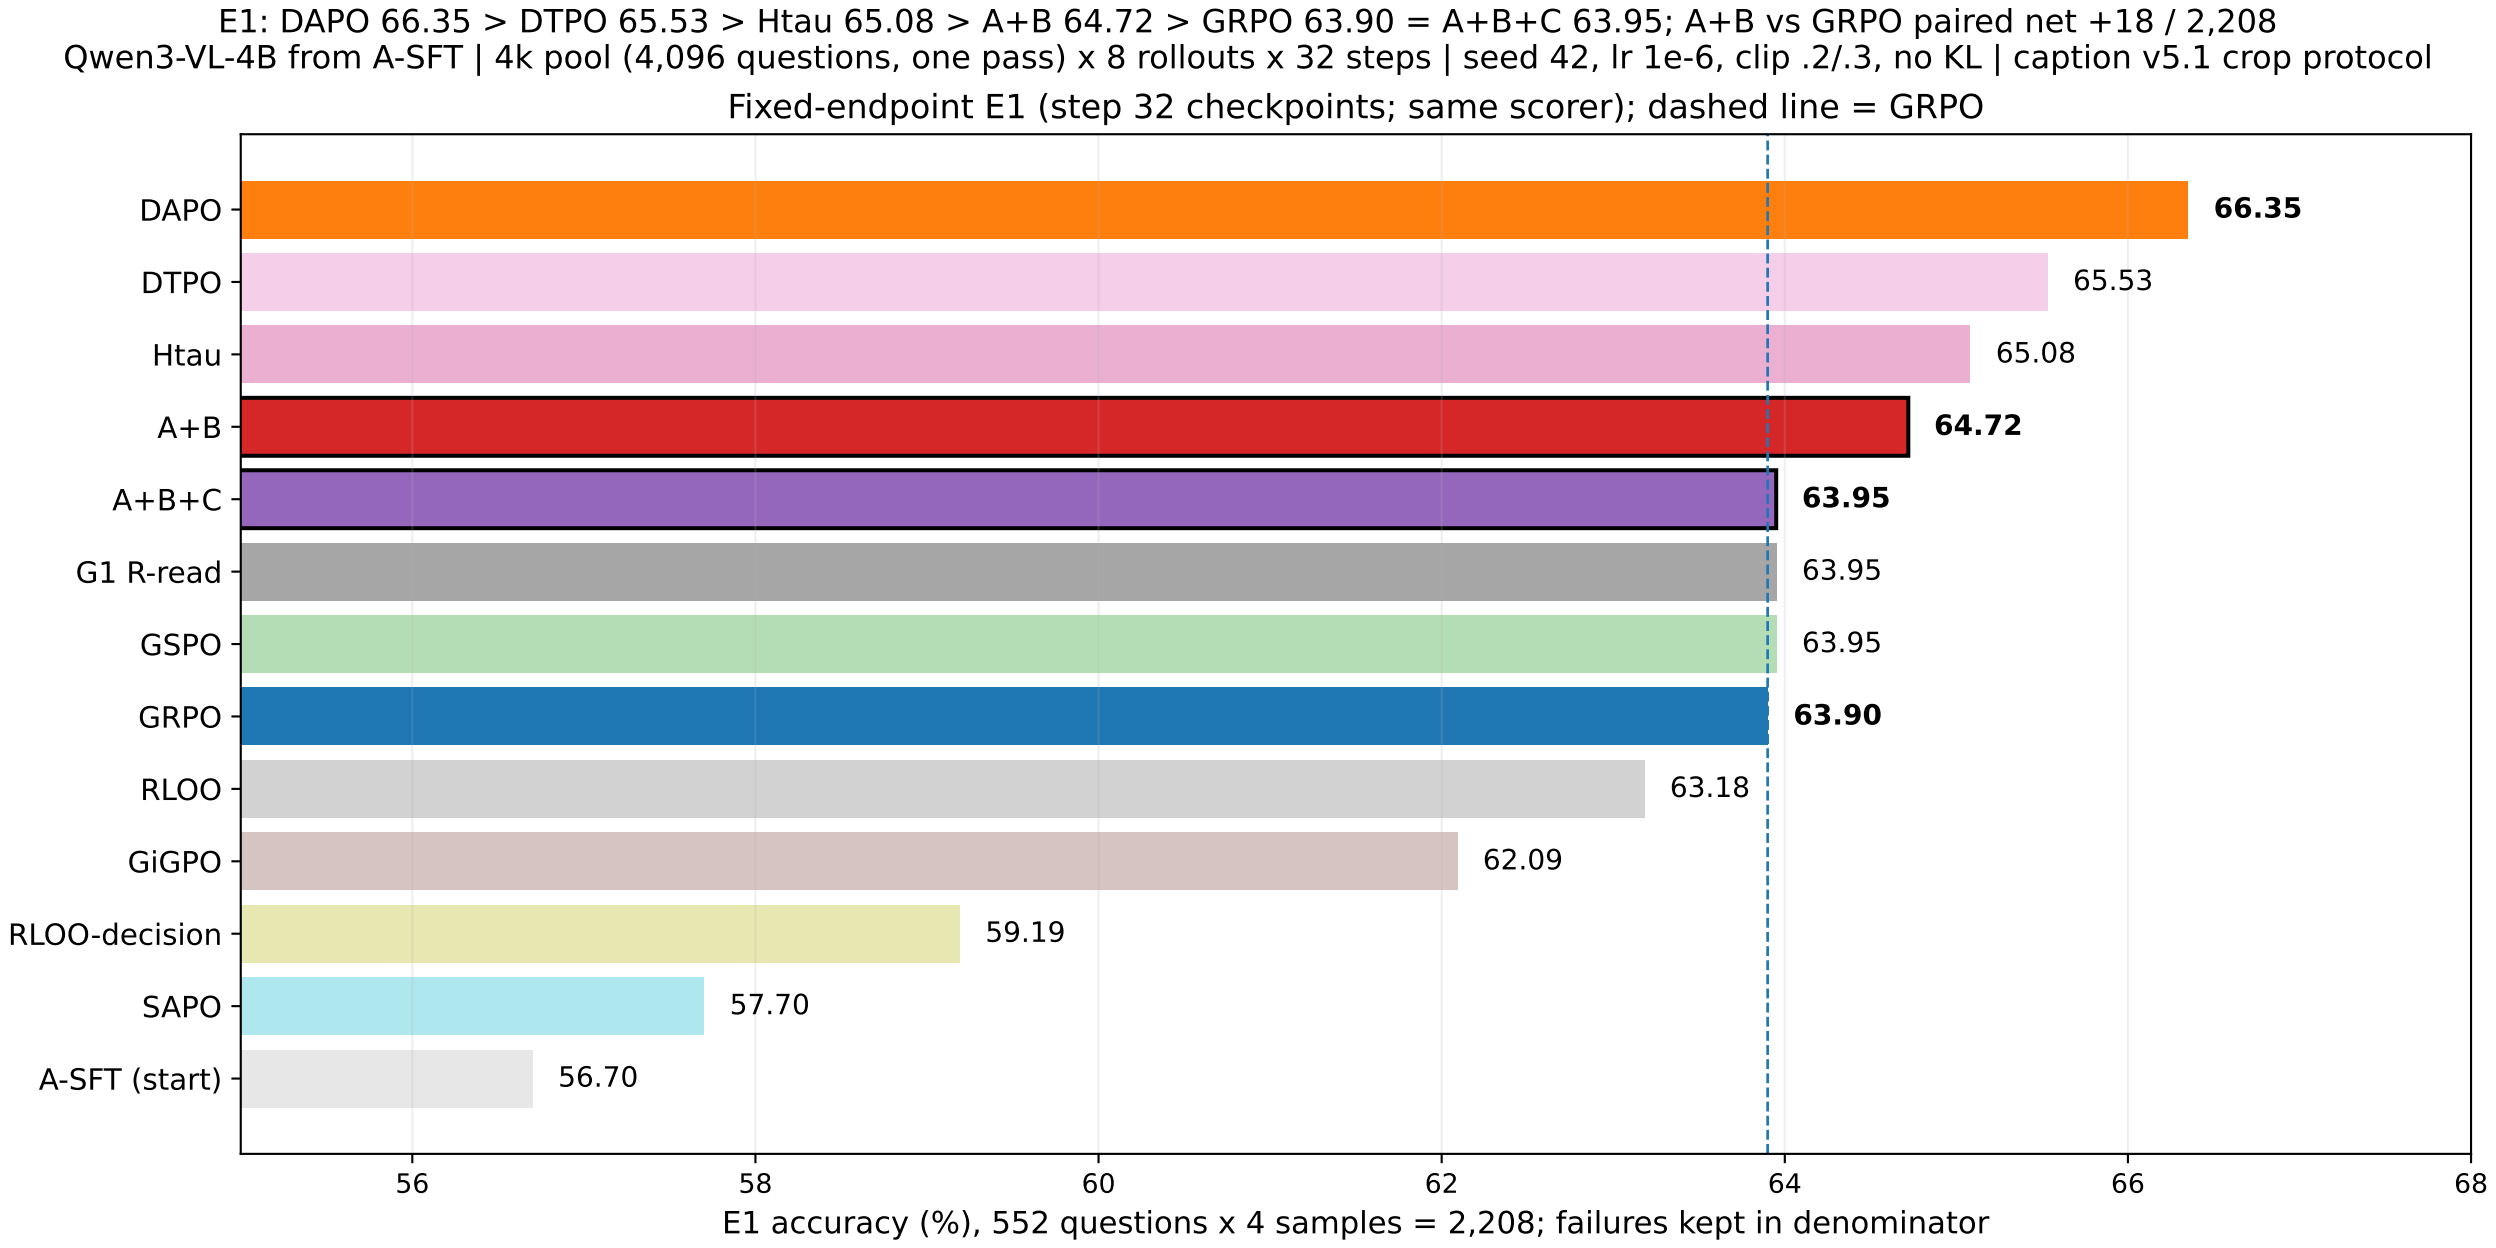
E1: DAPO 66.35 > DTPO 65.53 > Htau 65.08 > A+B 64.72 > GRPO 63.90 = A+B+C 63.95; A+B vs GRPO paired net +18 / 2,208
Qwen3-VL-4B from A-SFT | 4k pool (4,096 questions, one pass) x 8 rollouts x 32 steps | seed 42, lr 1e-6, clip .2/.3, no KL | caption v5.1 crop protocol
Fixed-endpoint E1 (step 32 checkpoints; same scorer); dashed line = GRPO
DAPO
DTPO
Htau
A+B
A+B+C
G1 R-read
GSPO
GRPO
RLOO
GiGPO
RLOO-decision
SAPO
A-SFT (start)
66.35
65.53
65.08
64.72
63.95
63.95
63.95
63.90
63.18
62.09
59.19
57.70
56.70
56
58
60
62
64
66
68
E1 accuracy (%), 552 questions x 4 samples = 2,208; failures kept in denominator