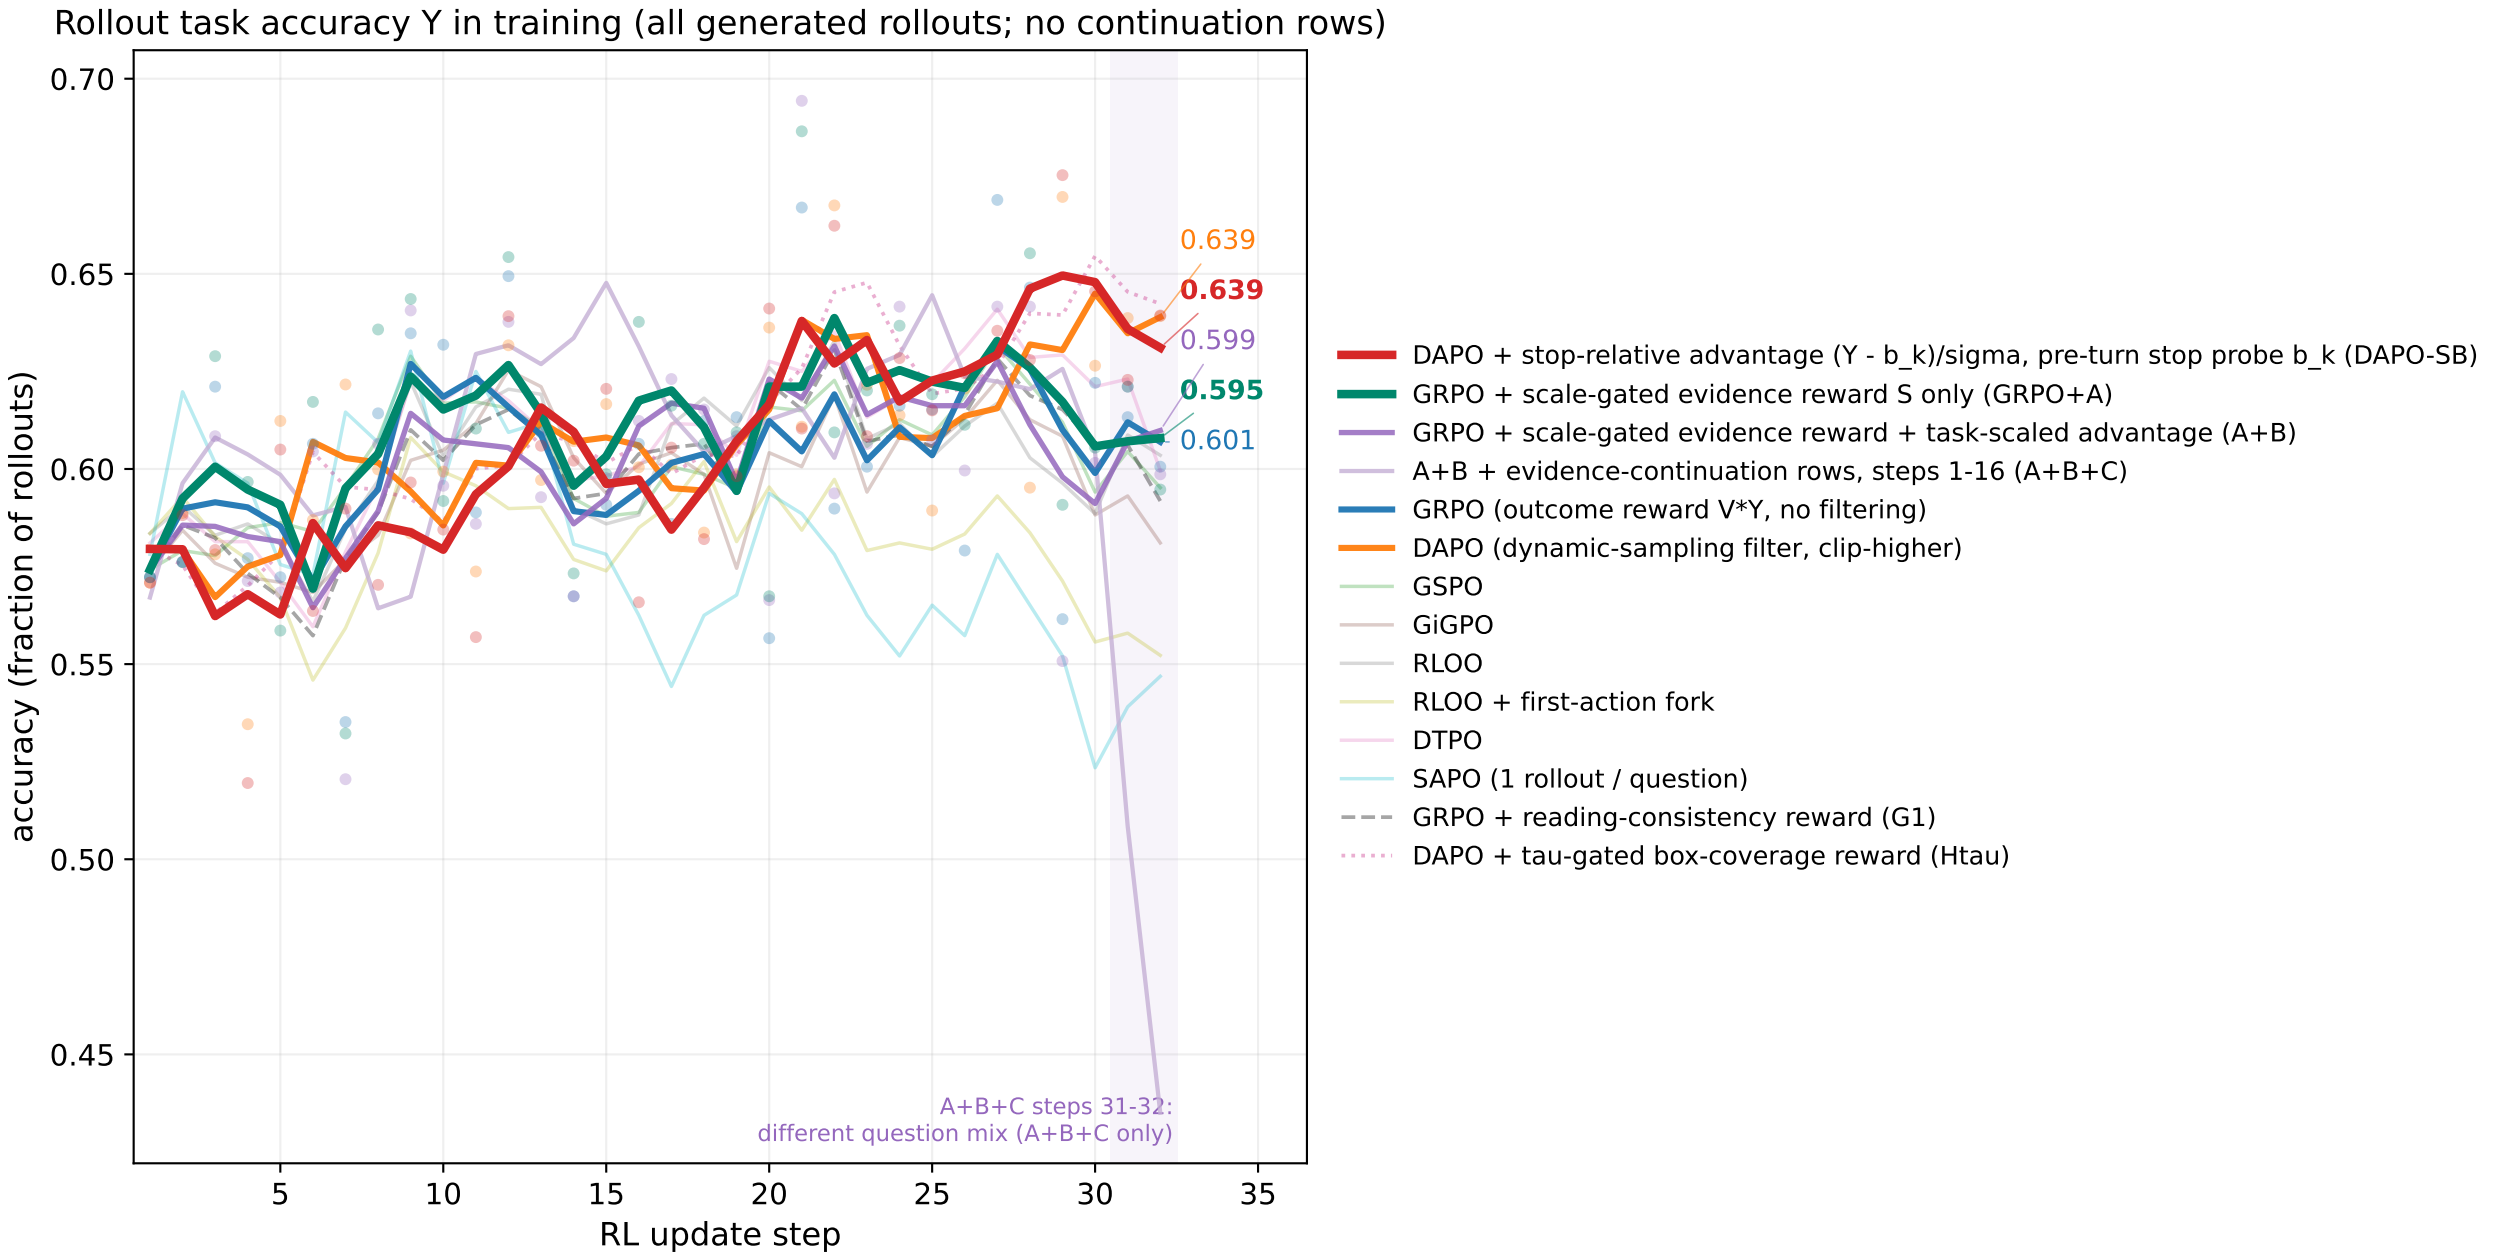Training accuracy, steps 25-32 mean: DAPO-SB 0.636 vs DAPO 0.626; GRPO+A 0.616 vs GRPO 0.613; A+B 0.608
Qwen3-VL-4B from A-SFT | 4k pool (one pass) x 8 rollouts x 32 steps | seed 42, lr 1e-6, no KL | caption v5.1 crop protocol | one seed per run