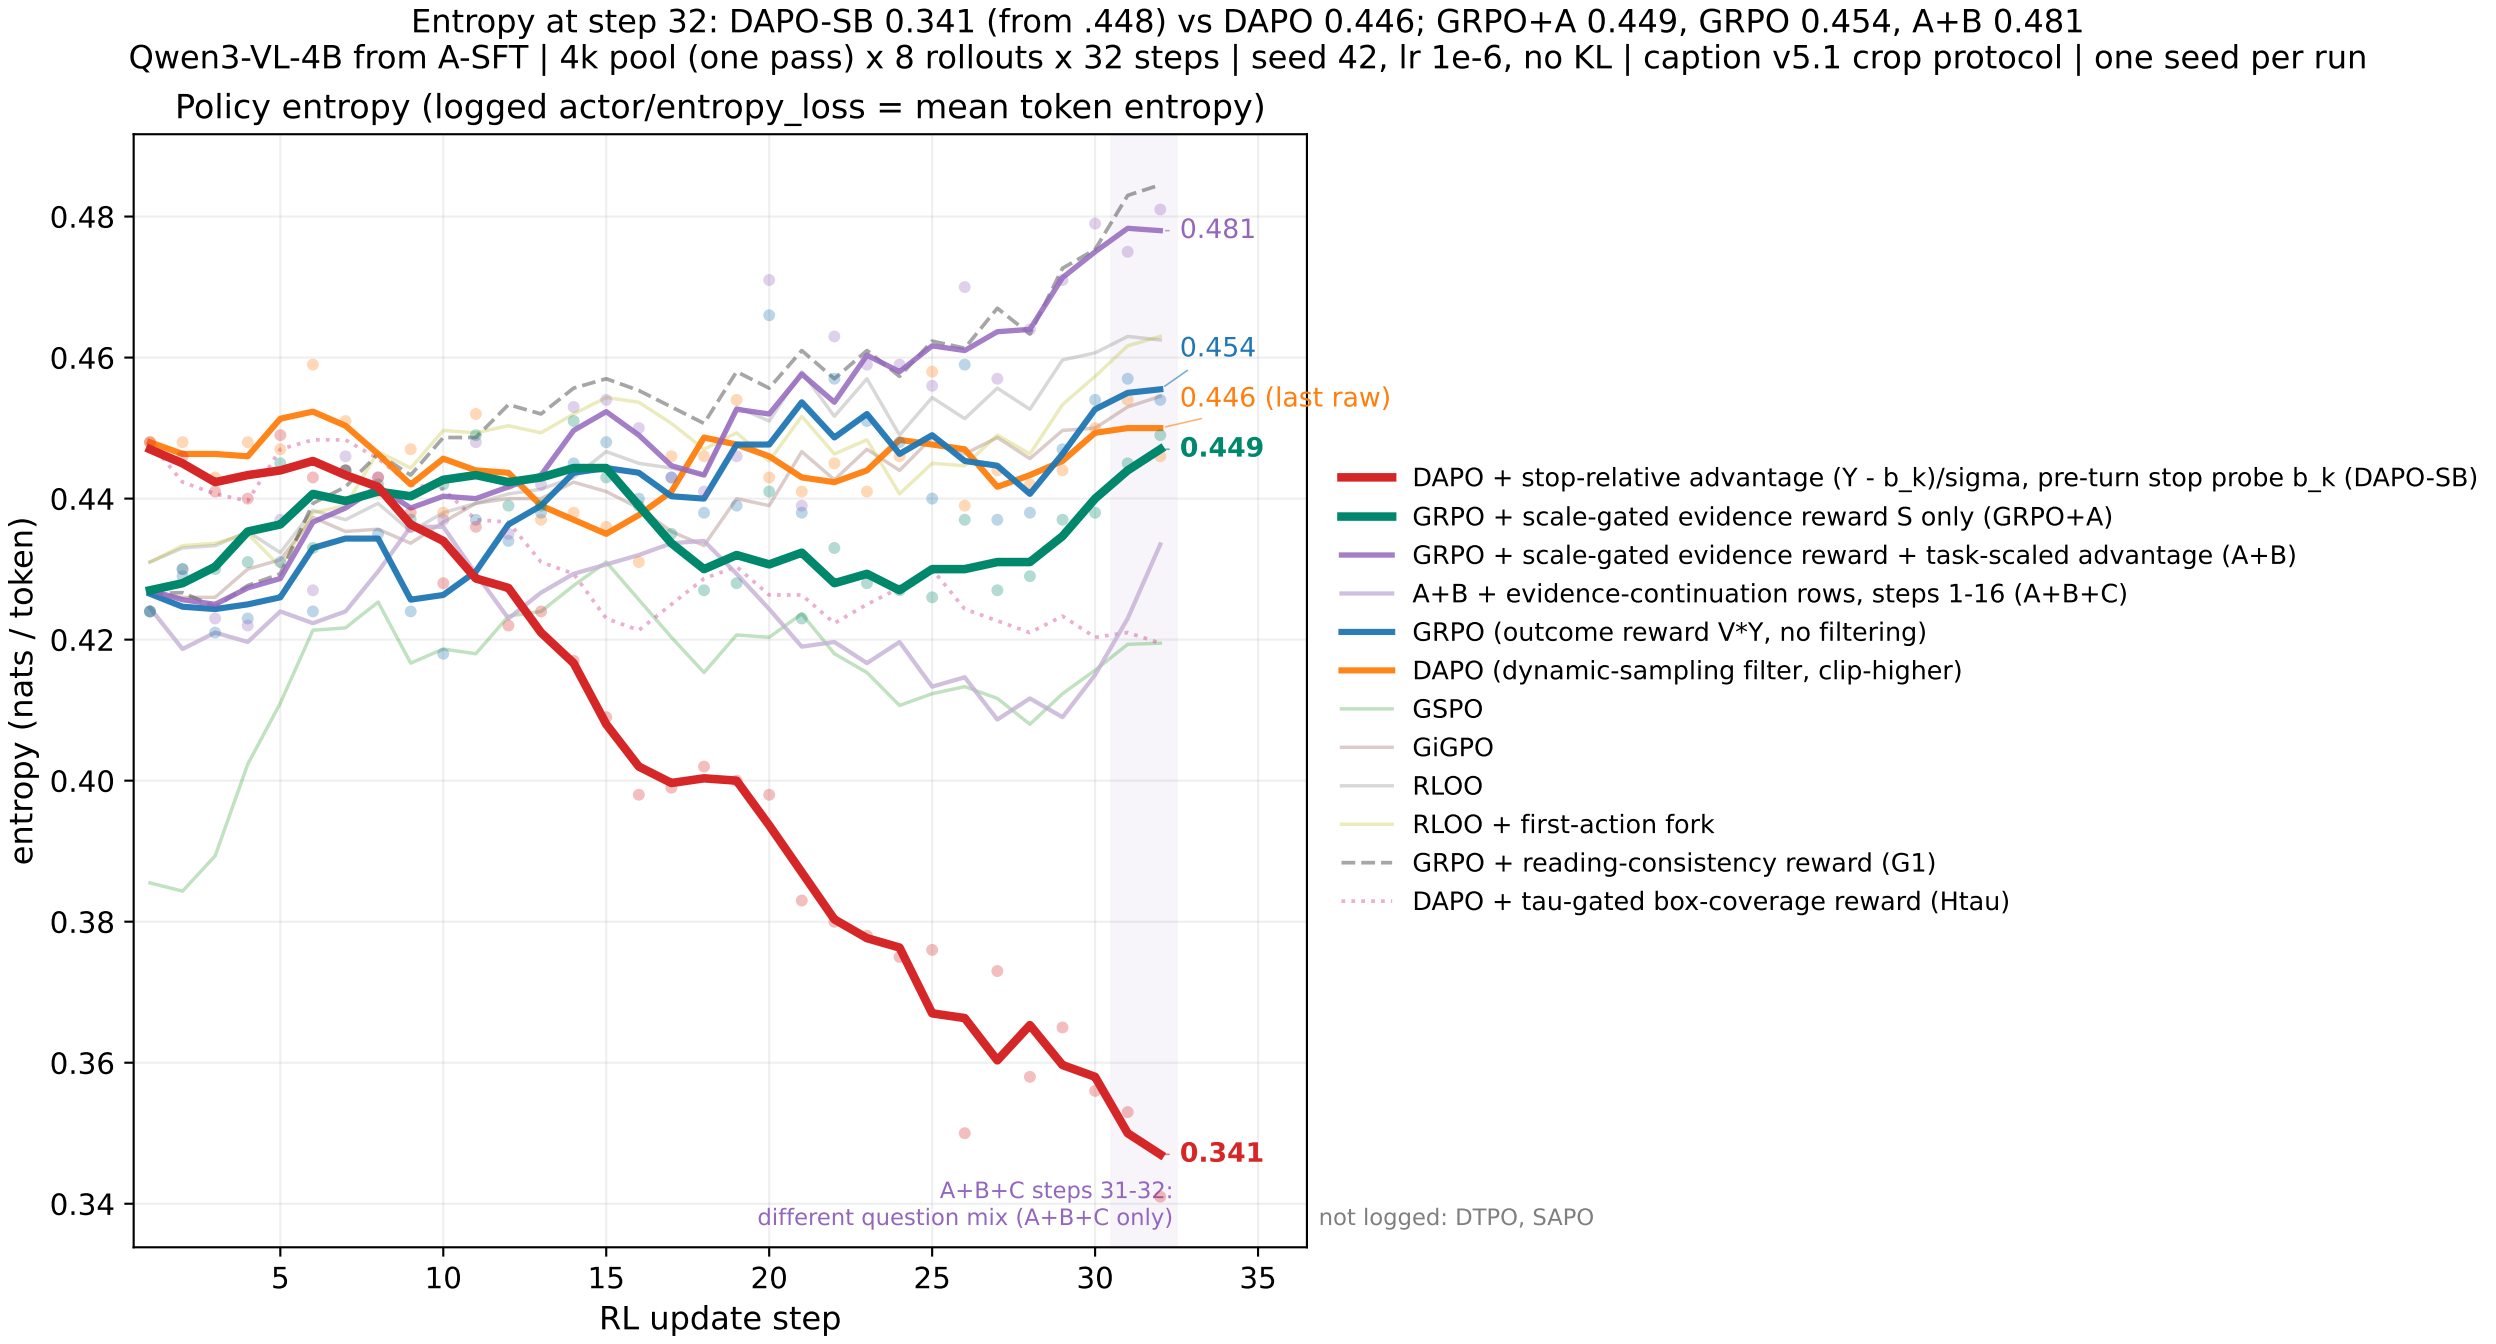
Entropy at step 32: DAPO-SB 0.341 (from .448) vs DAPO 0.446; GRPO+A 0.449, GRPO 0.454, A+B 0.481
Qwen3-VL-4B from A-SFT | 4k pool (one pass) x 8 rollouts x 32 steps | seed 42, lr 1e-6, no KL | caption v5.1 crop protocol | one seed per run
Policy entropy (logged actor/entropy_loss = mean token entropy)
entropy (nats / token)
RL update step
0.481
0.454
0.446 (last raw)
0.449
0.341
A+B+C steps 31-32:
different question mix (A+B+C only)
DAPO + stop-relative advantage (Y - b_k)/sigma, pre-turn stop probe b_k (DAPO-SB)
GRPO + scale-gated evidence reward S only (GRPO+A)
GRPO + scale-gated evidence reward + task-scaled advantage (A+B)
A+B + evidence-continuation rows, steps 1-16 (A+B+C)
GRPO (outcome reward V*Y, no filtering)
DAPO (dynamic-sampling filter, clip-higher)
GSPO
GiGPO
RLOO
RLOO + first-action fork
GRPO + reading-consistency reward (G1)
DAPO + tau-gated box-coverage reward (Htau)
not logged: DTPO, SAPO